

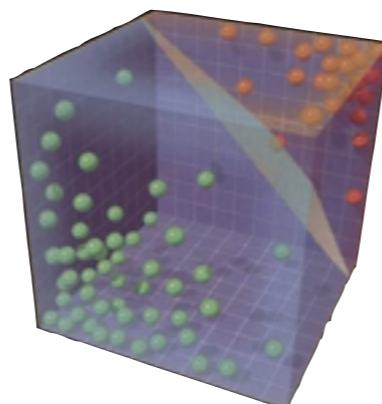


Classification

16-385 Computer Vision (Kris Kitani)
Carnegie Mellon University

typical perception pipeline

representation

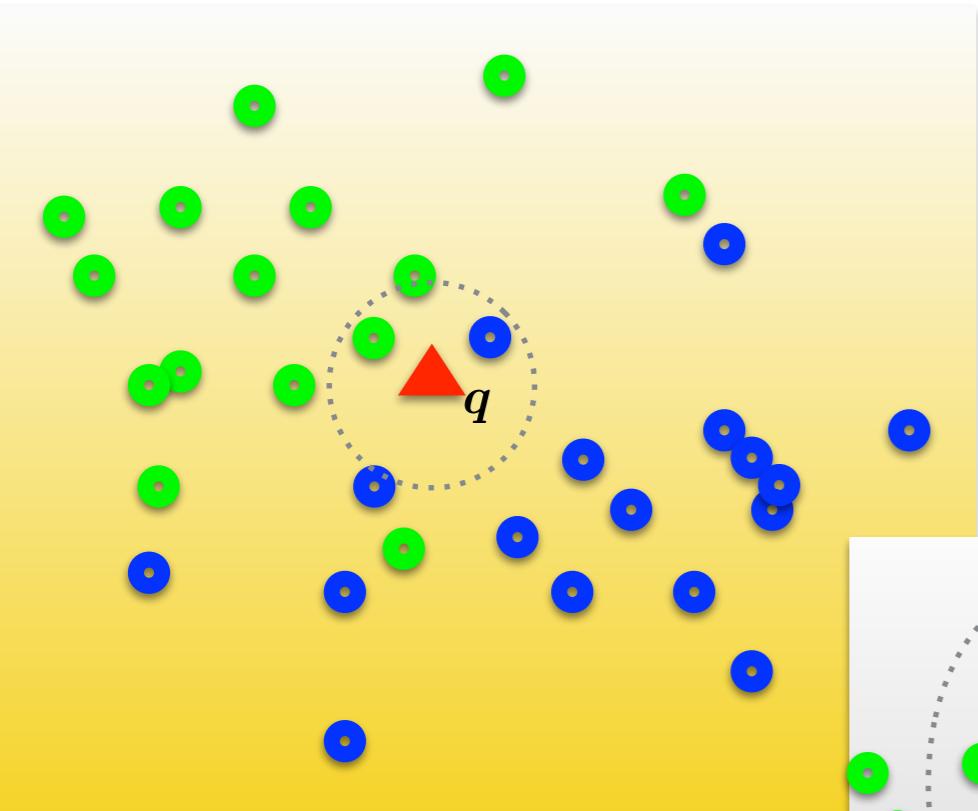


classifier

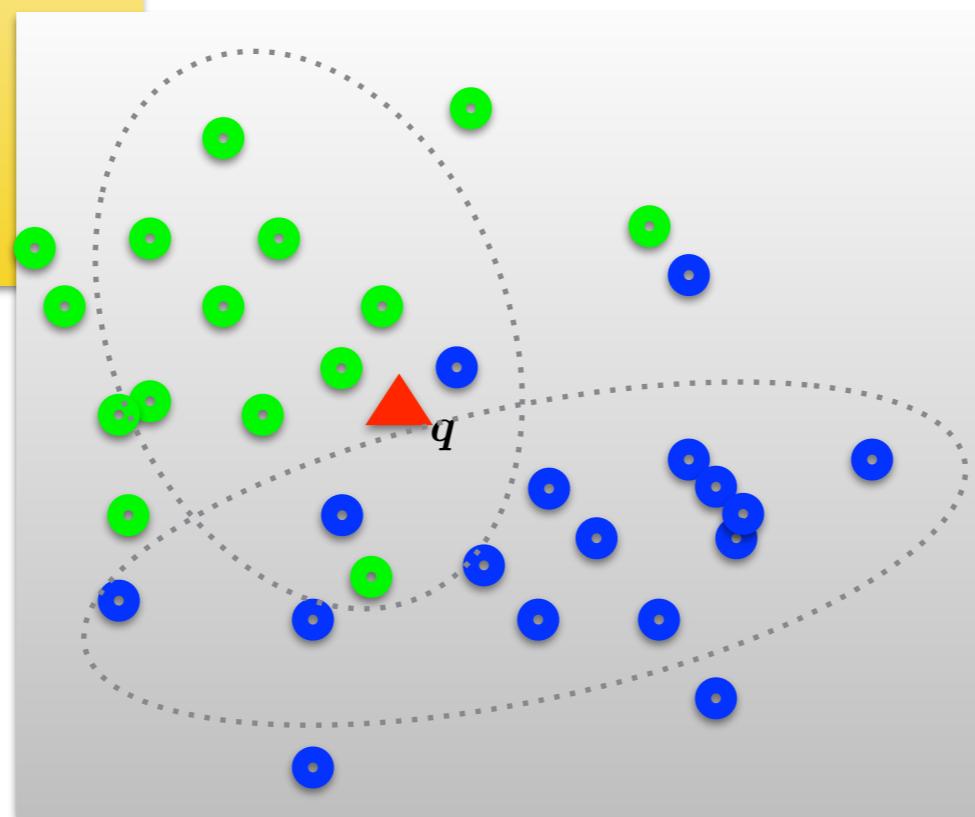


output

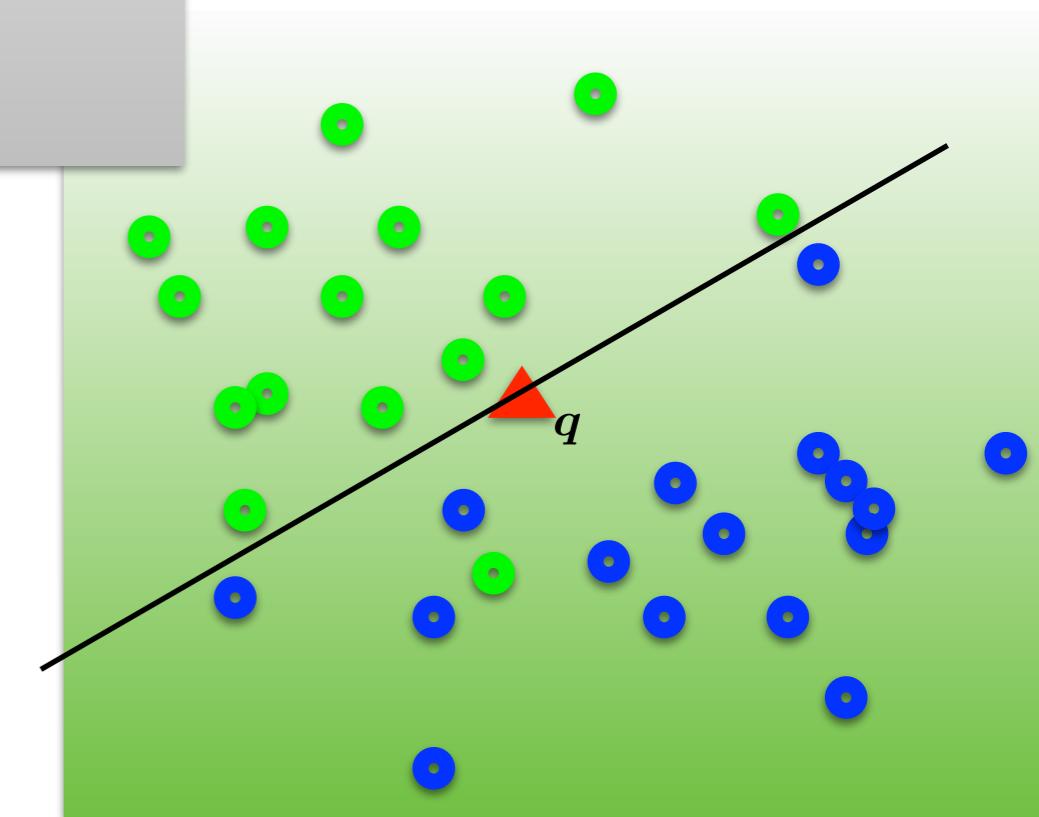
- Nearest Neighbor classifier
- Naive Bayes classifier
- Support Vector Machine



Nearest Neighbor

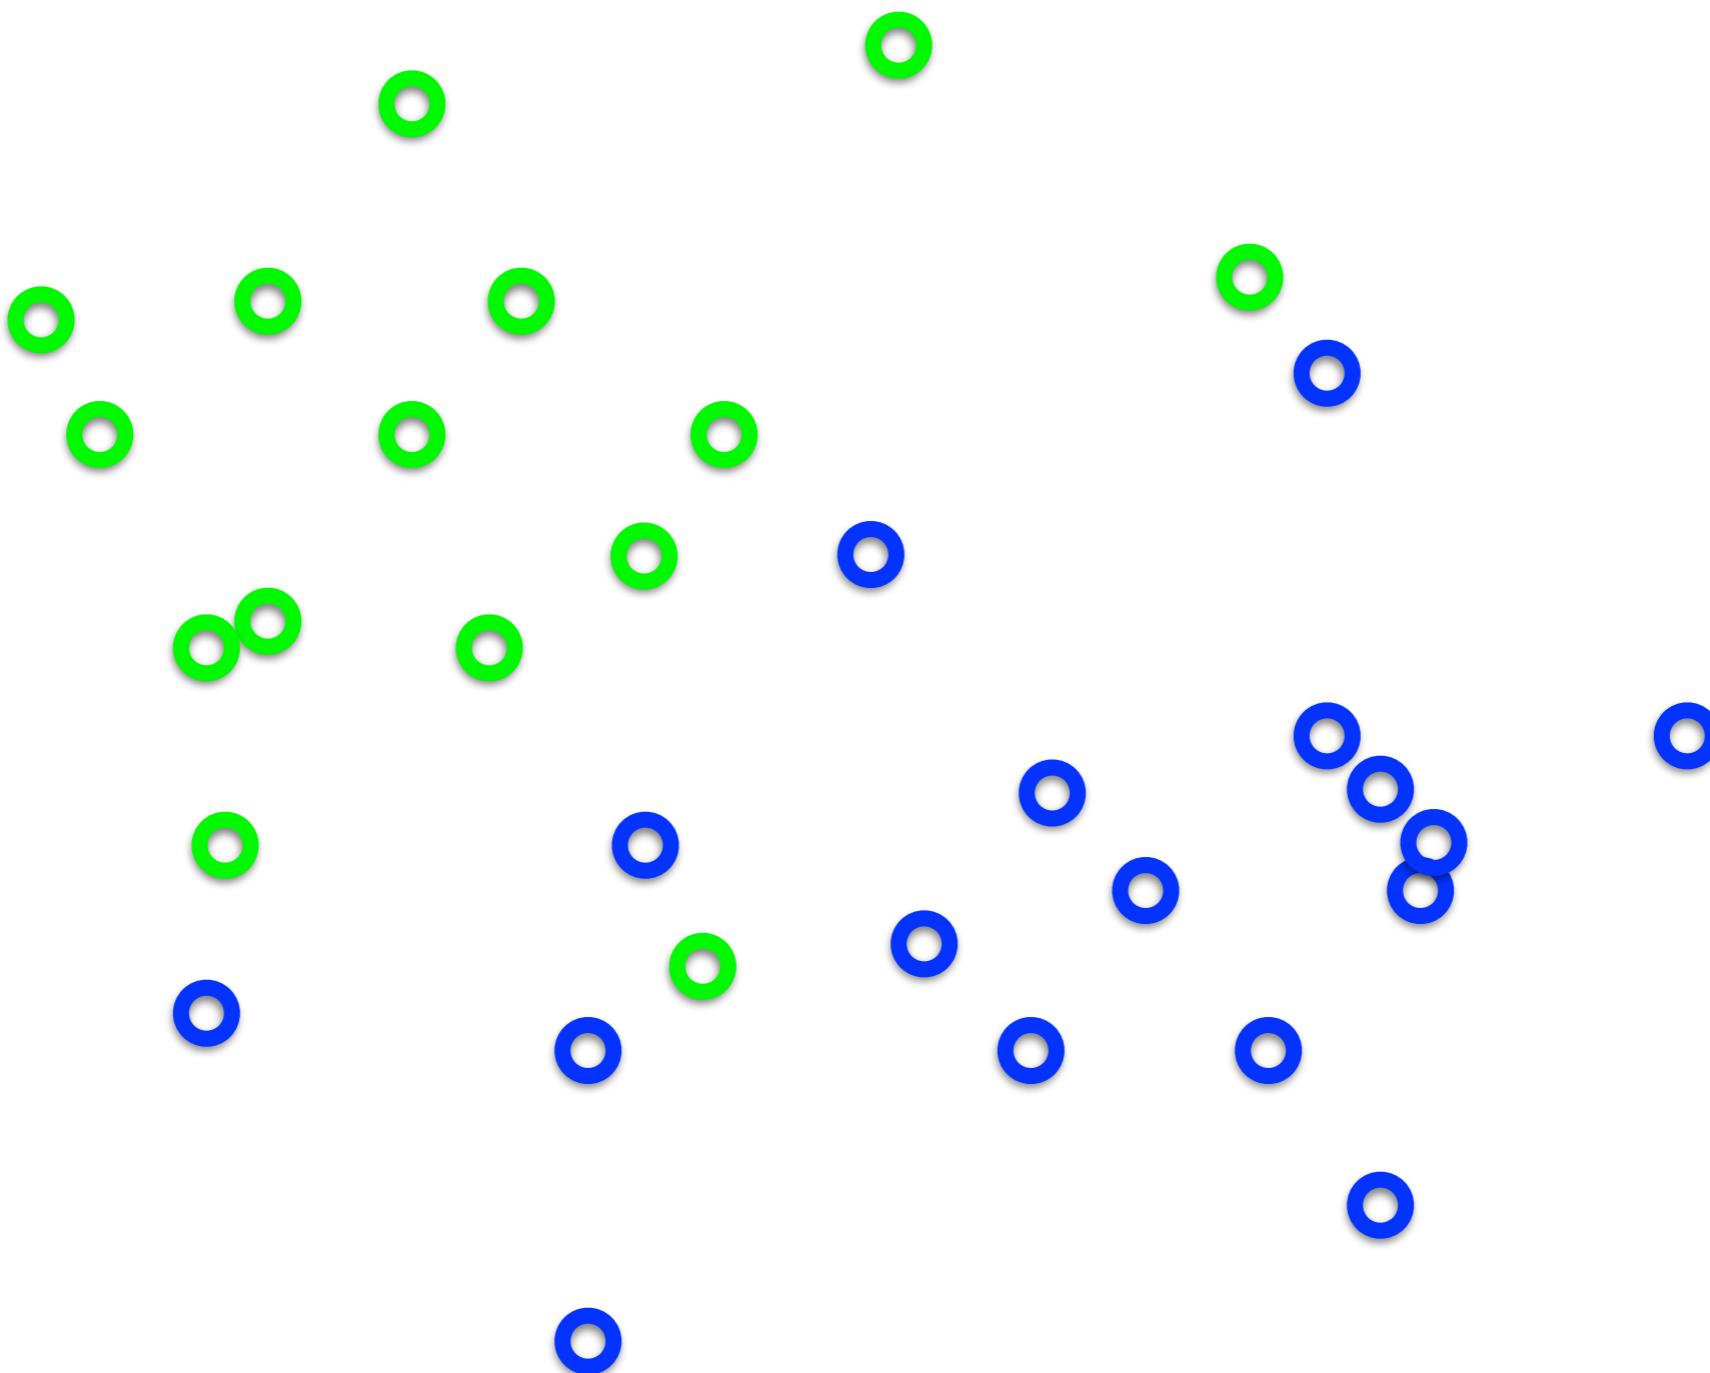


Naive Bayes

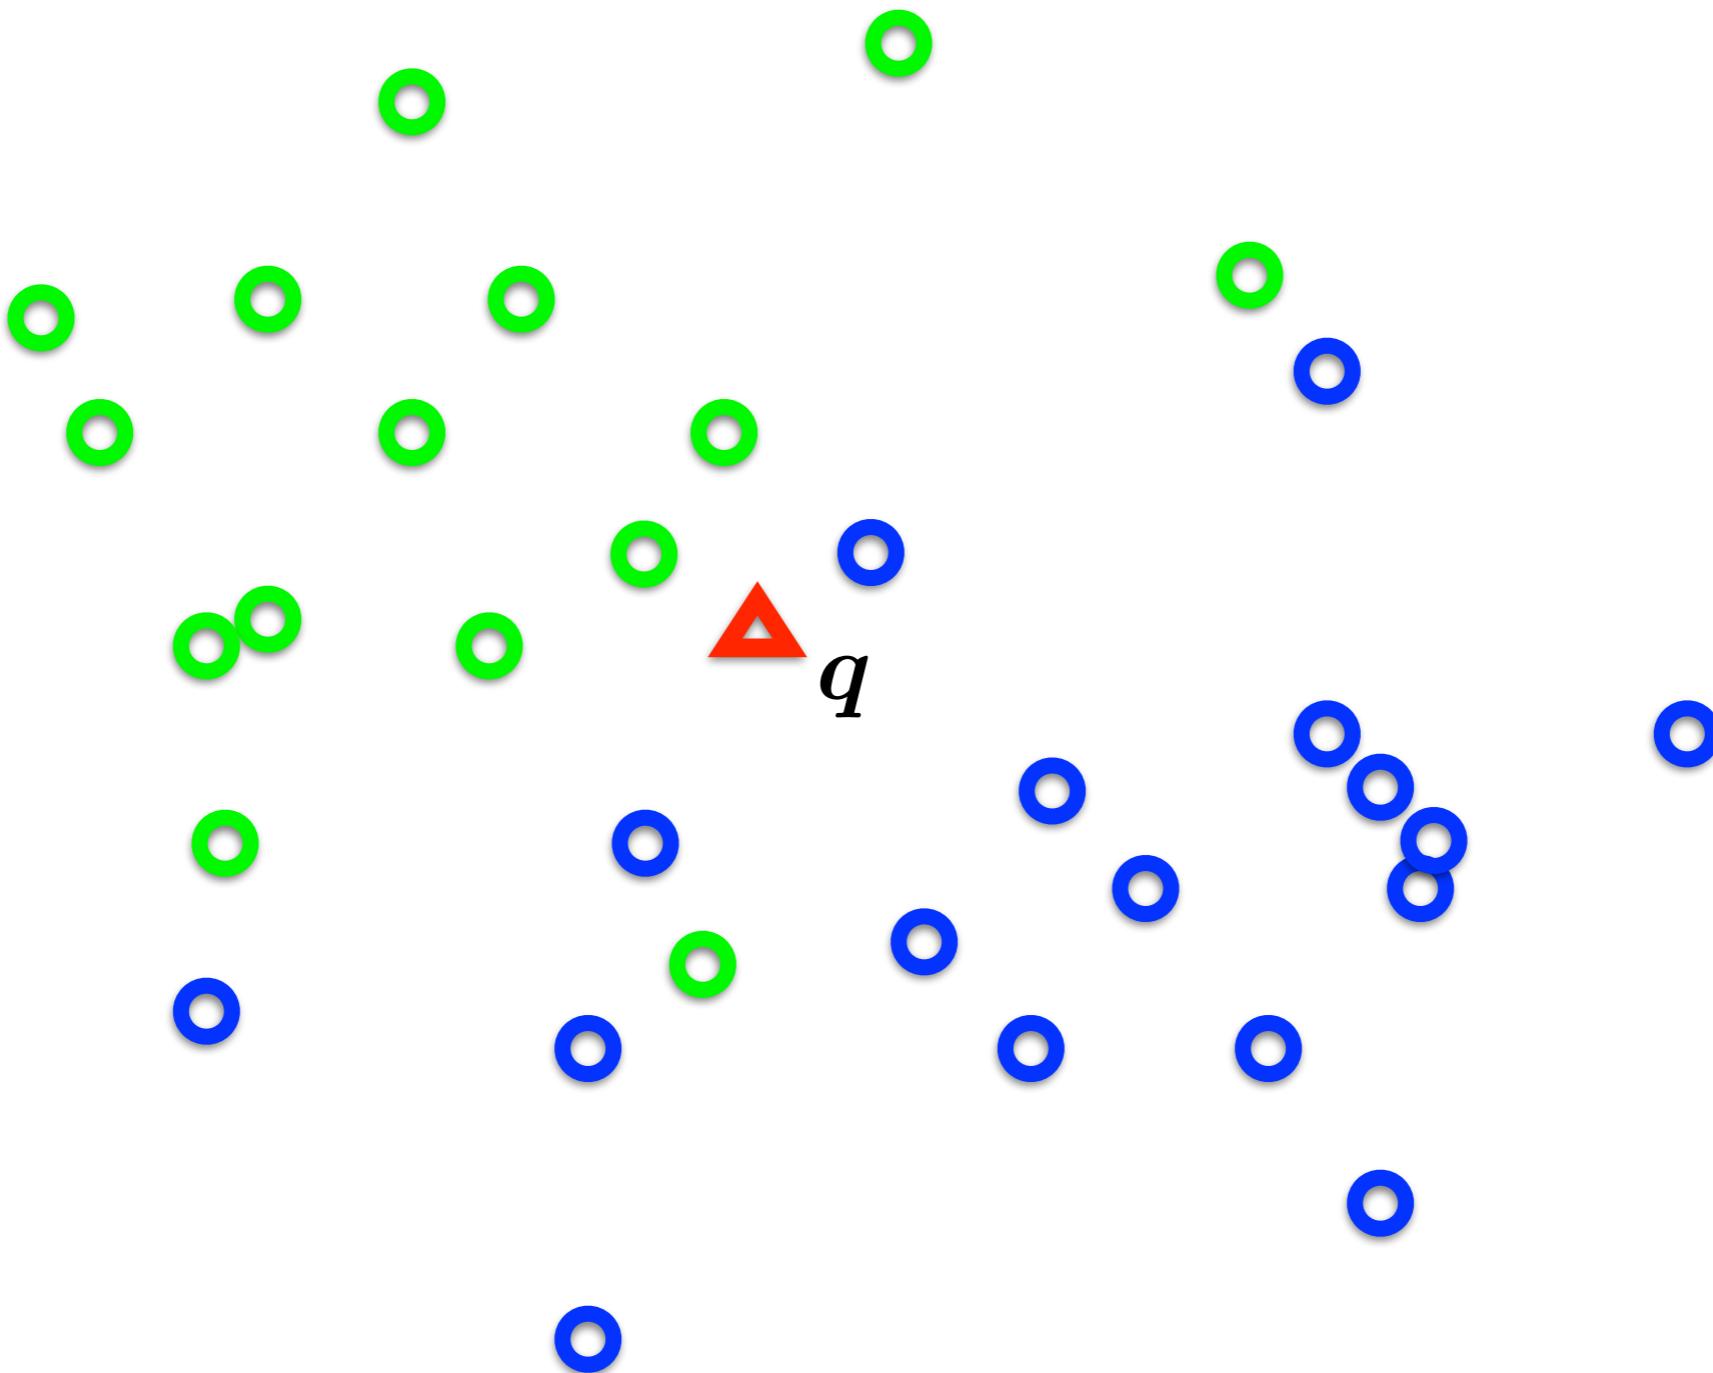


Support Vector Machine

Distribution of data from two classes

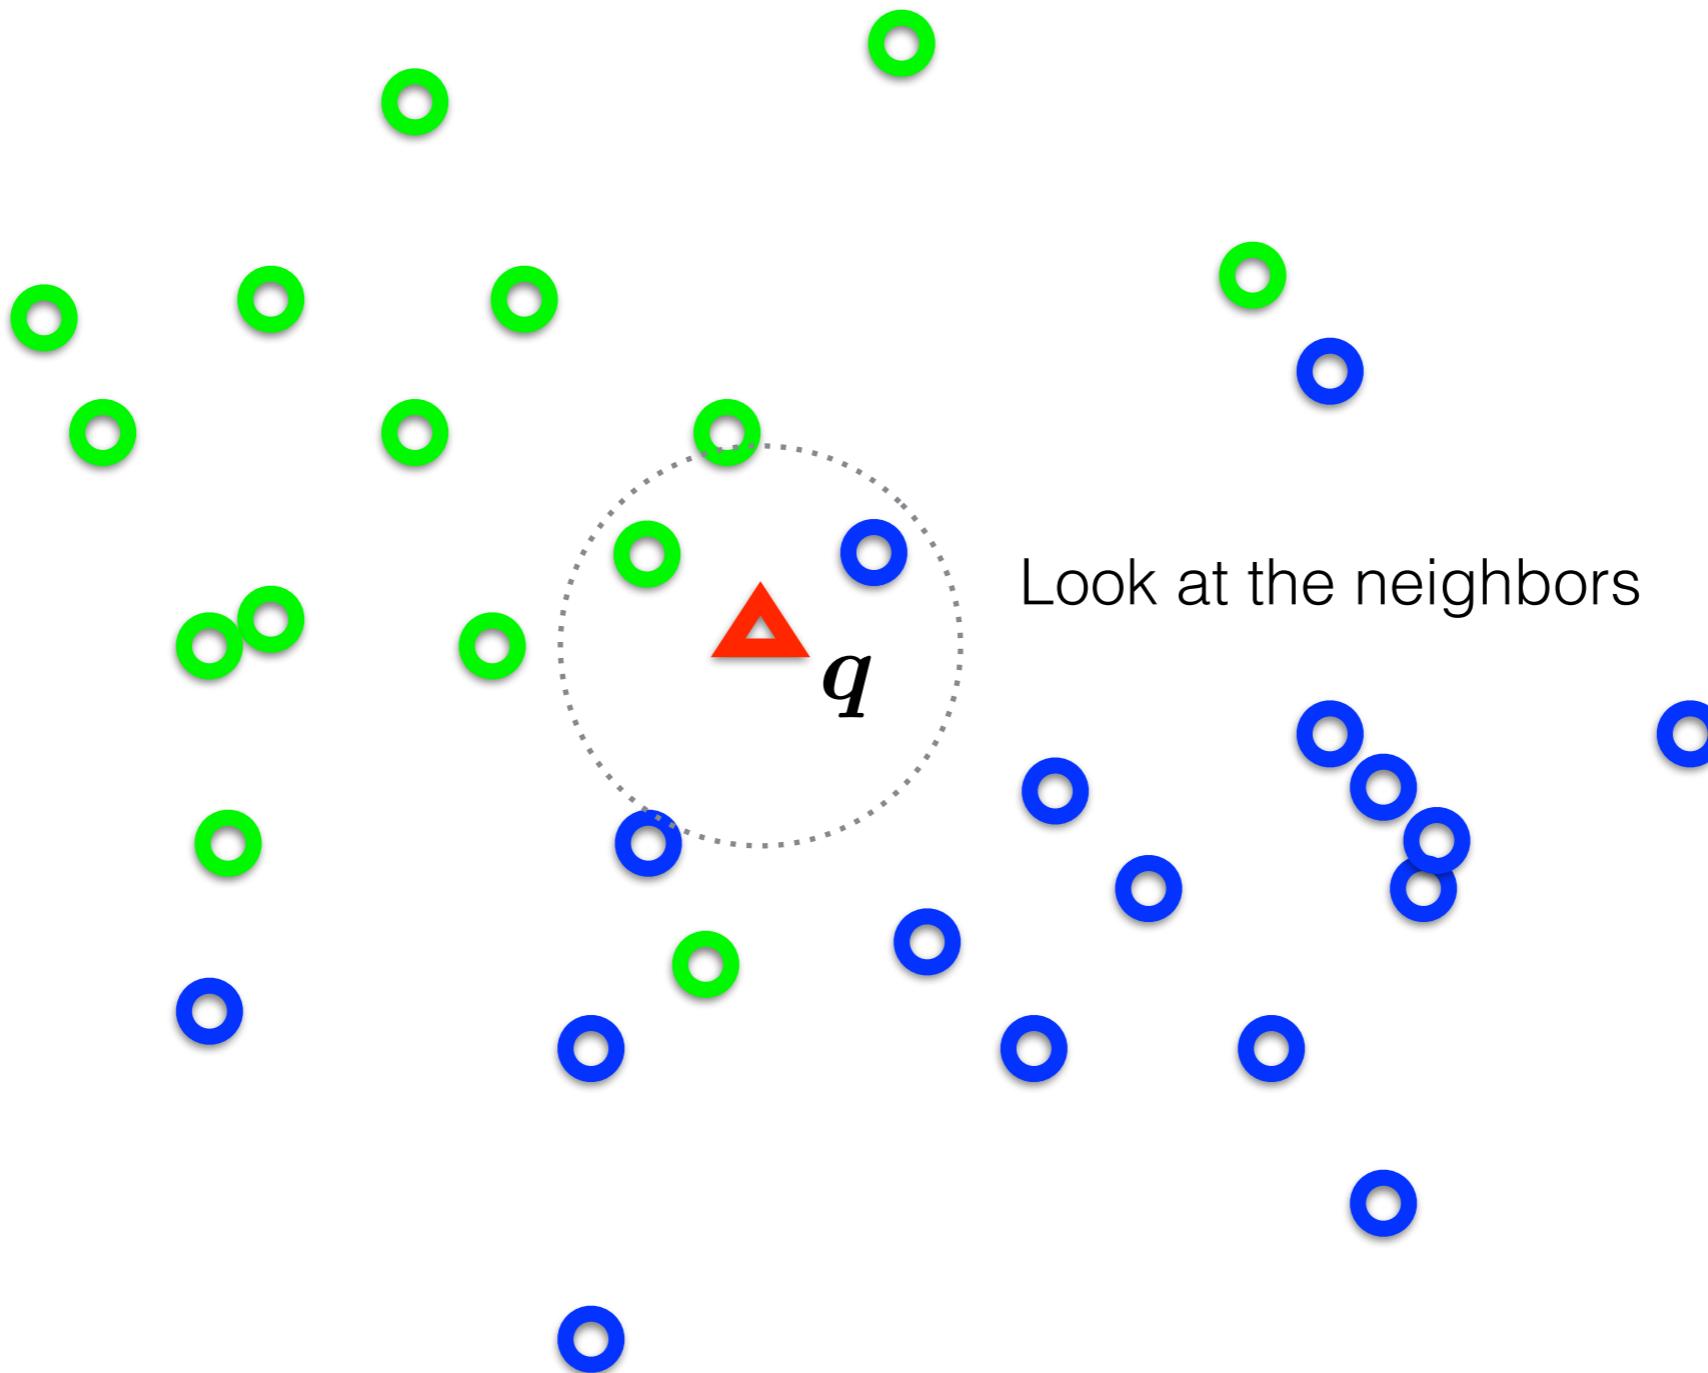


Distribution of data from two classes



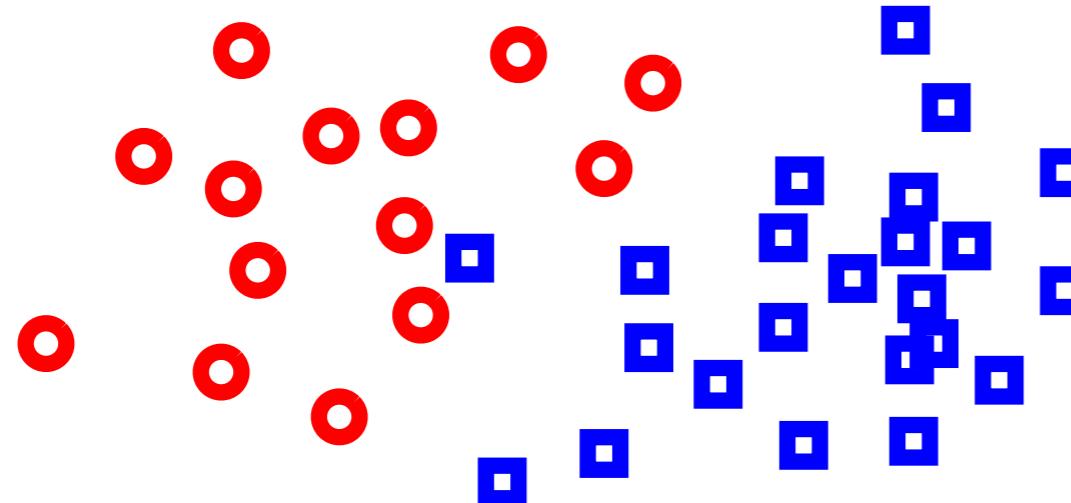
Which class does q belong to?

Distribution of data from two classes



K-nearest neighbor

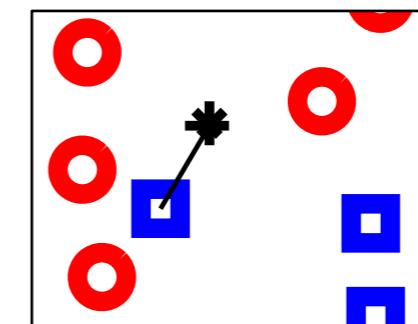
K-Nearest Neighbor (KNN) Classifier



Non-parametric pattern classification approach
Consider a two class problem where each sample consists of two measurements (x,y) .

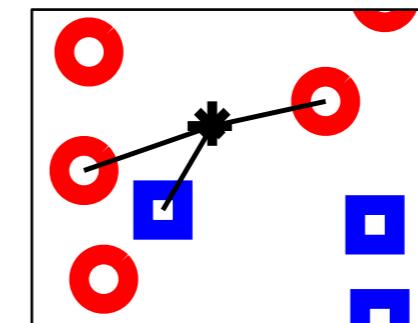
For a given query point q , assign the class of the nearest neighbor

$k = 1$



Compute the k nearest neighbors and assign the class by majority vote.

$k = 3$



Nearest Neighbor is competitive

4028150880327726495529284686500876/71127400776386420140578214711366
50711167679664143112241087634006330117113109975414895351982739901029
8468682467933943144705960444612336459685600641865284554770782237018
76953465018828357808571101378507110114527623028596972136418240510226
93977149069842728100783331376131605747595849916501320348220251514889
82049962335648092836457294912860704116759914592504108408989425198980
35517216919955162286714604033223689853854520563283995794671313660901
94368160413174951001162198403649071657525185470670258104571851900607
8857389886823975629288168879180172075190209862393802111142972512199
148534347507488153959369036398212868553941251514414435912233029009
9319097549201051493361525220266012030255795808950325908846884546549
69285457999216340783934656219260061287992047750564674307507426899404
12845278113035703193631773084826529739099642972116747590821445161325
90662367728608302983253980019513960141712379749939282718091017796999
2101045282835177112976405078847858498138031715516574935471208160734
283087884084458566309376893495891288681379011470817457121130621280766
41992780136134111560707232522949810/61278000822722799275134941756283

MNIST Digit Recognition

- Handwritten digits
- 28x28 pixel images: $d = 784$
- 60,000 training samples
- 10,000 test samples

Yann LeCunn

	Test Error Rate (%)
Linear classifier (1-layer NN)	12.0
K-nearest-neighbors, Euclidean	5.0
K-nearest-neighbors, Euclidean, deskewed	2.4
K-NN, Tangent Distance, 16x16	1.1
K-NN, shape context matching	0.67
1000 RBF + linear classifier	3.6
SVM deg 4 polynomial	1.1
2-layer NN, 300 hidden units	4.7
2-layer NN, 300 HU, [deskewing]	1.6
LeNet-5, [distortions]	0.8
Boosted LeNet-4, [distortions]	0.7

Pros

- simple yet effective

Cons

- search is expensive (can be sped-up)
- storage requirements
- difficulties with high-dimensional data

What is the best distance metric between data points?

- Typically Euclidean distance
- Locality sensitive distance metrics
- Important to normalize.
Dimensions have different scales

How many K?

- Typically $k=1$ is good
- Cross-validation

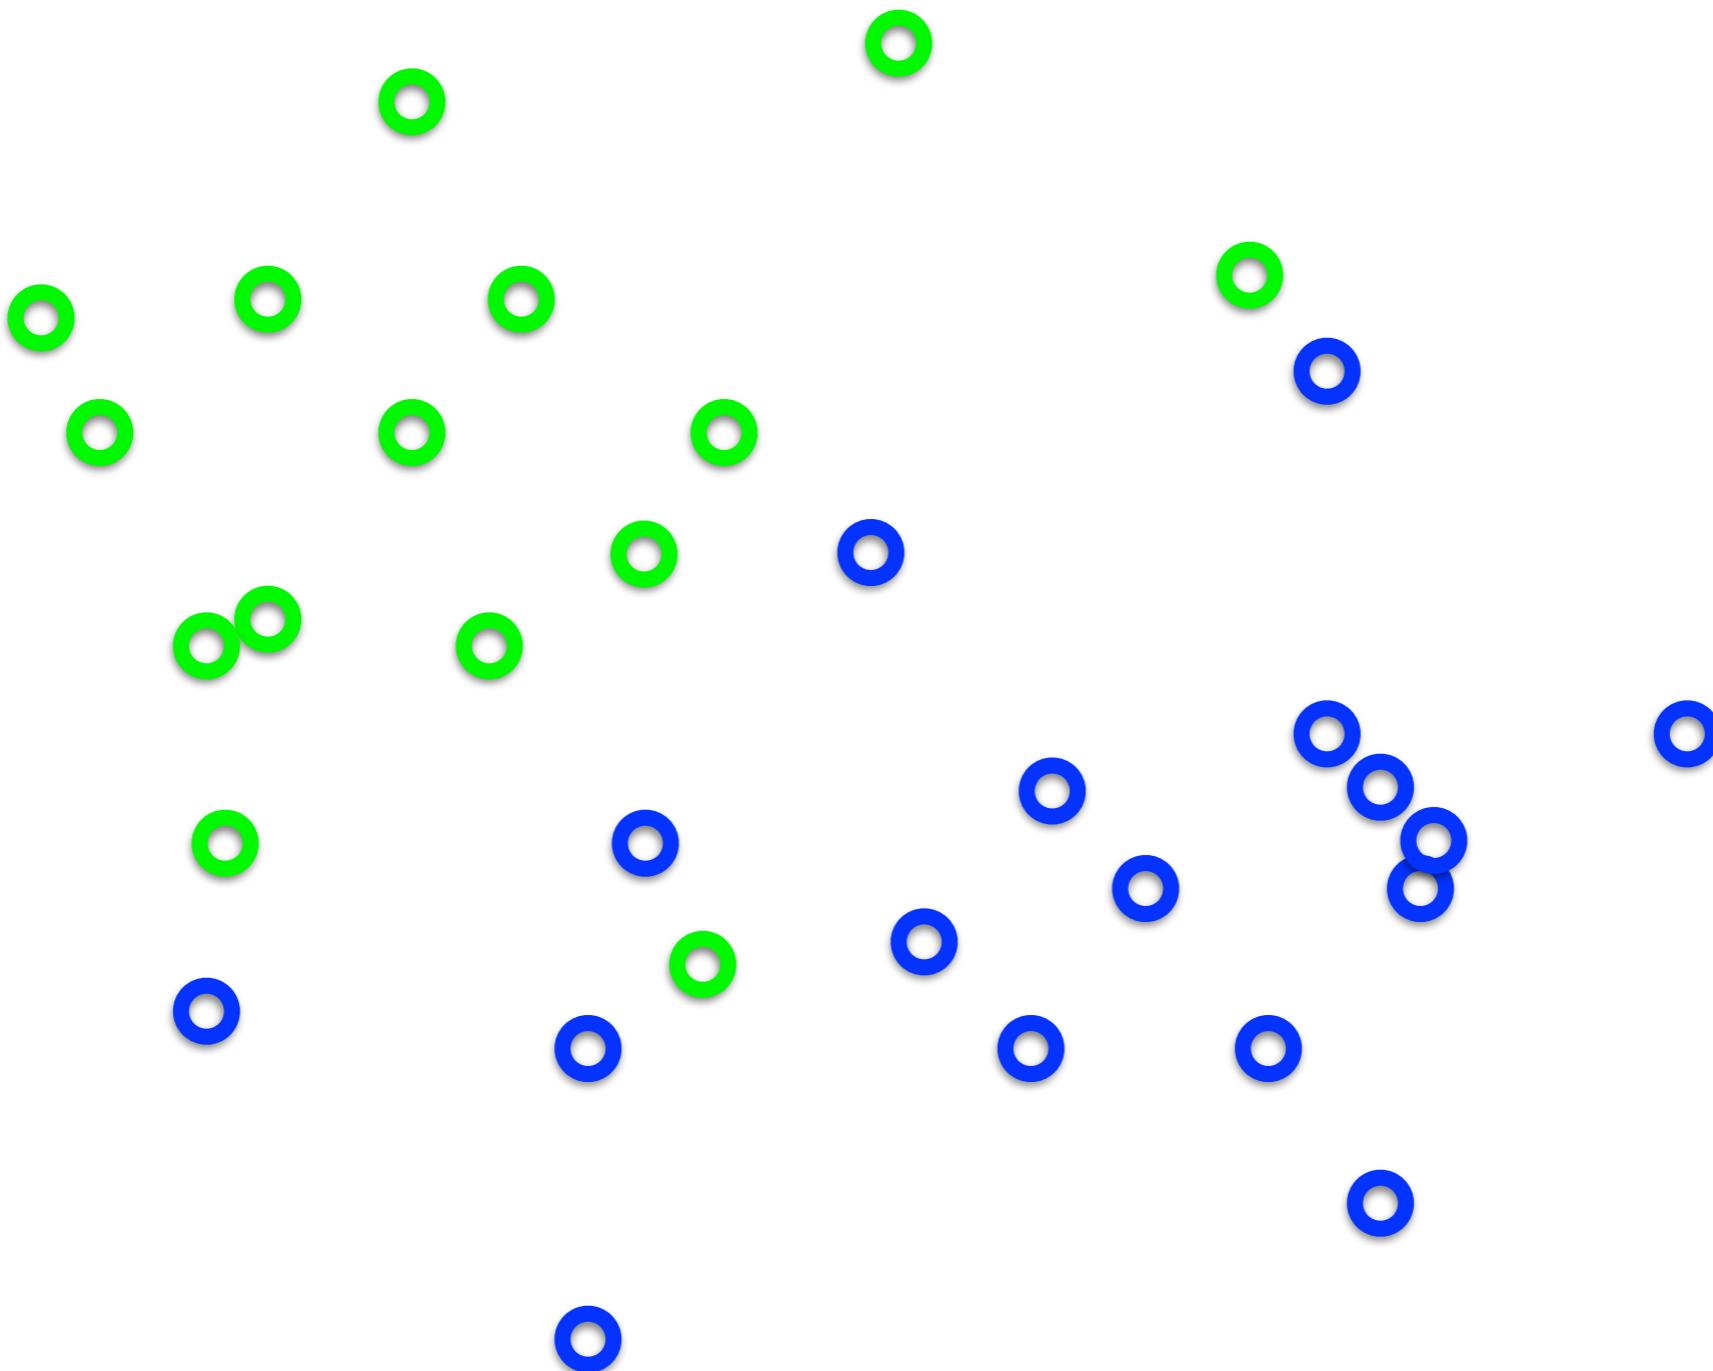
Distance metrics

$$D(\mathbf{x}, \mathbf{y}) = \sqrt{(x_1 - y_1)^2 + \cdots + (x_N - y_N)^2} \quad \text{Euclidean}$$

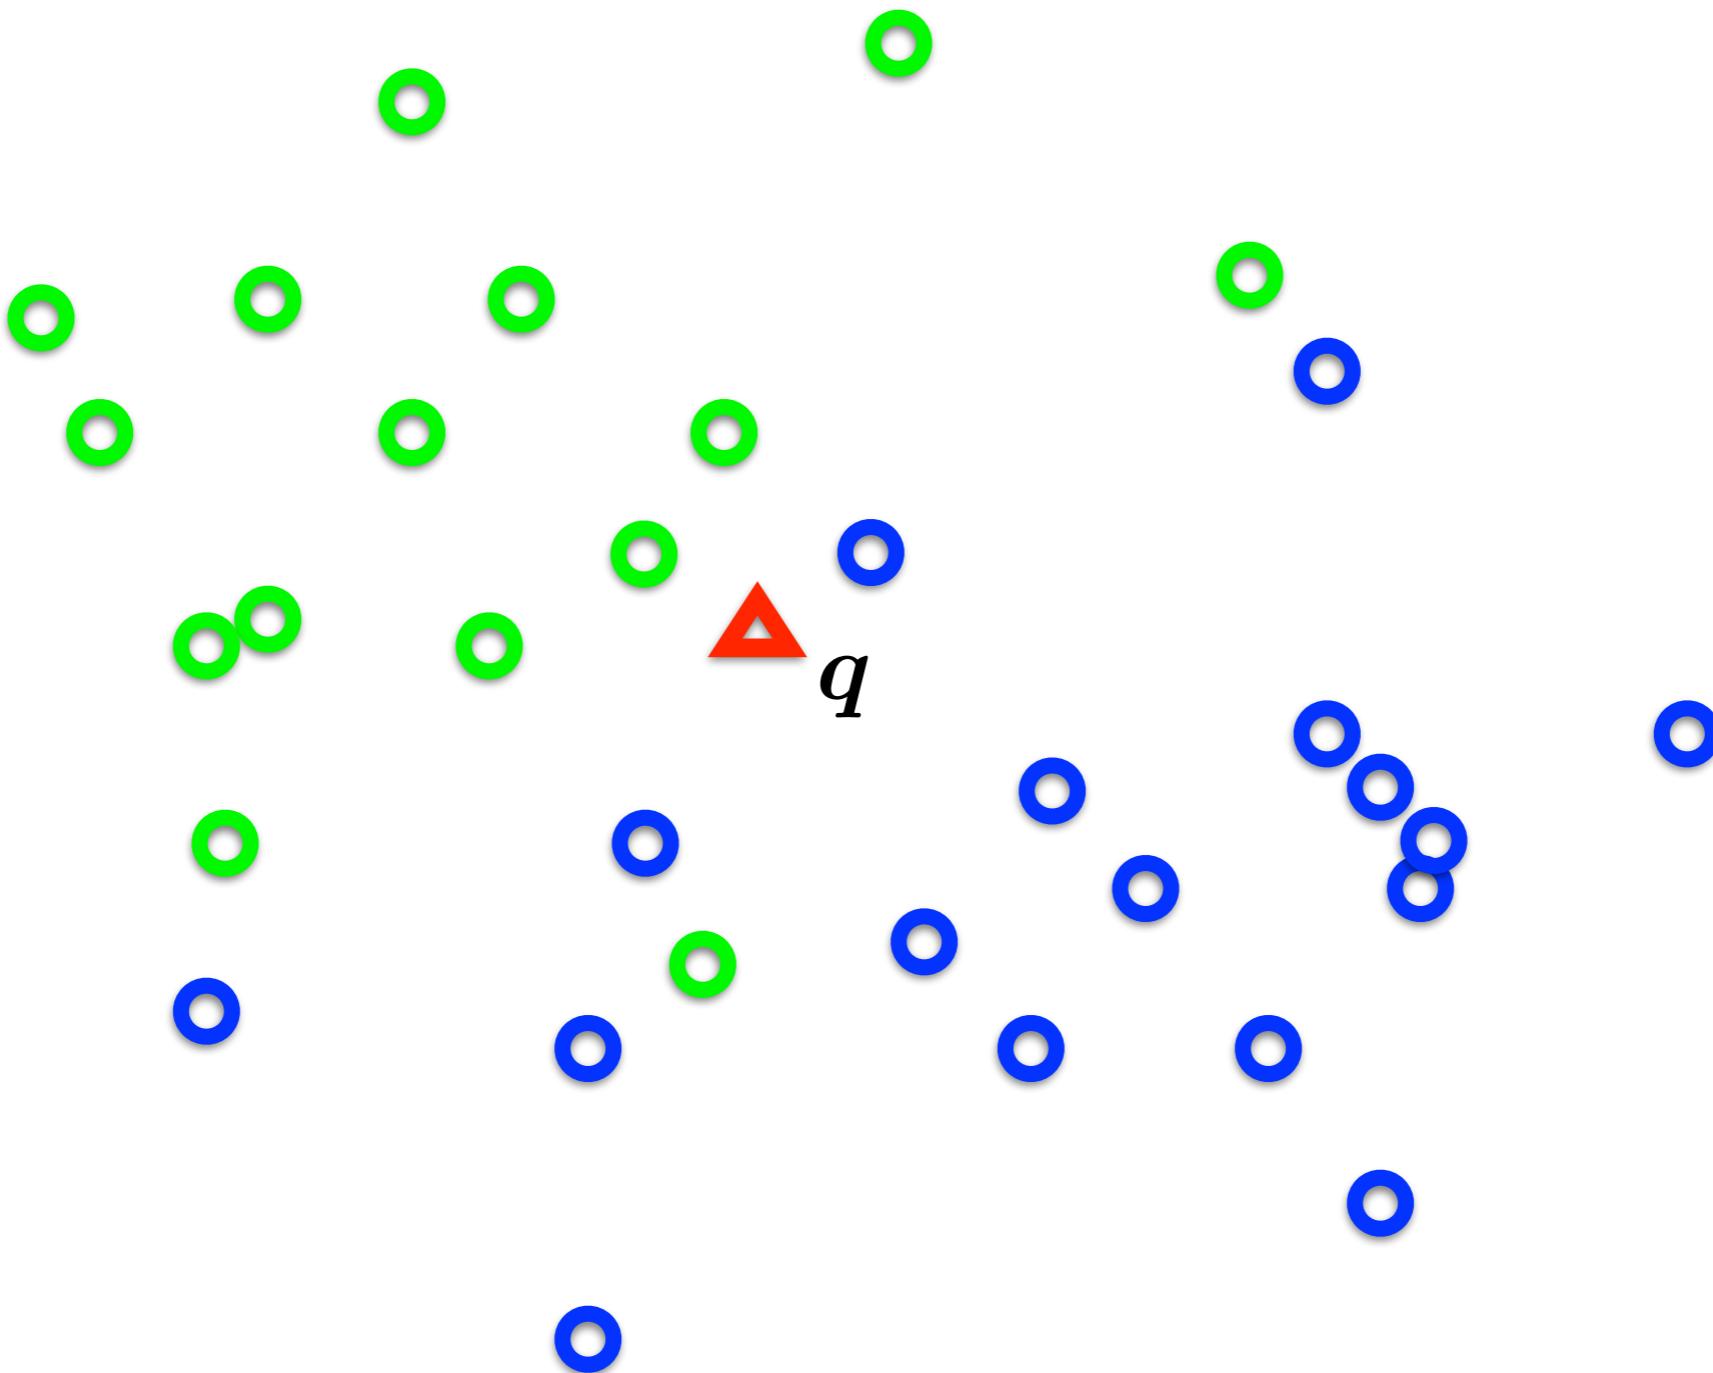
$$D(\mathbf{x}, \mathbf{y}) = \frac{\mathbf{x} \cdot \mathbf{y}}{\|\mathbf{x}\| \|\mathbf{y}\|} = \frac{x_1 y_1 + \cdots + x_N y_N}{\sqrt{\sum_n x_n^2} \sqrt{\sum_n y_n^2}} \quad \text{Cosine}$$

$$D(\mathbf{x}, \mathbf{y}) = \frac{1}{2} \sum_n \frac{(x_n - y_n)^2}{(x_n + y_n)} \quad \text{Chi-squared}$$

Distribution of data from two classes

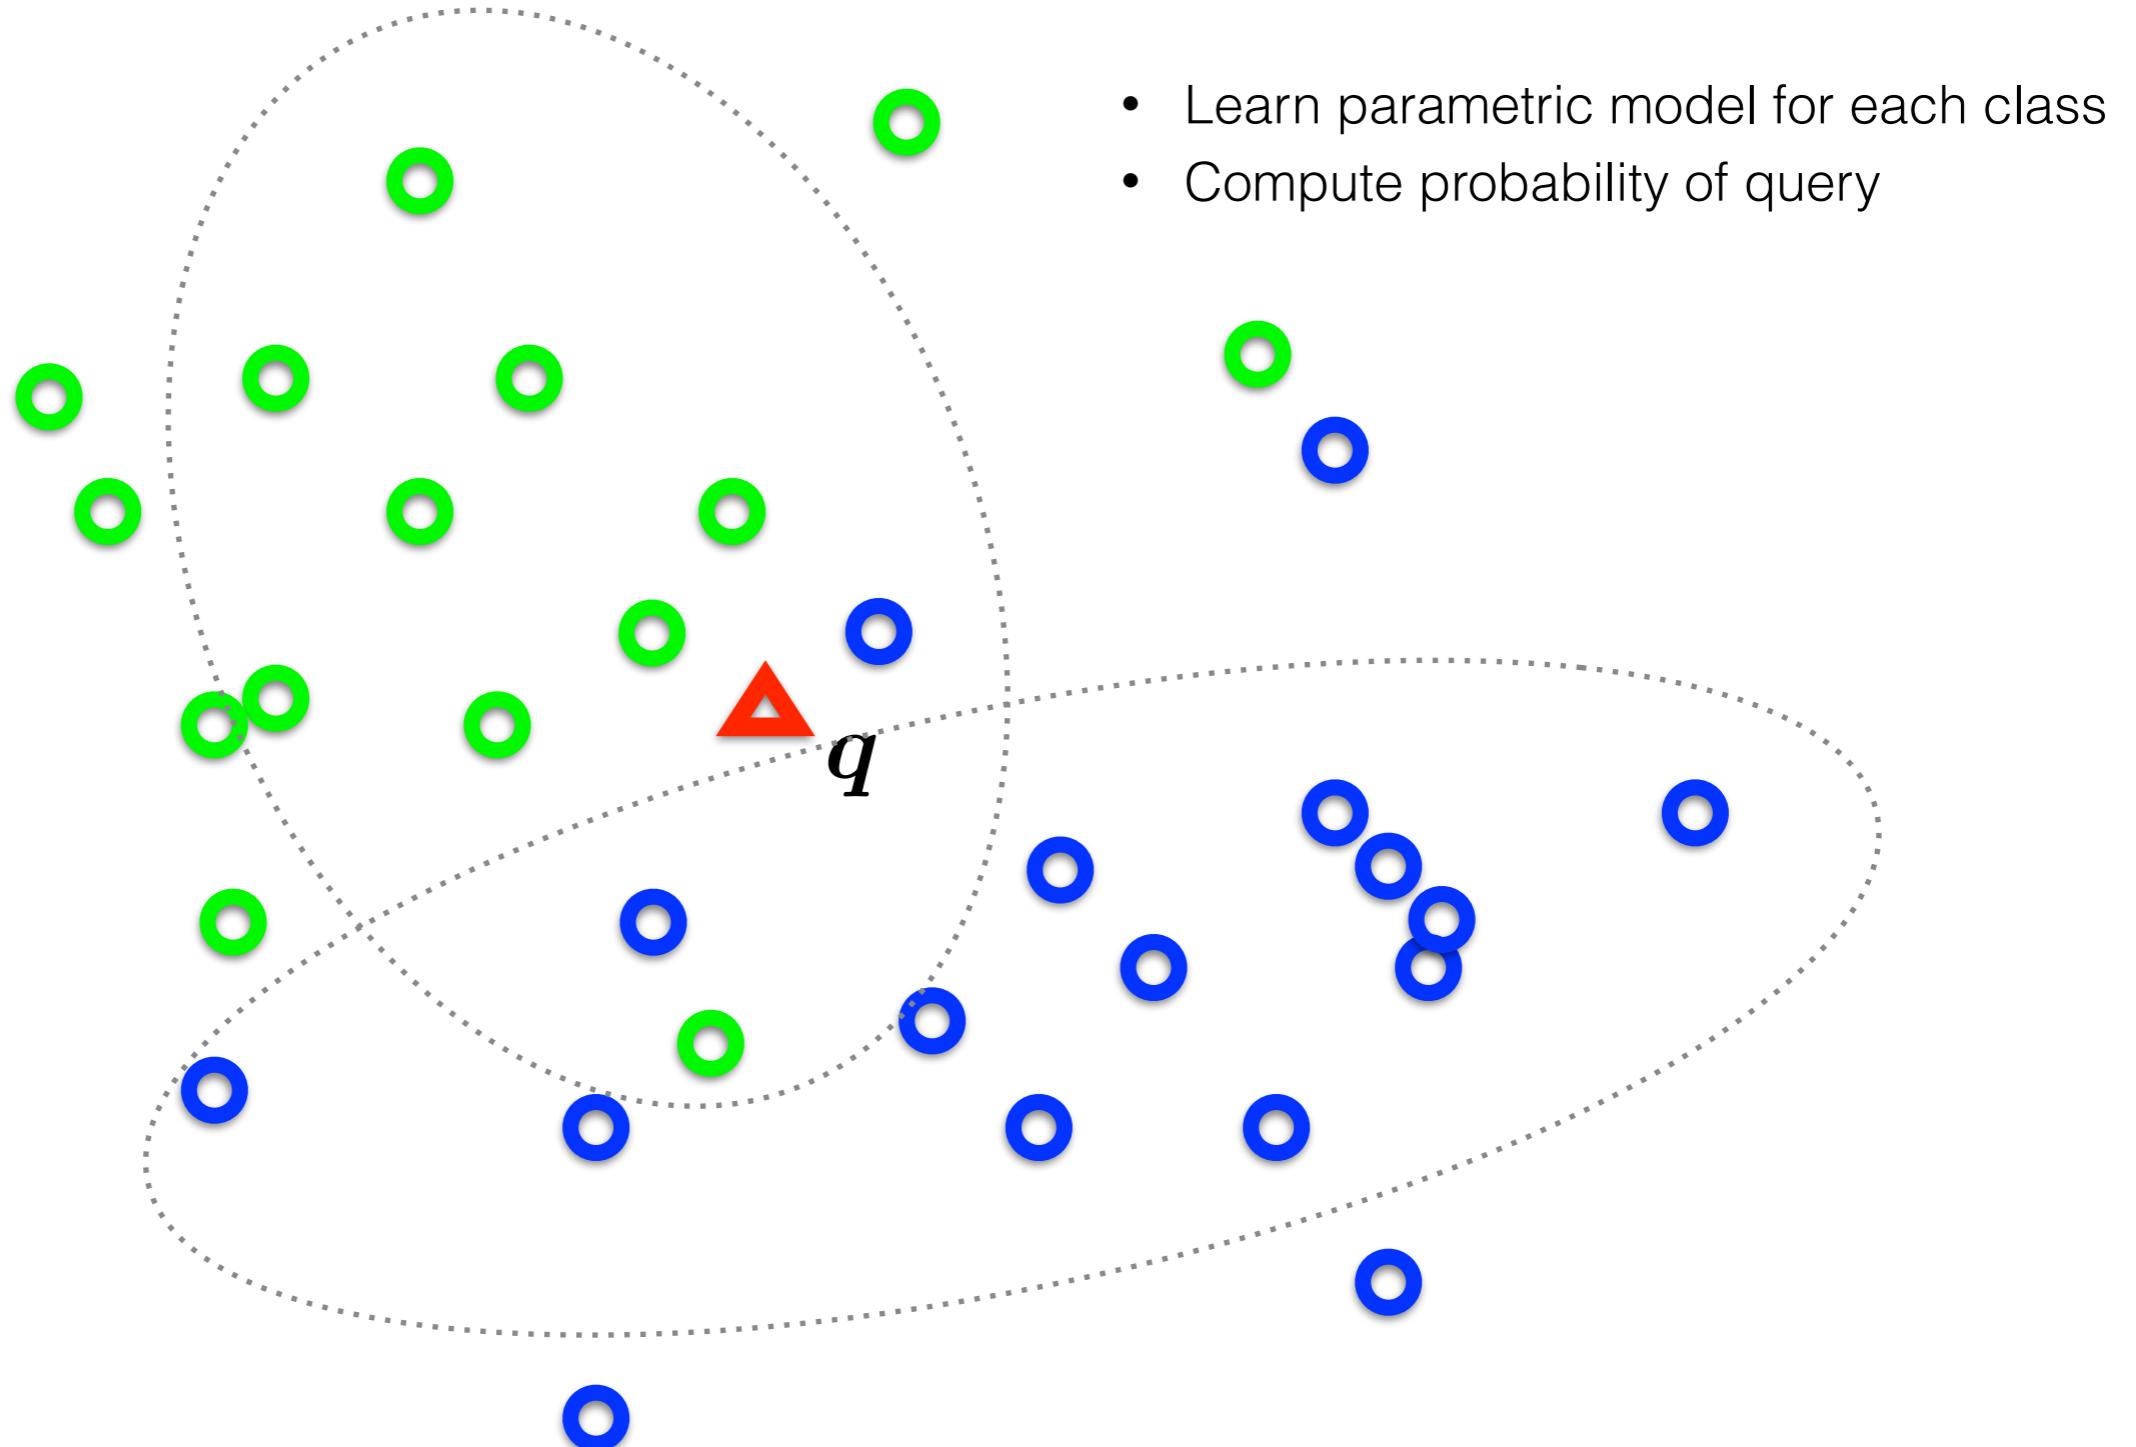


Distribution of data from two classes



Which class does q belong to?

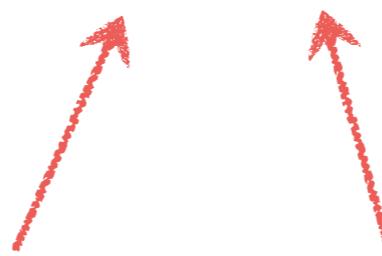
Distribution of data from two classes



Naive Bayes

This is called the posterior:
the probability of a class z given the observed features X

$$p(z|X)$$



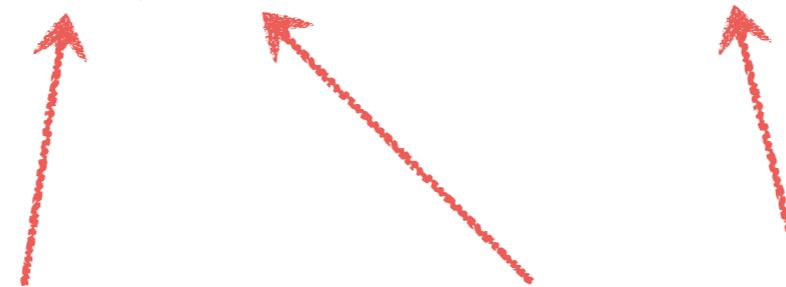
For classification, z is a
discrete random variable
(e.g., car, person, building)

X is a set of observed feature
(e.g., features from a single image)

(it's a function that returns a single probability value)

This is called the posterior:
the probability of a class z given the observed features X

$$p(z|x_1, \dots, x_N)$$



For classification, z is a
discrete random variable
(e.g., car, person, building)

Each x is an observed feature
(e.g., visual words)

(it's a function that returns a single probability value)

Recall:

The posterior can be decomposed according to
Bayes' Rule

$$p(A|B) = \frac{\text{likelihood} \quad \text{prior}}{p(B)}$$

posterior

In our context...

$$p(z|x_1, \dots, x_N) = \frac{p(x_1, \dots, x_N | z)p(z)}{p(x_1, \dots, x_N)}$$

The naive Bayes' classifier is solving this optimization

$$\hat{z} = \arg \max_{z \in \mathcal{Z}} p(z | \mathbf{X})$$

MAP (maximum a posteriori) estimate

$$\hat{z} = \arg \max_{z \in \mathcal{Z}} \frac{p(\mathbf{X}|z)p(z)}{p(\mathbf{X})} \quad \text{Bayes' Rule}$$

$$\hat{z} = \arg \max_{z \in \mathcal{Z}} p(\mathbf{X}|z)p(z) \quad \text{Remove constants}$$

To optimize this...

$$\hat{z} = \arg \max_{z \in \mathcal{Z}} p(z|X)$$

We need to compute this



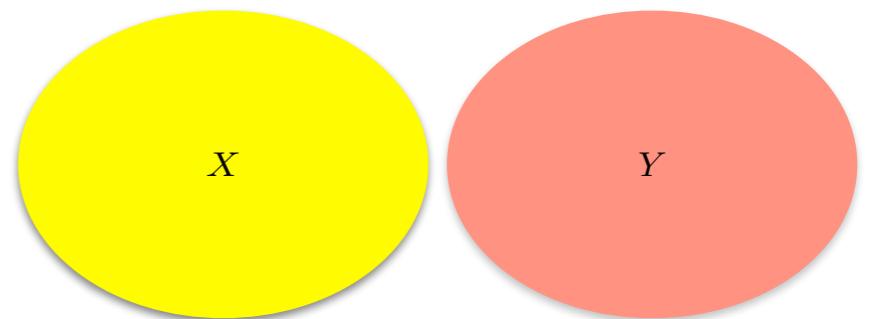
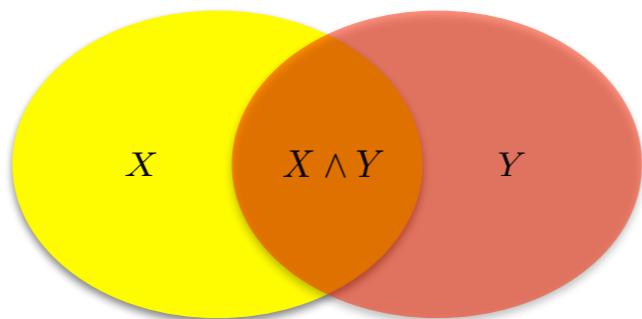
$$p(z|x_1, \dots, x_N) = \frac{p(x_1, \dots, x_N | z)p(z)}{p(x_1, \dots, x_N)}$$

Compute the likelihood...

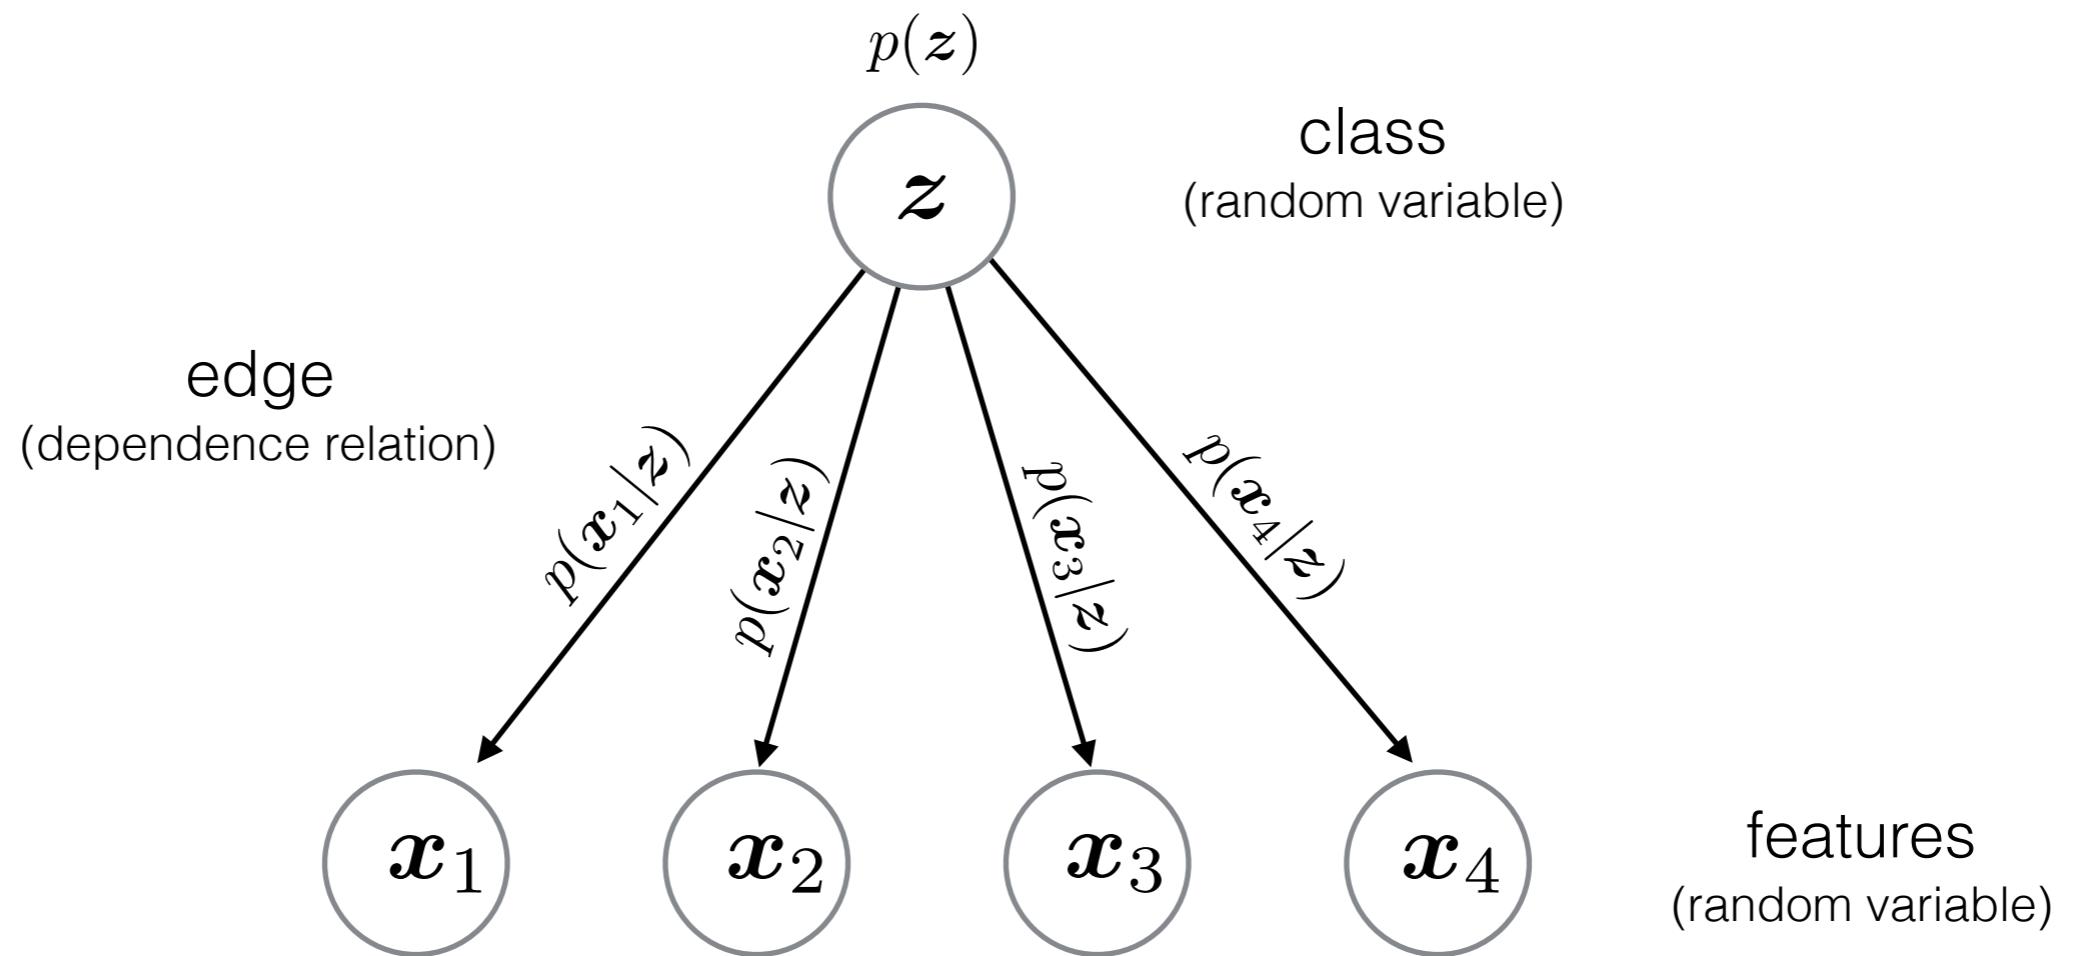
A naive Bayes' classifier assumes all features are
conditionally independent

$$\begin{aligned} p(\mathbf{x}_1, \dots, \mathbf{x}_N | \mathbf{z}) &= p(\mathbf{x}_1 | \mathbf{z})p(\mathbf{x}_2, \dots, \mathbf{x}_N | \mathbf{z}) \\ &= p(\mathbf{x}_1 | \mathbf{z})p(\mathbf{x}_2 | \mathbf{z})p(\mathbf{x}_3, \dots, \mathbf{x}_N | \mathbf{z}) \\ &= p(\mathbf{x}_1 | \mathbf{z})p(\mathbf{x}_2 | \mathbf{z}) \cdots p(\mathbf{x}_N | \mathbf{z}) \end{aligned}$$

Recall:



$$p(x, y) = p(x|y)p(y) \quad p(x, y) = p(x)p(y)$$



Graphical model visualization

To compute the MAP estimate

Given (1) a set of known parameters

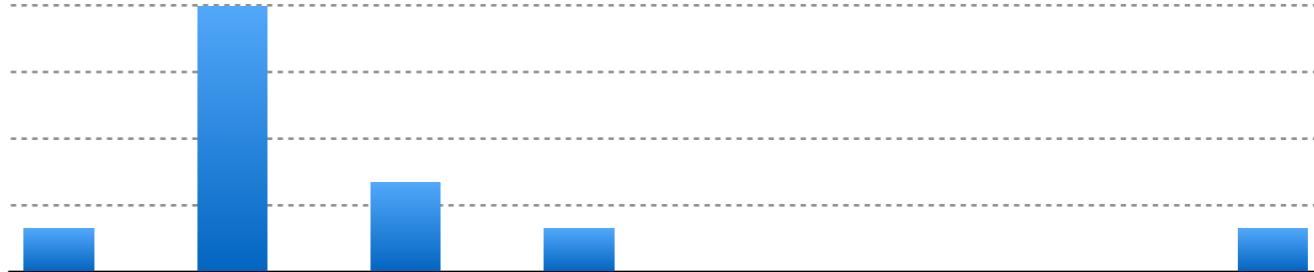
$$p(\mathbf{z}) \quad p(\mathbf{x}|\mathbf{z})$$

(2) observations

$$\{x_1, x_2, \dots, x_N\}$$

Compute which \mathbf{z} has the largest probability

$$\hat{\mathbf{z}} = \arg \max_{\mathbf{z} \in \mathcal{Z}} p(\mathbf{z}) \prod_n p(x_n | \mathbf{z})$$



count	1	6	2	1	0	0	0	1
word	Tartan	robot	CHIMP	CMU	bio	soft	ankle	sensor
p(x z)	0.09	0.55	0.18	0.09	0.0	0.0	0.0	0.09

$$\begin{aligned}
 p(X|z) &= \prod_v p(x_v|z)^{c(w_v)} \\
 &= (0.09)^1 (0.55)^6 \cdots (0.09)^1
 \end{aligned}$$

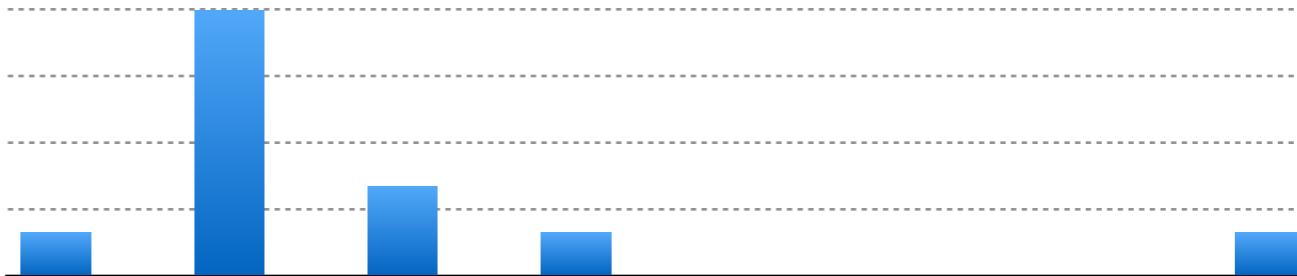
Numbers get really small so use log probabilities

$$\log p(X|z = \text{'grandchallenge'}) = -2.42 - 3.68 - 3.43 - 2.42 - 0.07 - 0.07 - 0.07 - 2.42 = -14.58$$

$$\log p(X|z = \text{'softrobot'}) = -7.63 - 9.37 - 15.18 - 2.97 - 0.02 - 0.01 - 0.02 - 2.27 = -37.48$$

* typically add pseudo-counts (0.001)

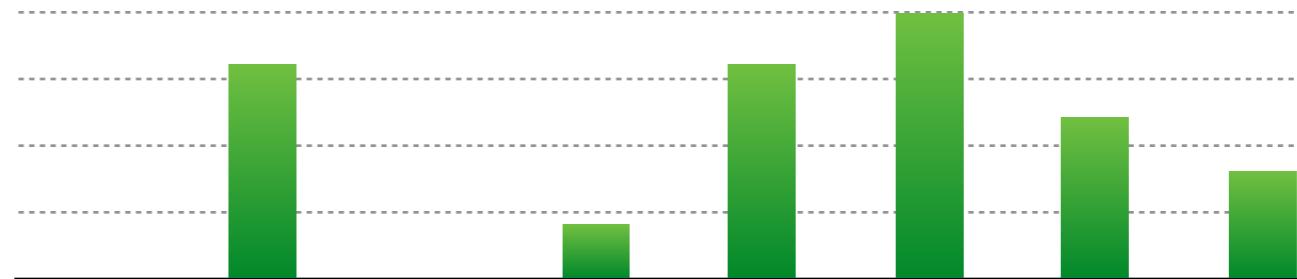
** this is an example for computing the likelihood, need to multiply times **prior** to get posterior



count	1	6	2	1	0	0	0	1
word	Tartan	robot	CHIMP	CMU	bio	soft	ankle	sensor
p(x z)	0.09	0.55	0.18	0.09	0.0	0.0	0.0	0.09

$$\log p(X|z=\text{grand challenge}) = -14.58$$

$$\log p(X|z=\text{bio inspired}) = -37.48$$



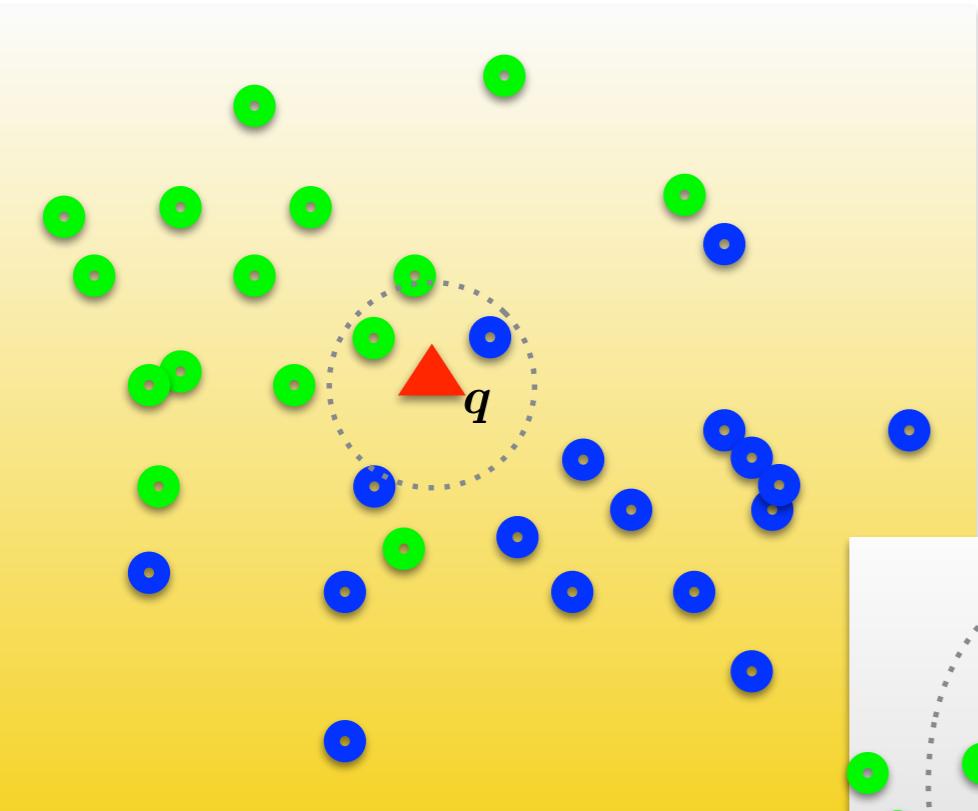
count	0	4	0	1	4	5	3	2
word	Tartan	robot	CHIMP	CMU	bio	soft	ankle	sensor
p(x z)	0.0	0.21	0.0	0.05	0.21	0.26	0.16	0.11

$$\log p(X|z=\text{grand challenge}) = -94.06$$

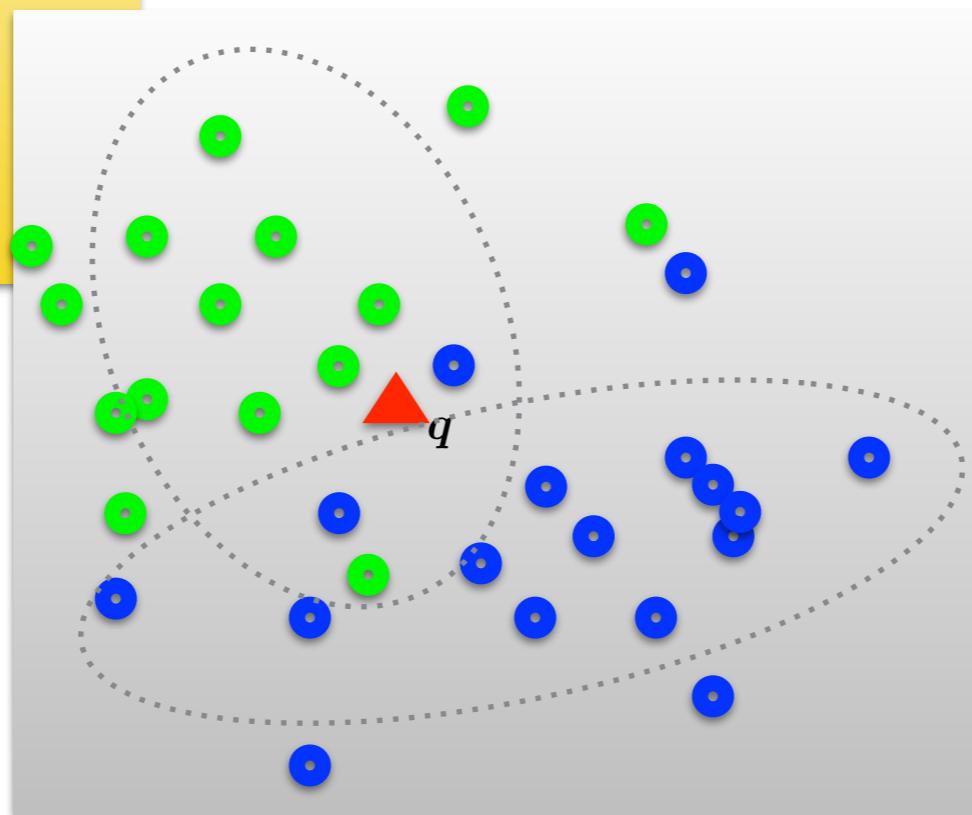
$$\log p(X|z=\text{bio inspired}) = -32.41$$

* typically add pseudo-counts (0.001)

** this is an example for computing the likelihood, need to multiply times prior to get posterior

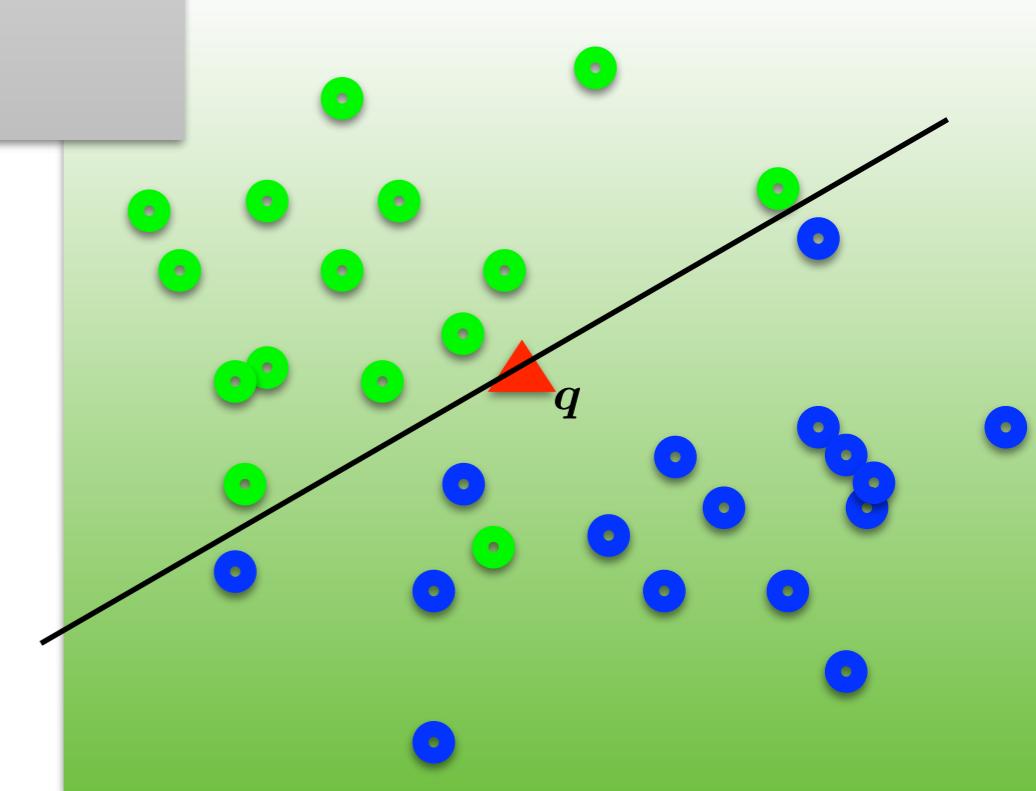


Nearest Neighbor

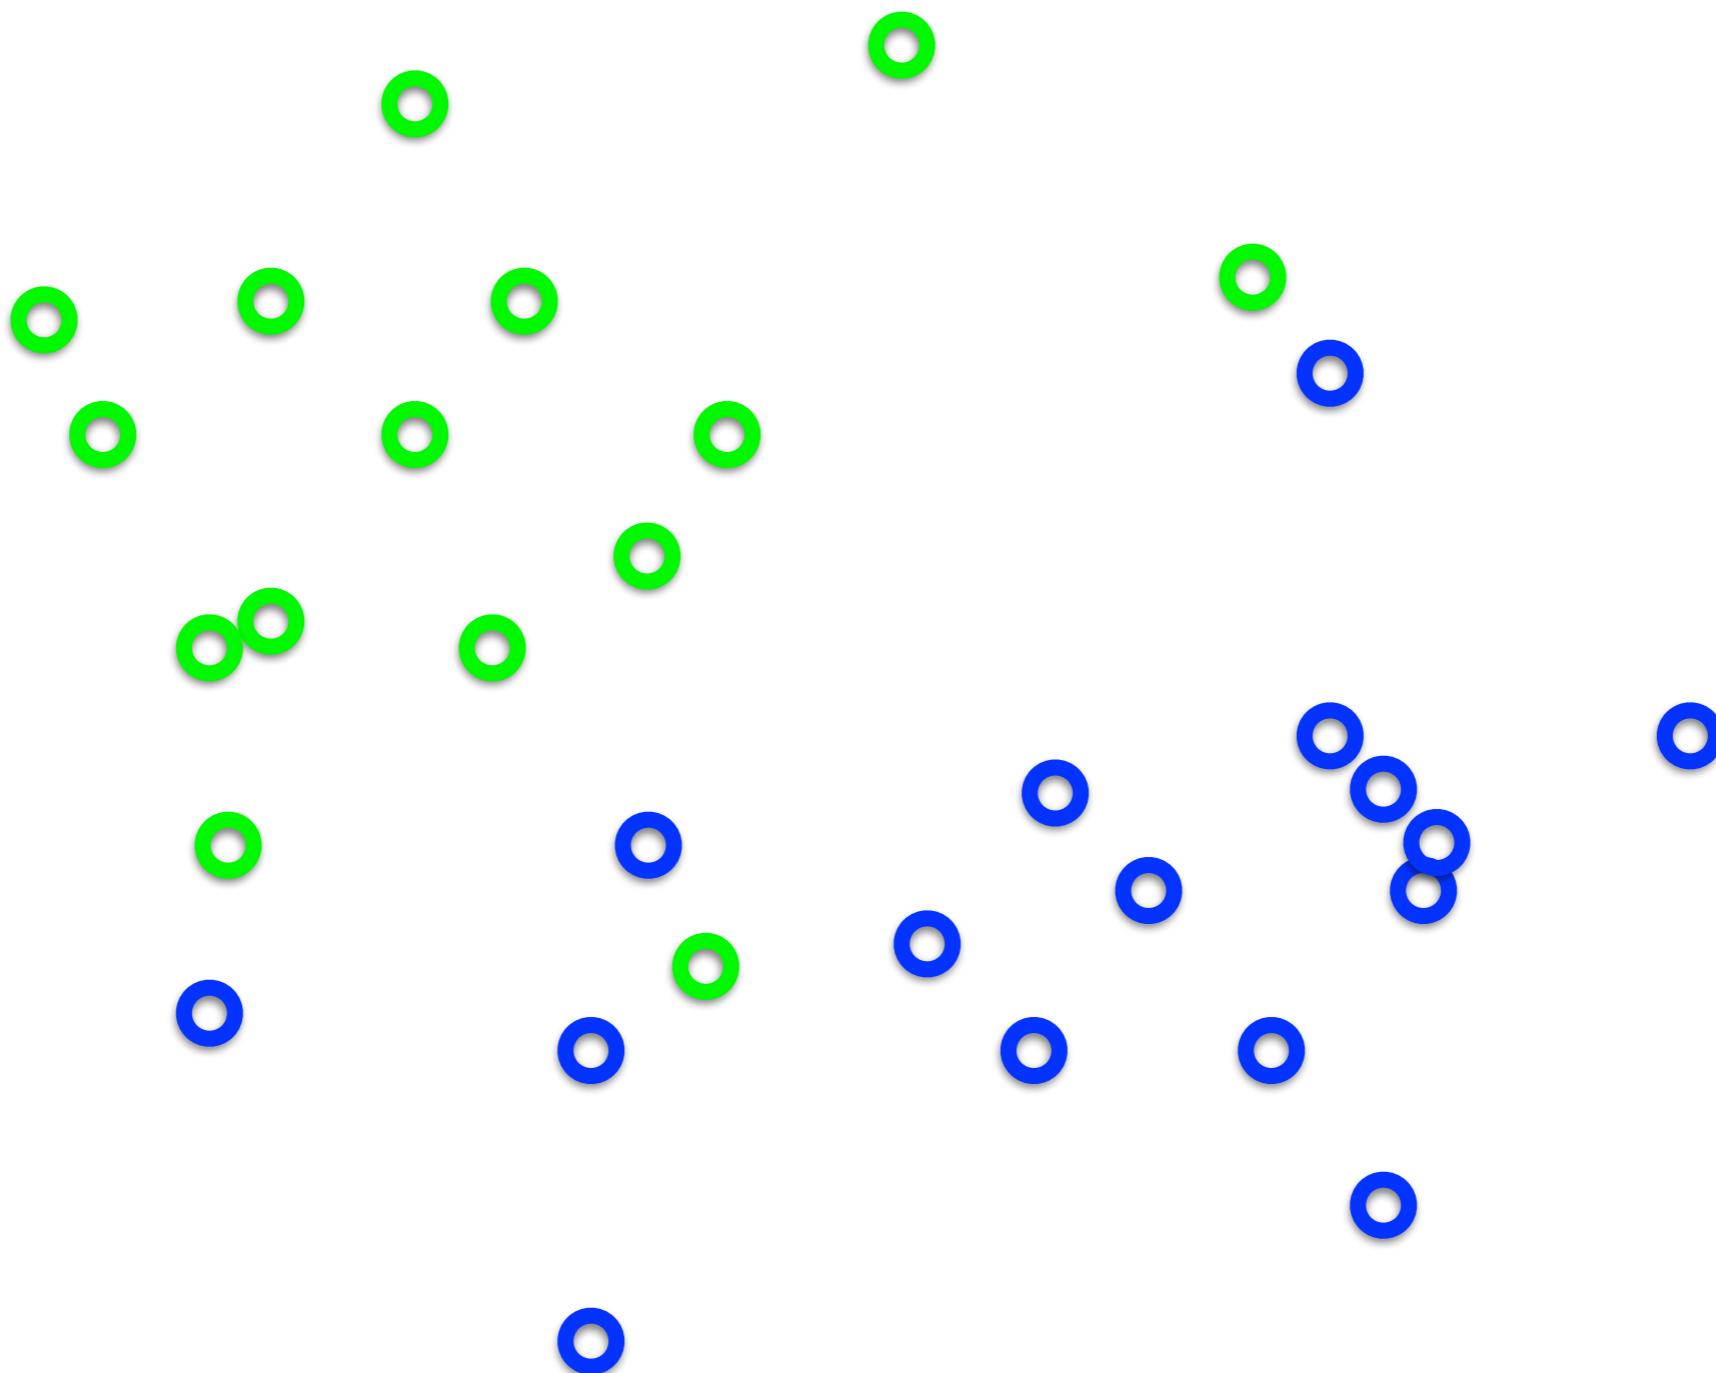


Naive Bayes

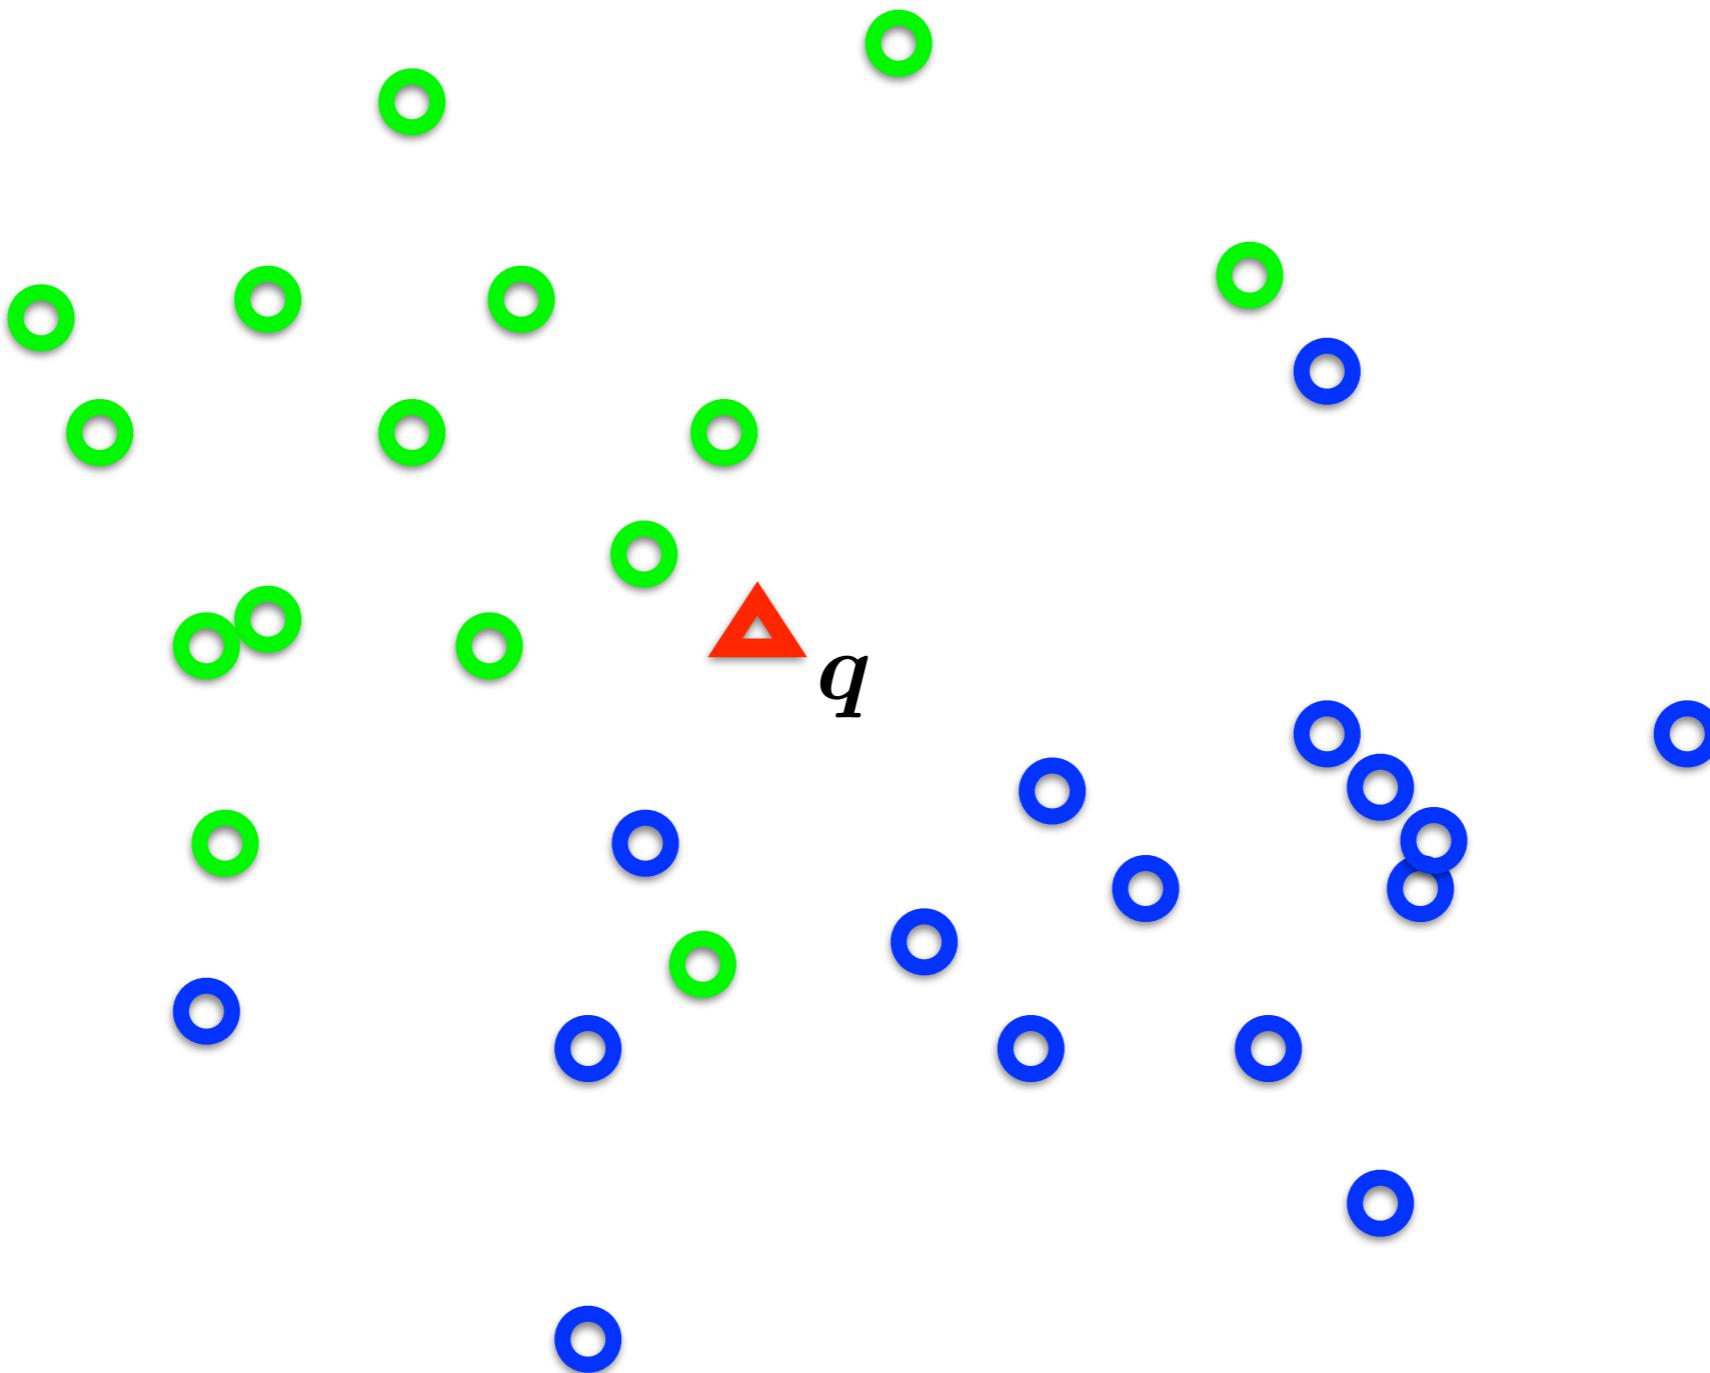
Support Vector Machine



Distribution of data from two classes

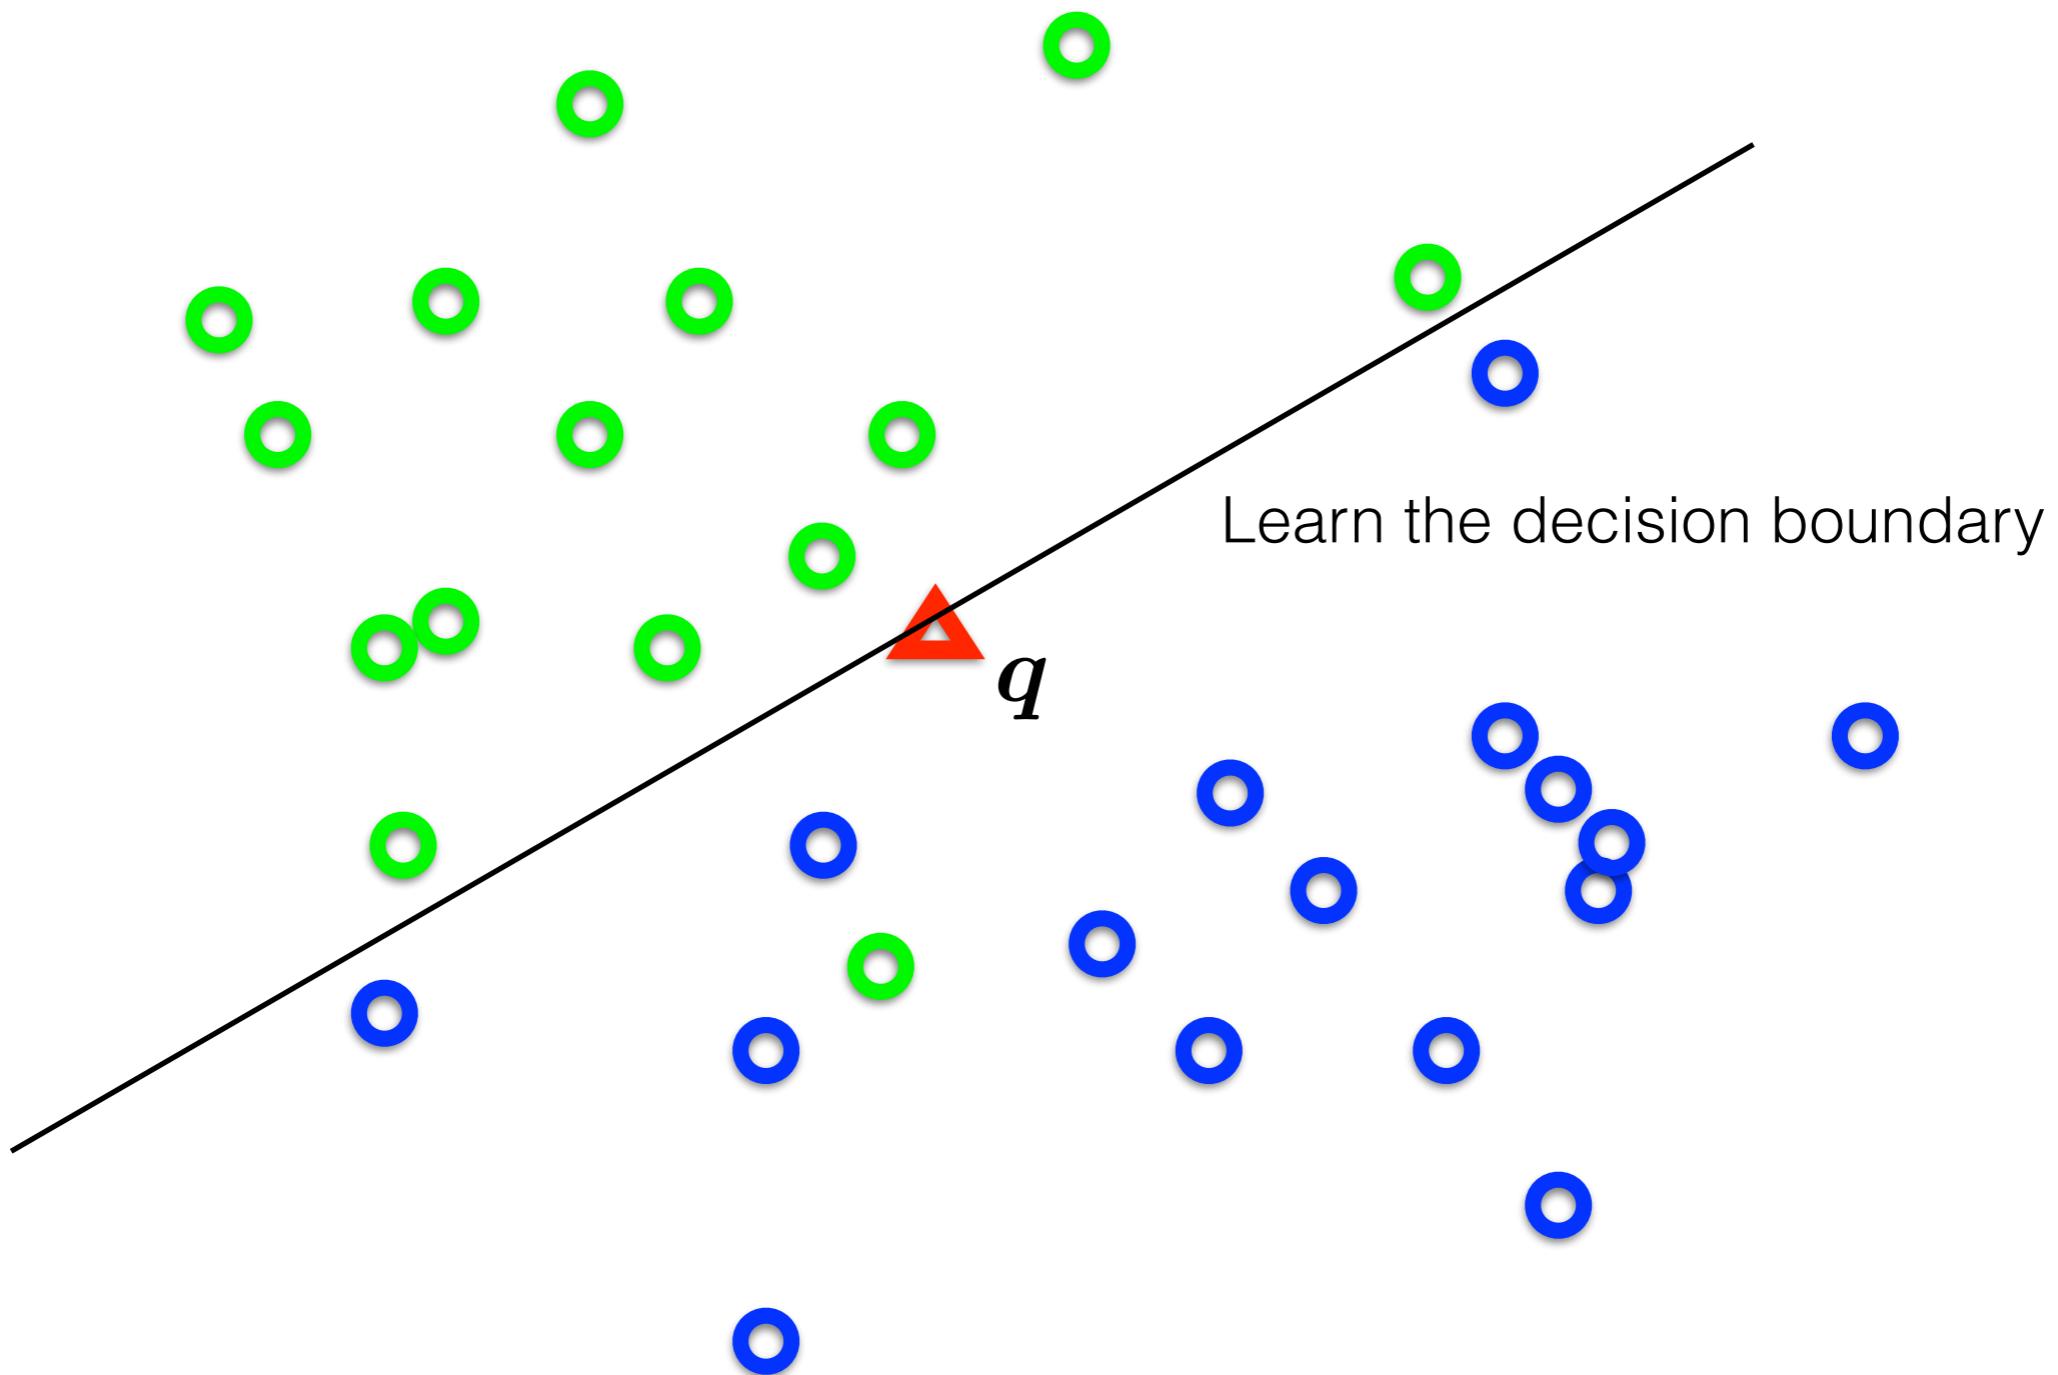


Distribution of data from two classes



Which class does q belong to?

Distribution of data from two classes

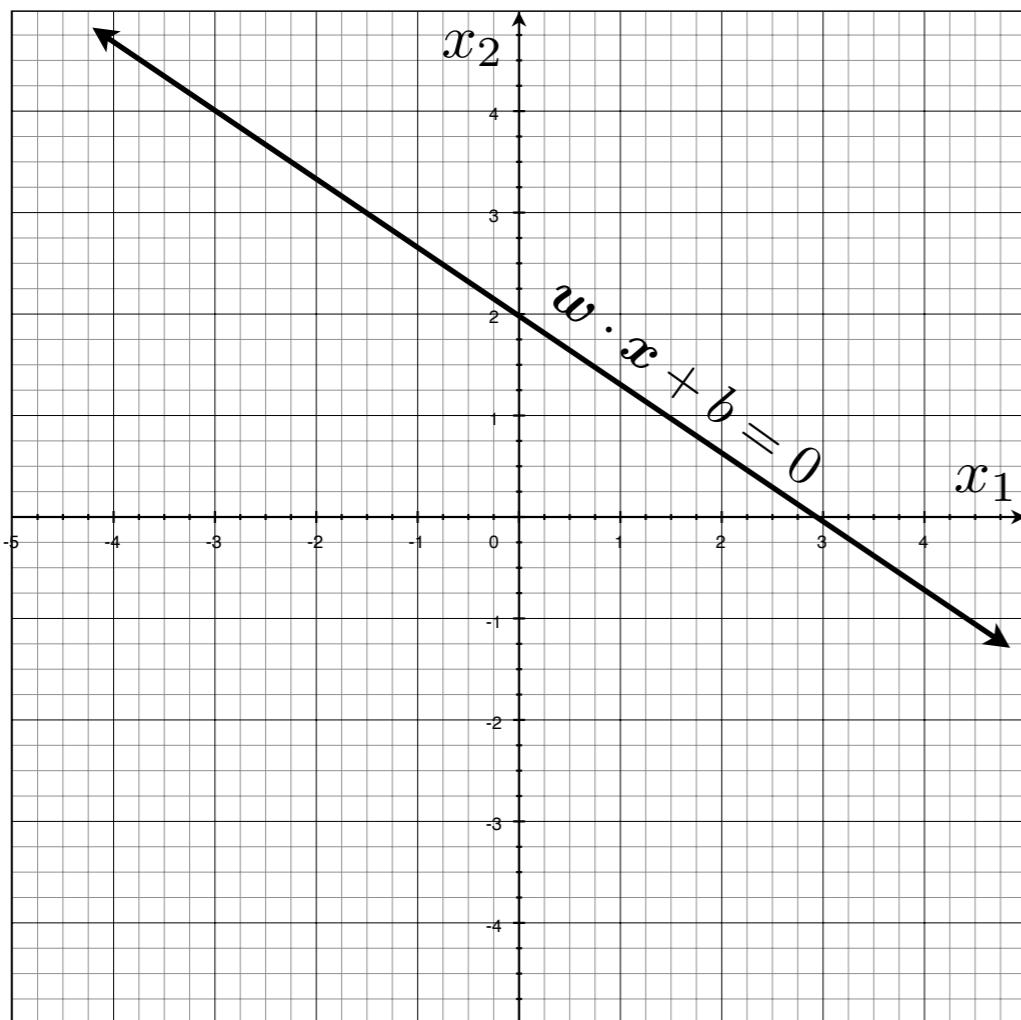


Support Vector Machine

First we need to understand hyperplanes...

Hyperplanes (lines) in 2D

$$w_1x_1 + w_2x_2 + b = 0$$



a line can be written as
dot product plus a bias

$$\mathbf{w} \cdot \mathbf{x} + b = 0$$

$$\mathbf{w} \in \mathcal{R}^2$$

another version, add a weight 1
and push the bias inside

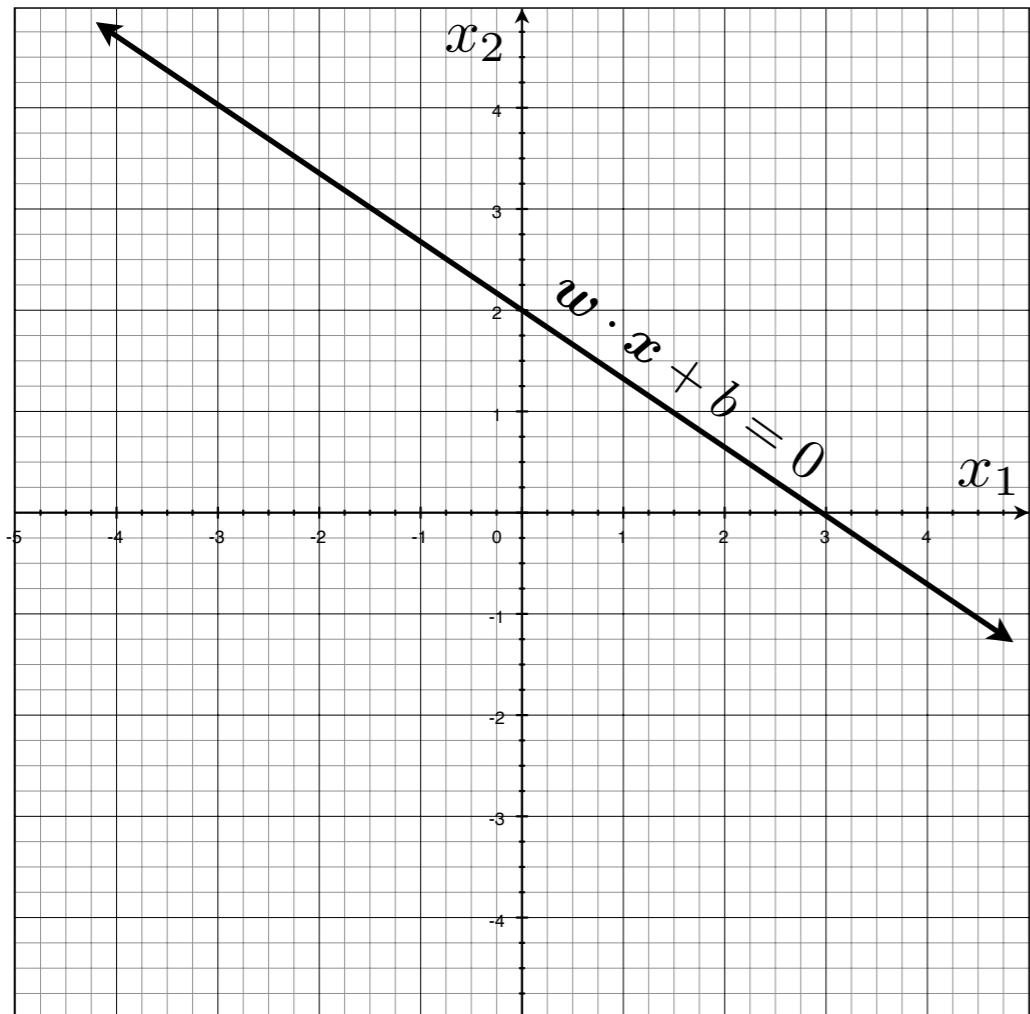
$$\mathbf{w} \cdot \mathbf{x} = 0$$

$$\mathbf{w} \in \mathcal{R}^3$$

Hyperplanes (lines) in 2D

$$\mathbf{w} \cdot \mathbf{x} + b = 0 \quad (\text{offset/bias outside}) \quad \mathbf{w} \cdot \mathbf{x} = 0 \quad (\text{offset/bias inside})$$

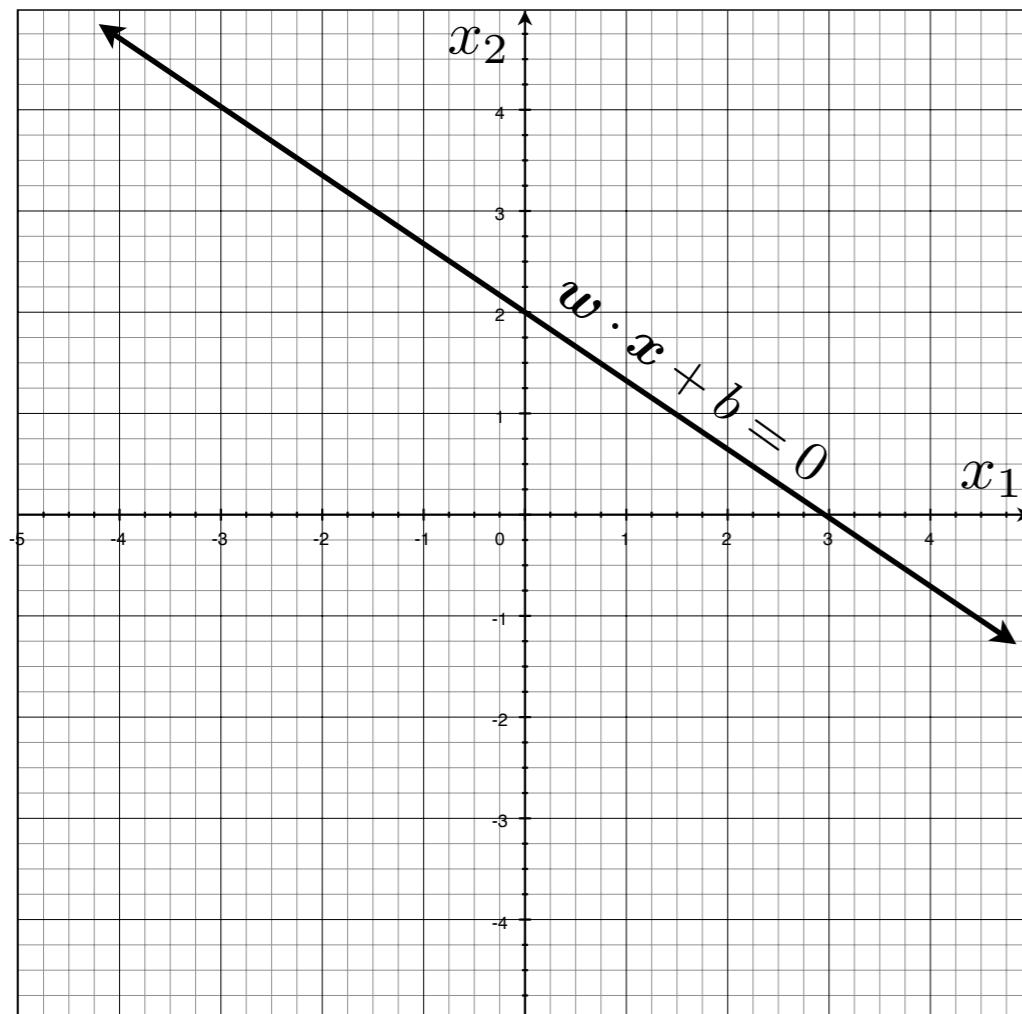
$$w_1x_1 + w_2x_2 + b = 0$$



Hyperplanes (lines) in 2D

$$\mathbf{w} \cdot \mathbf{x} + b = 0 \quad (\text{offset/bias outside}) \quad \mathbf{w} \cdot \mathbf{x} = 0 \quad (\text{offset/bias inside})$$

$$w_1x_1 + w_2x_2 + b = 0$$



Important property:
Free to choose any normalization of w

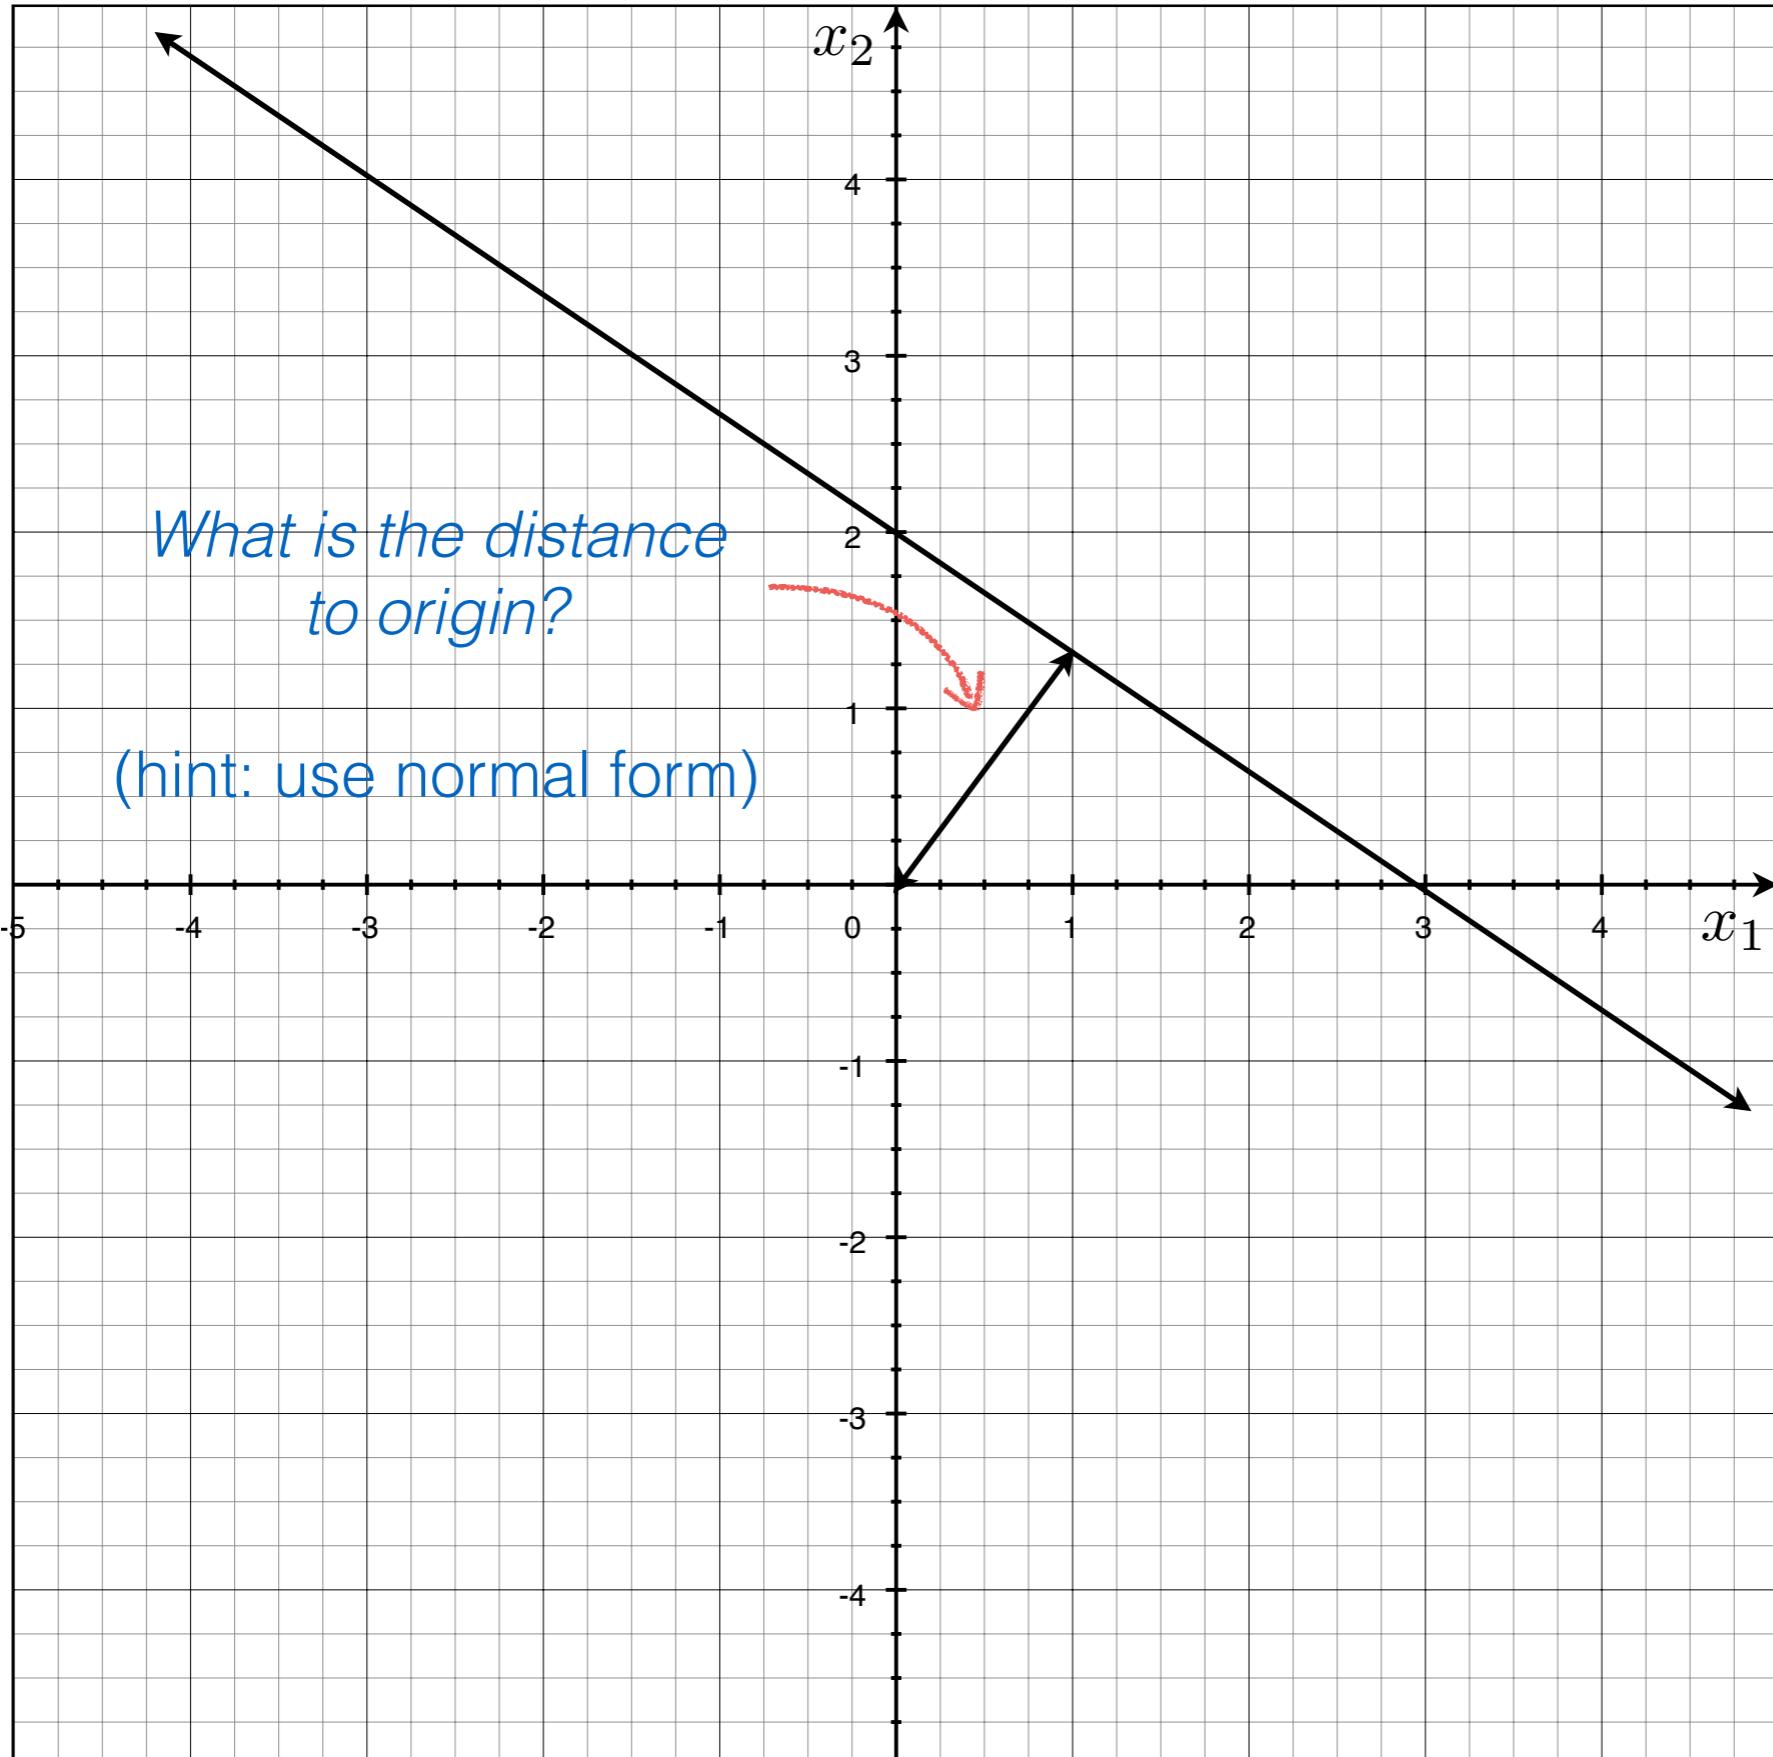
The line

$$w_1x_1 + w_2x_2 + b = 0$$

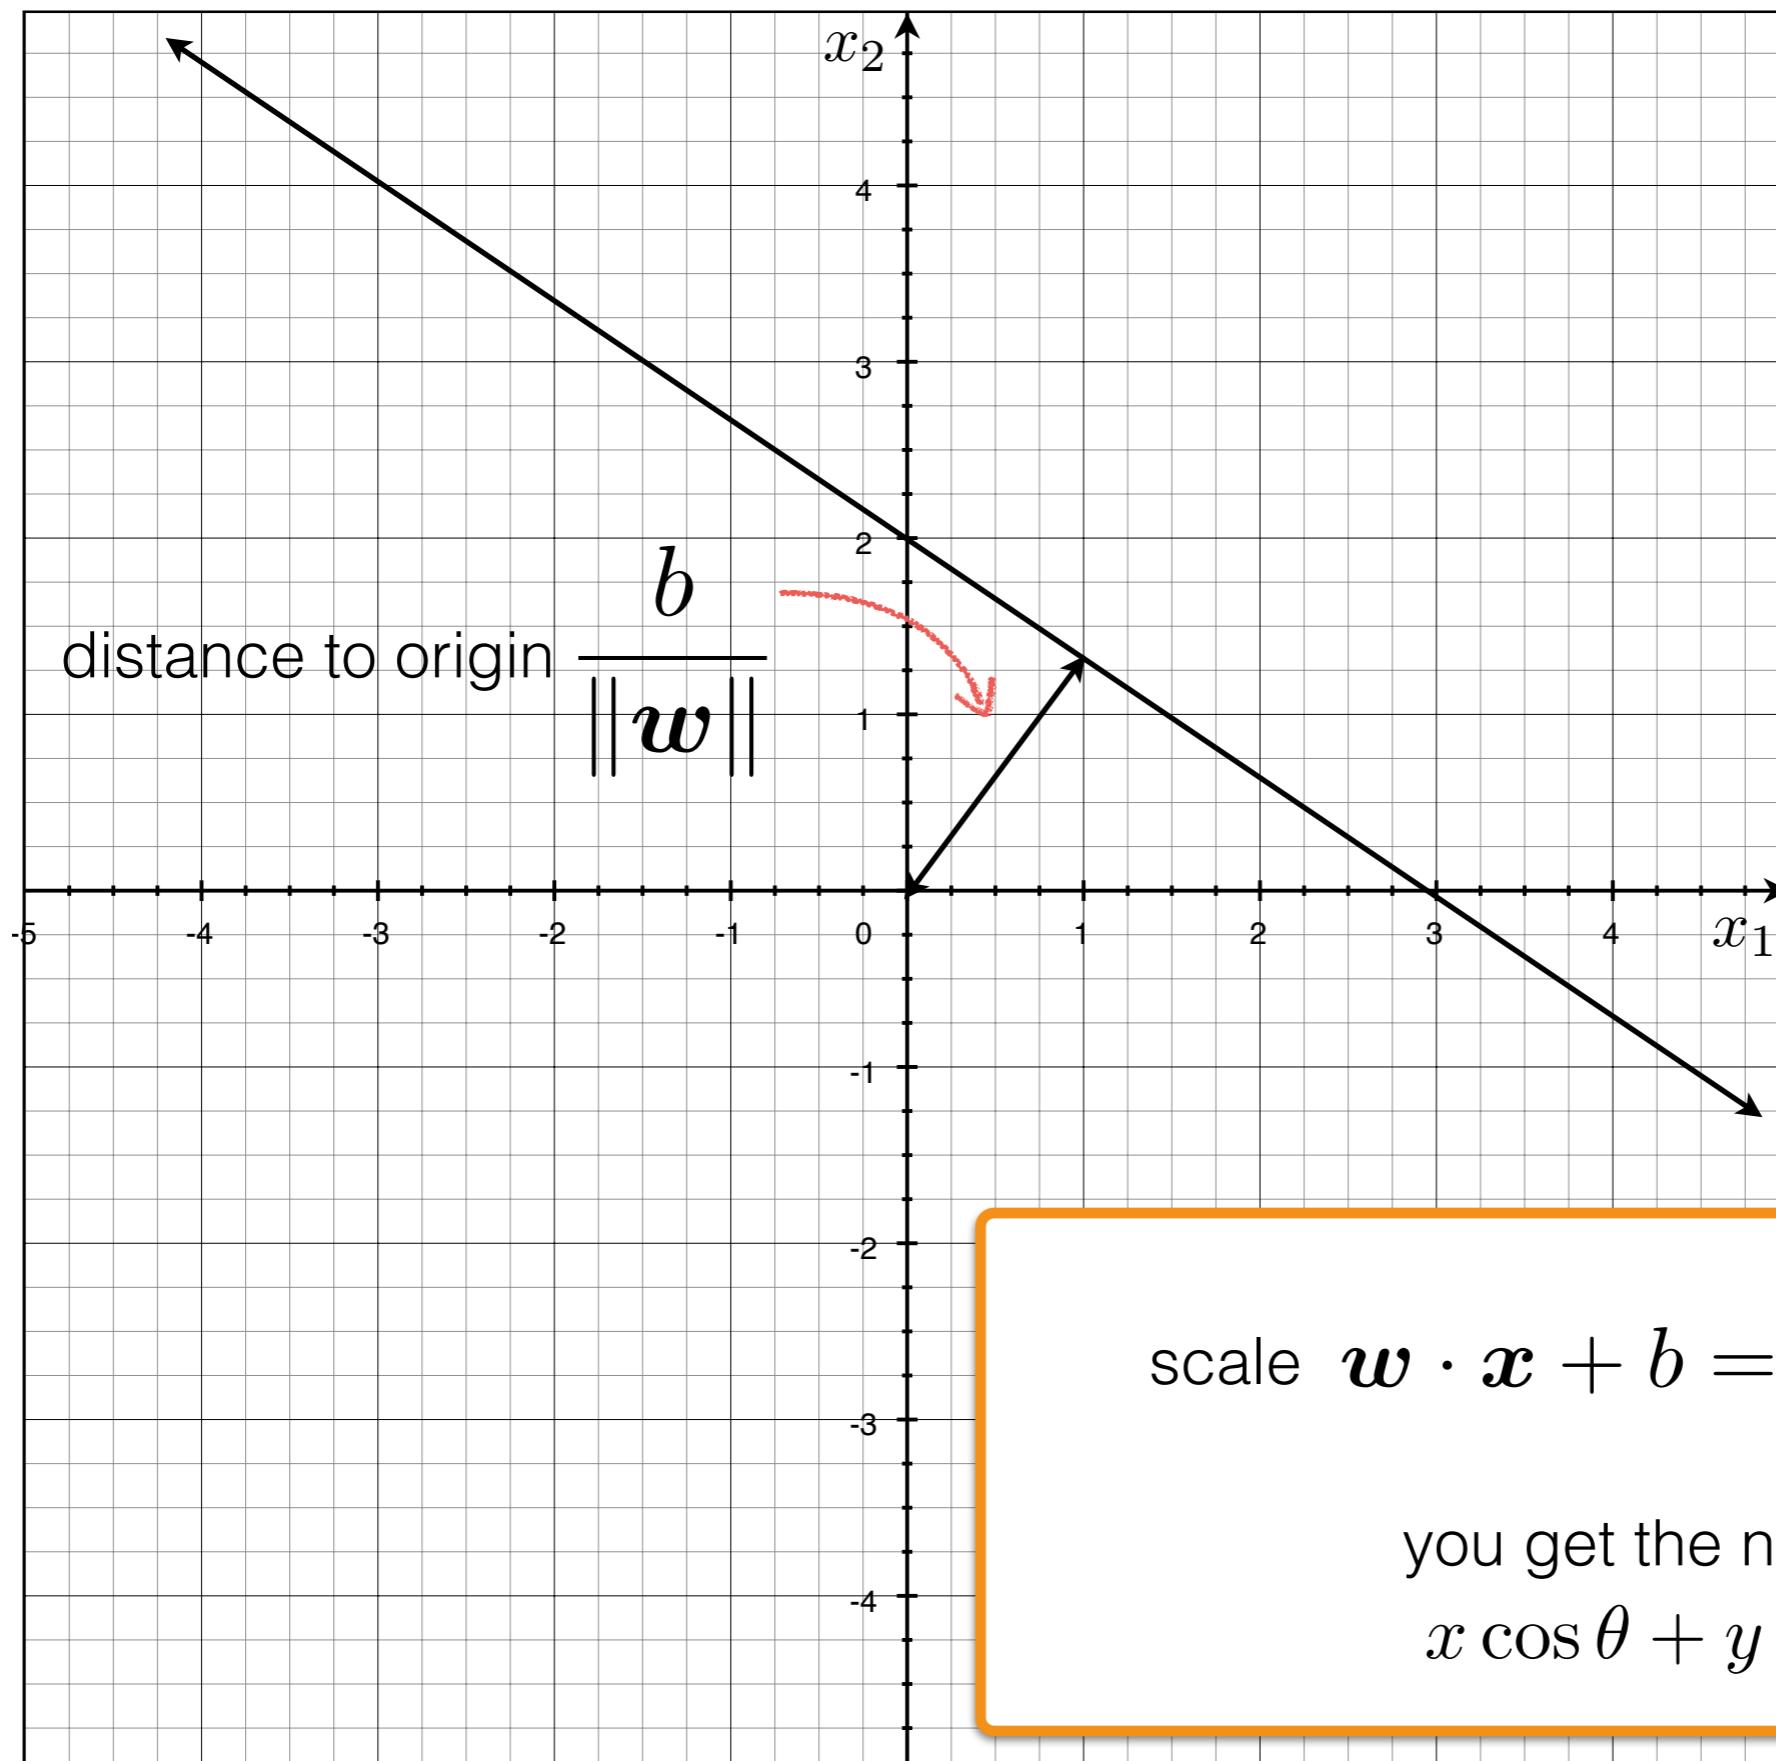
and the line

$$\lambda(w_1x_1 + w_2x_2 + b) = 0$$

define the same line

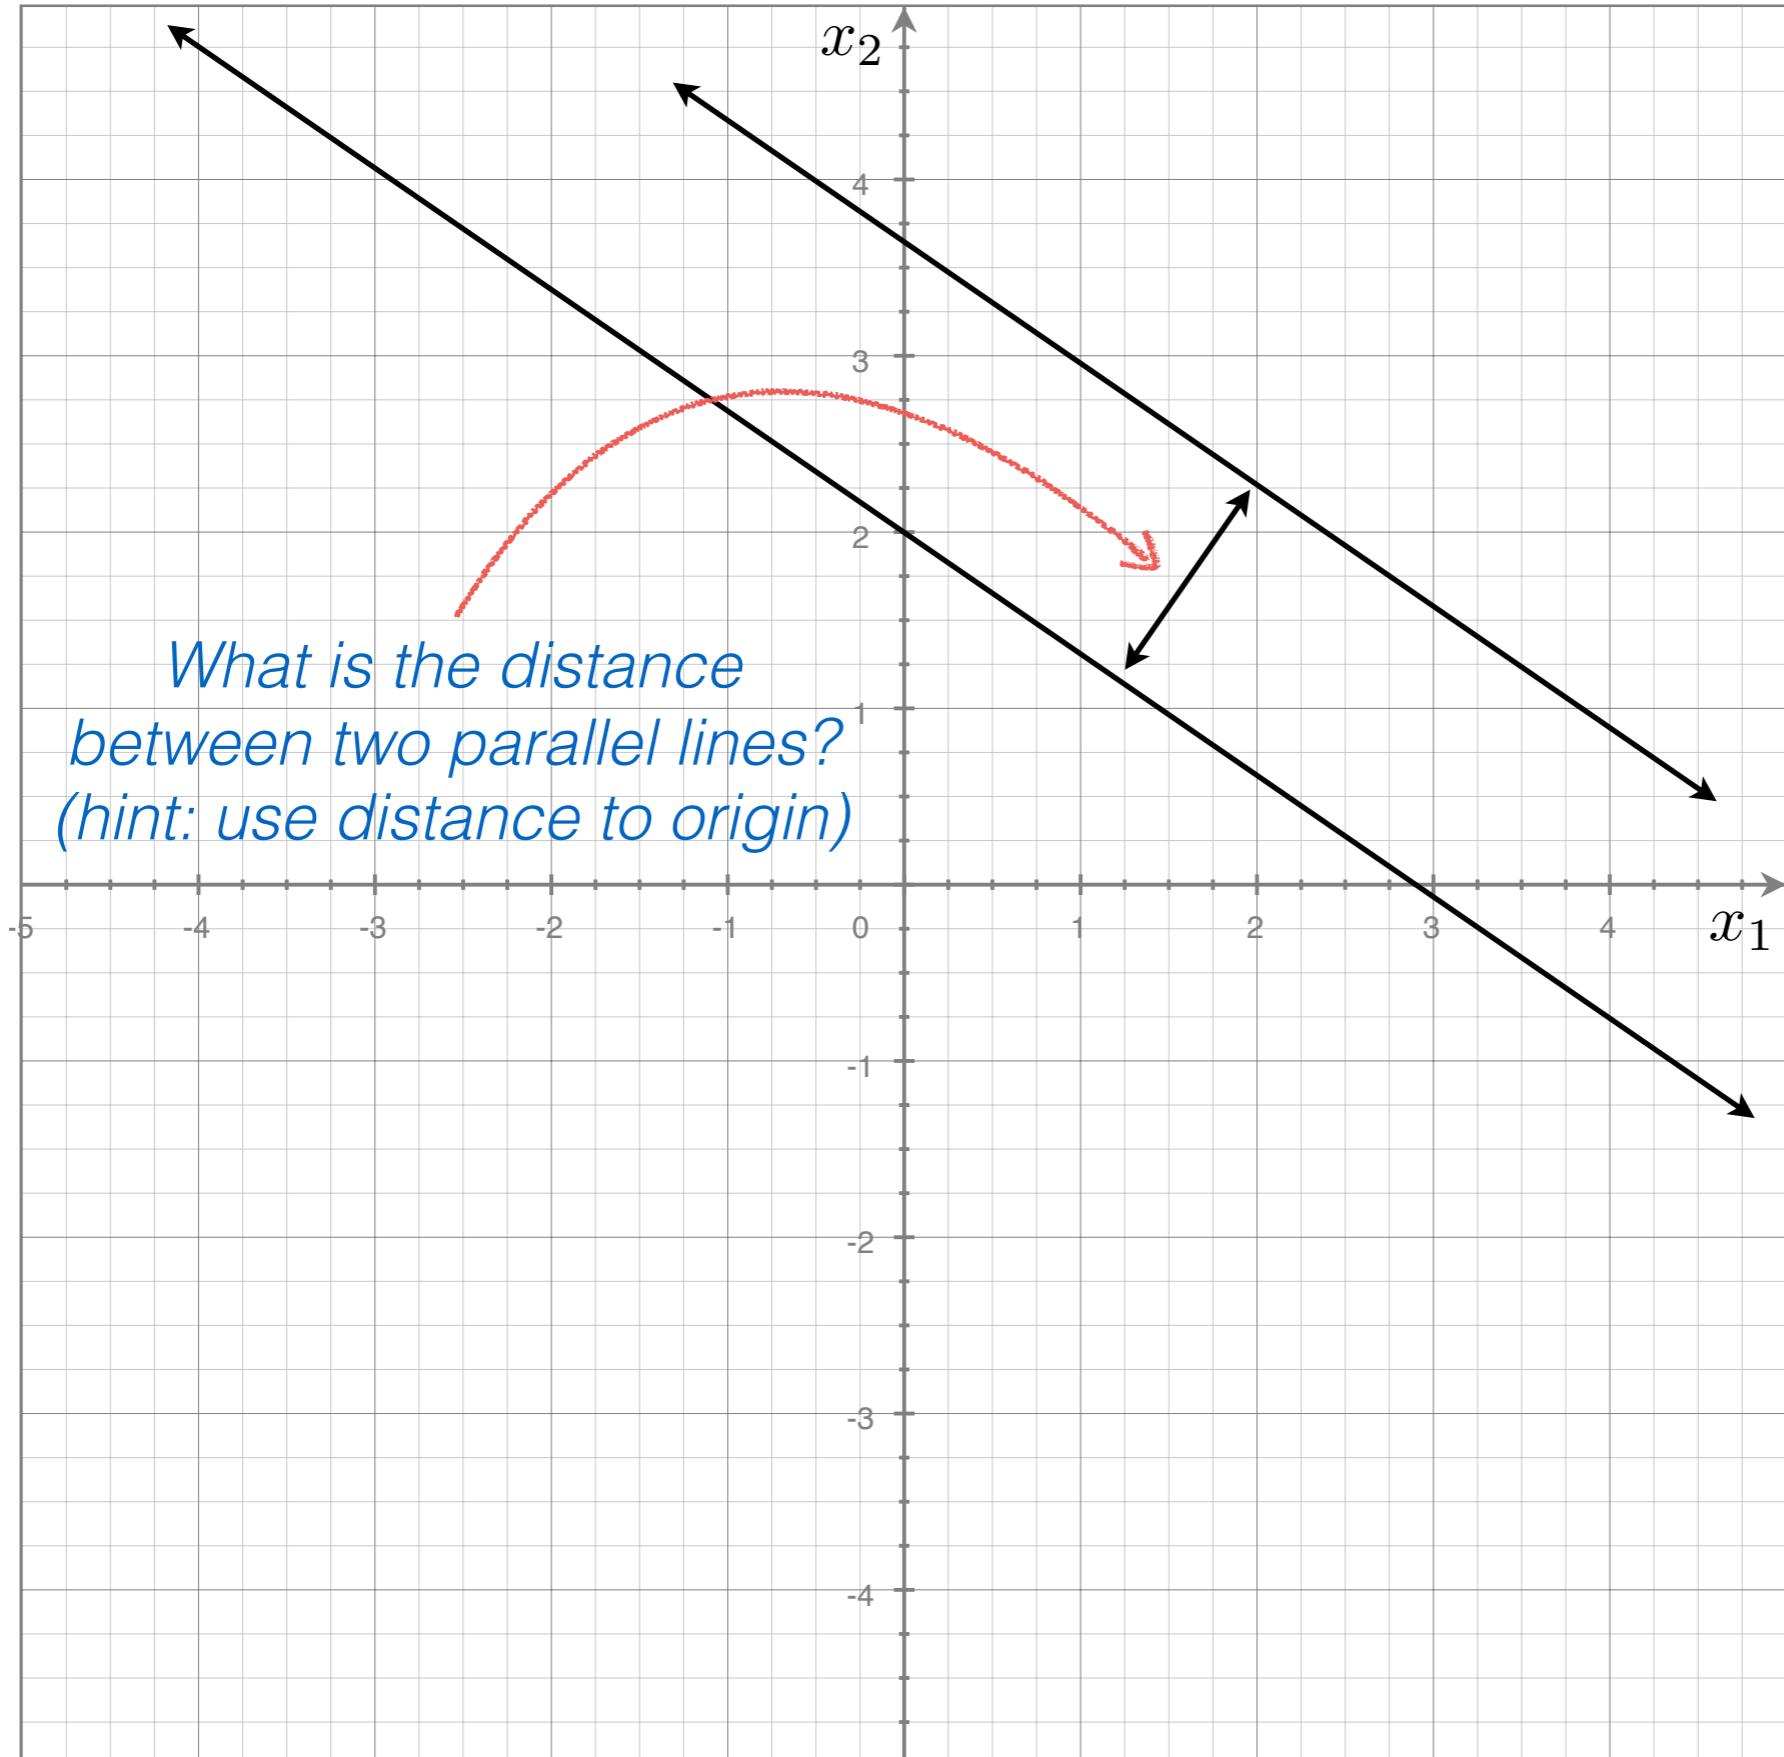


$$\mathbf{w} \cdot \mathbf{x} + b = 0$$



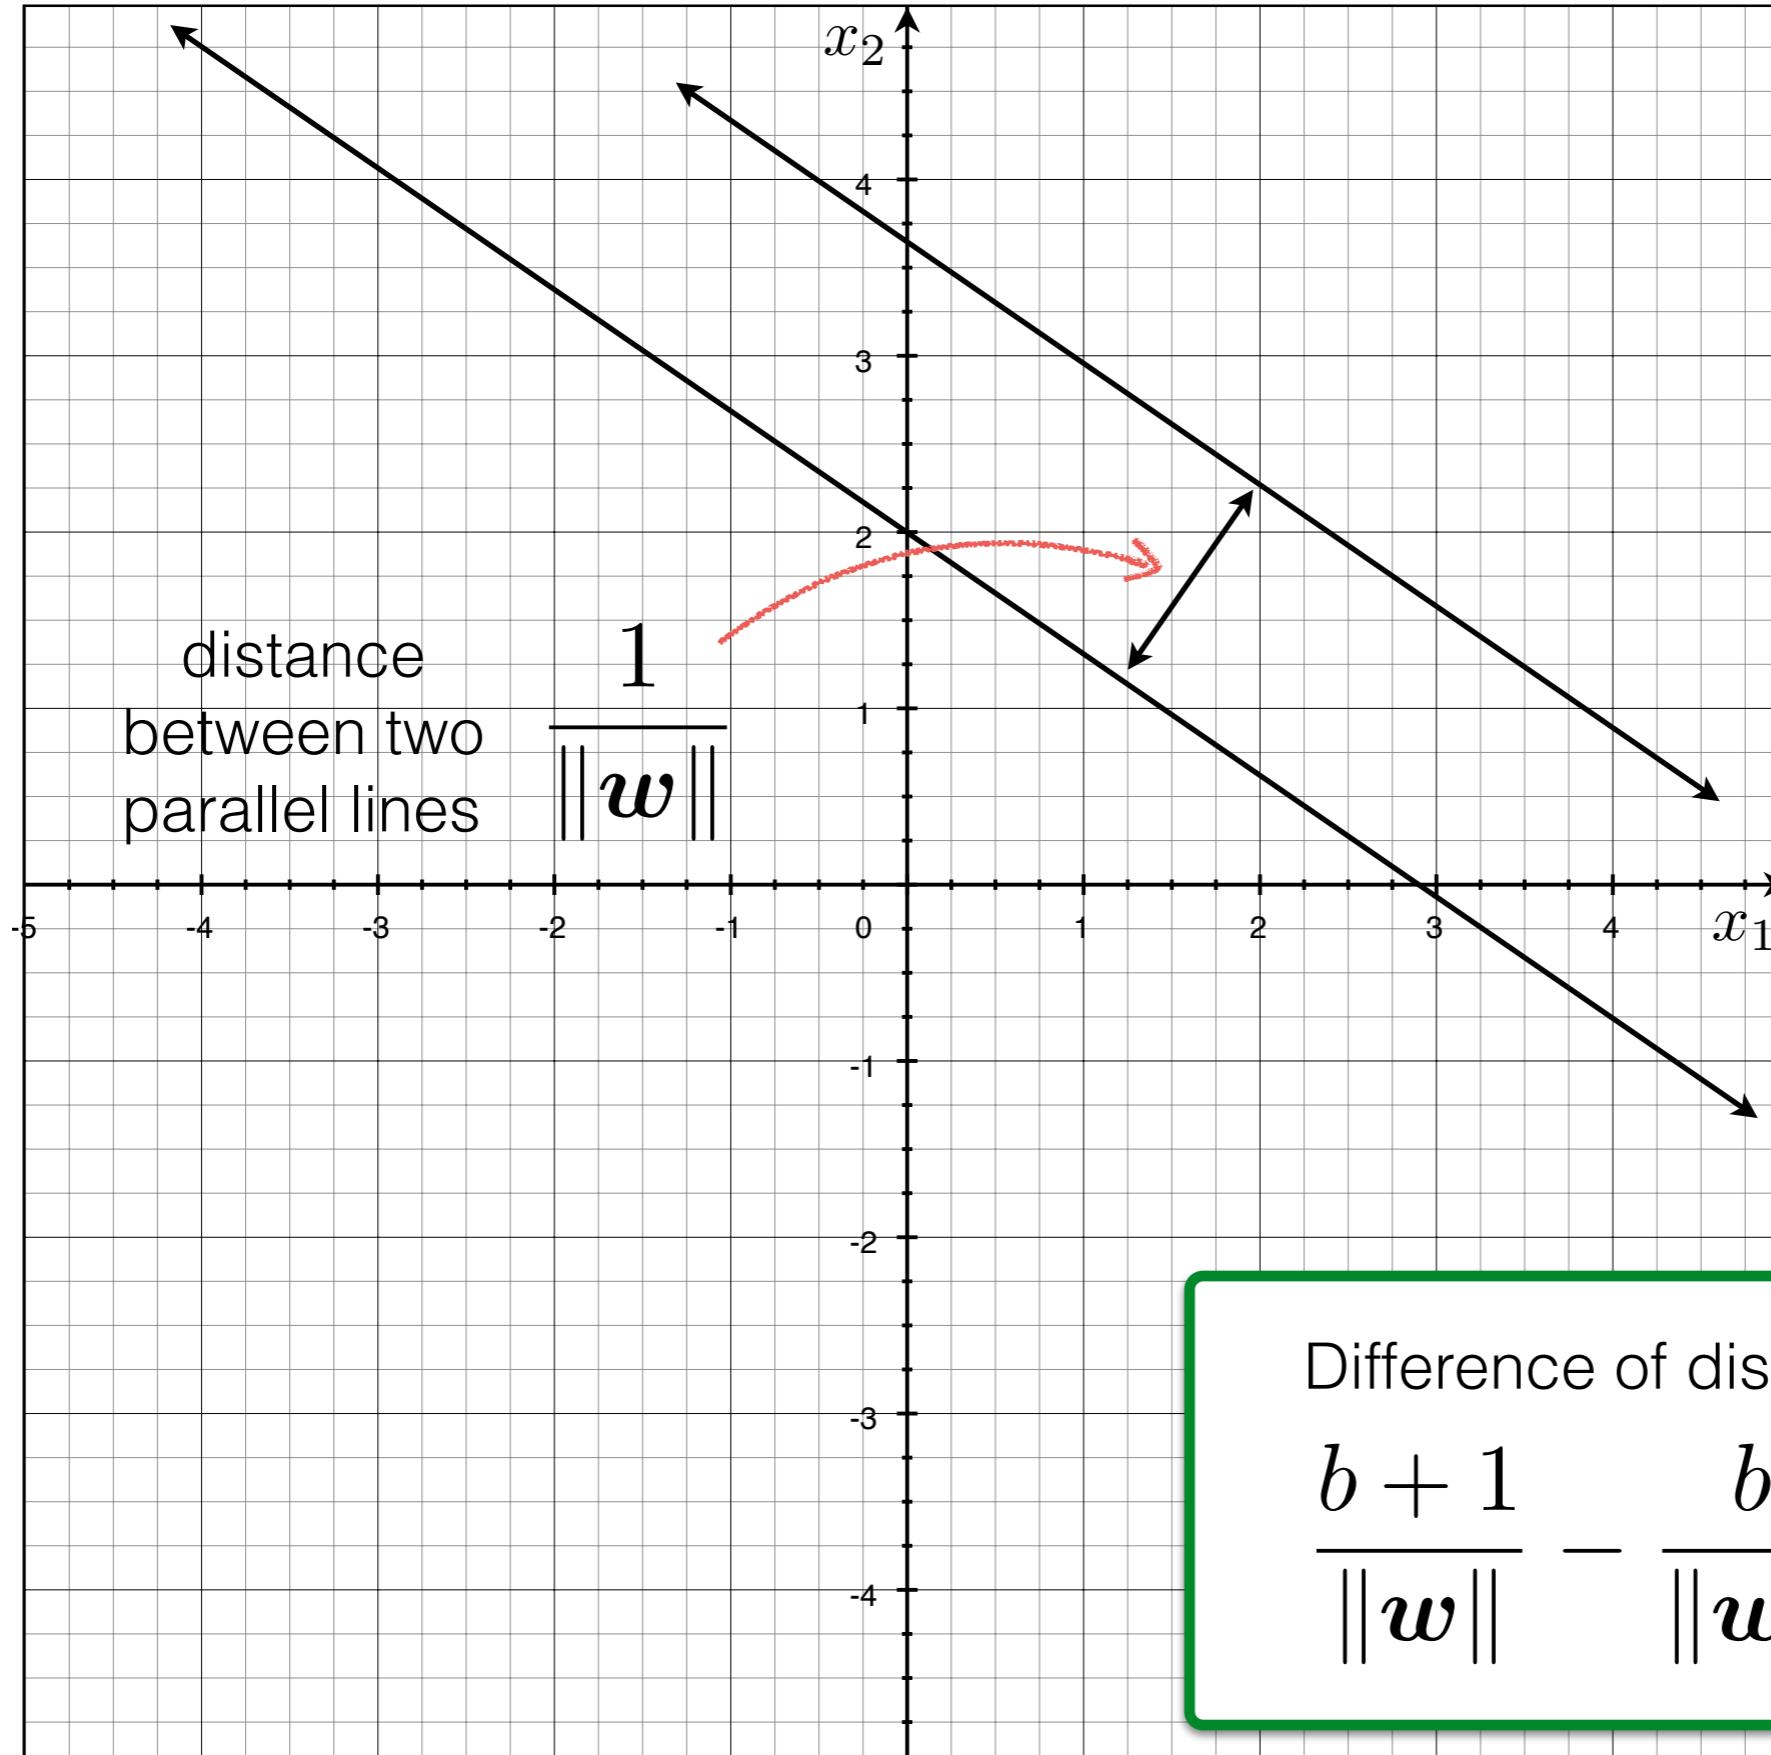
scale $w \cdot x + b = 0$ by $\frac{1}{\|w\|}$

you get the normal form
 $x \cos \theta + y \sin \theta = \rho$



$$\mathbf{w} \cdot \mathbf{x} + b = -1$$

$$\mathbf{w} \cdot \mathbf{x} + b = 0$$

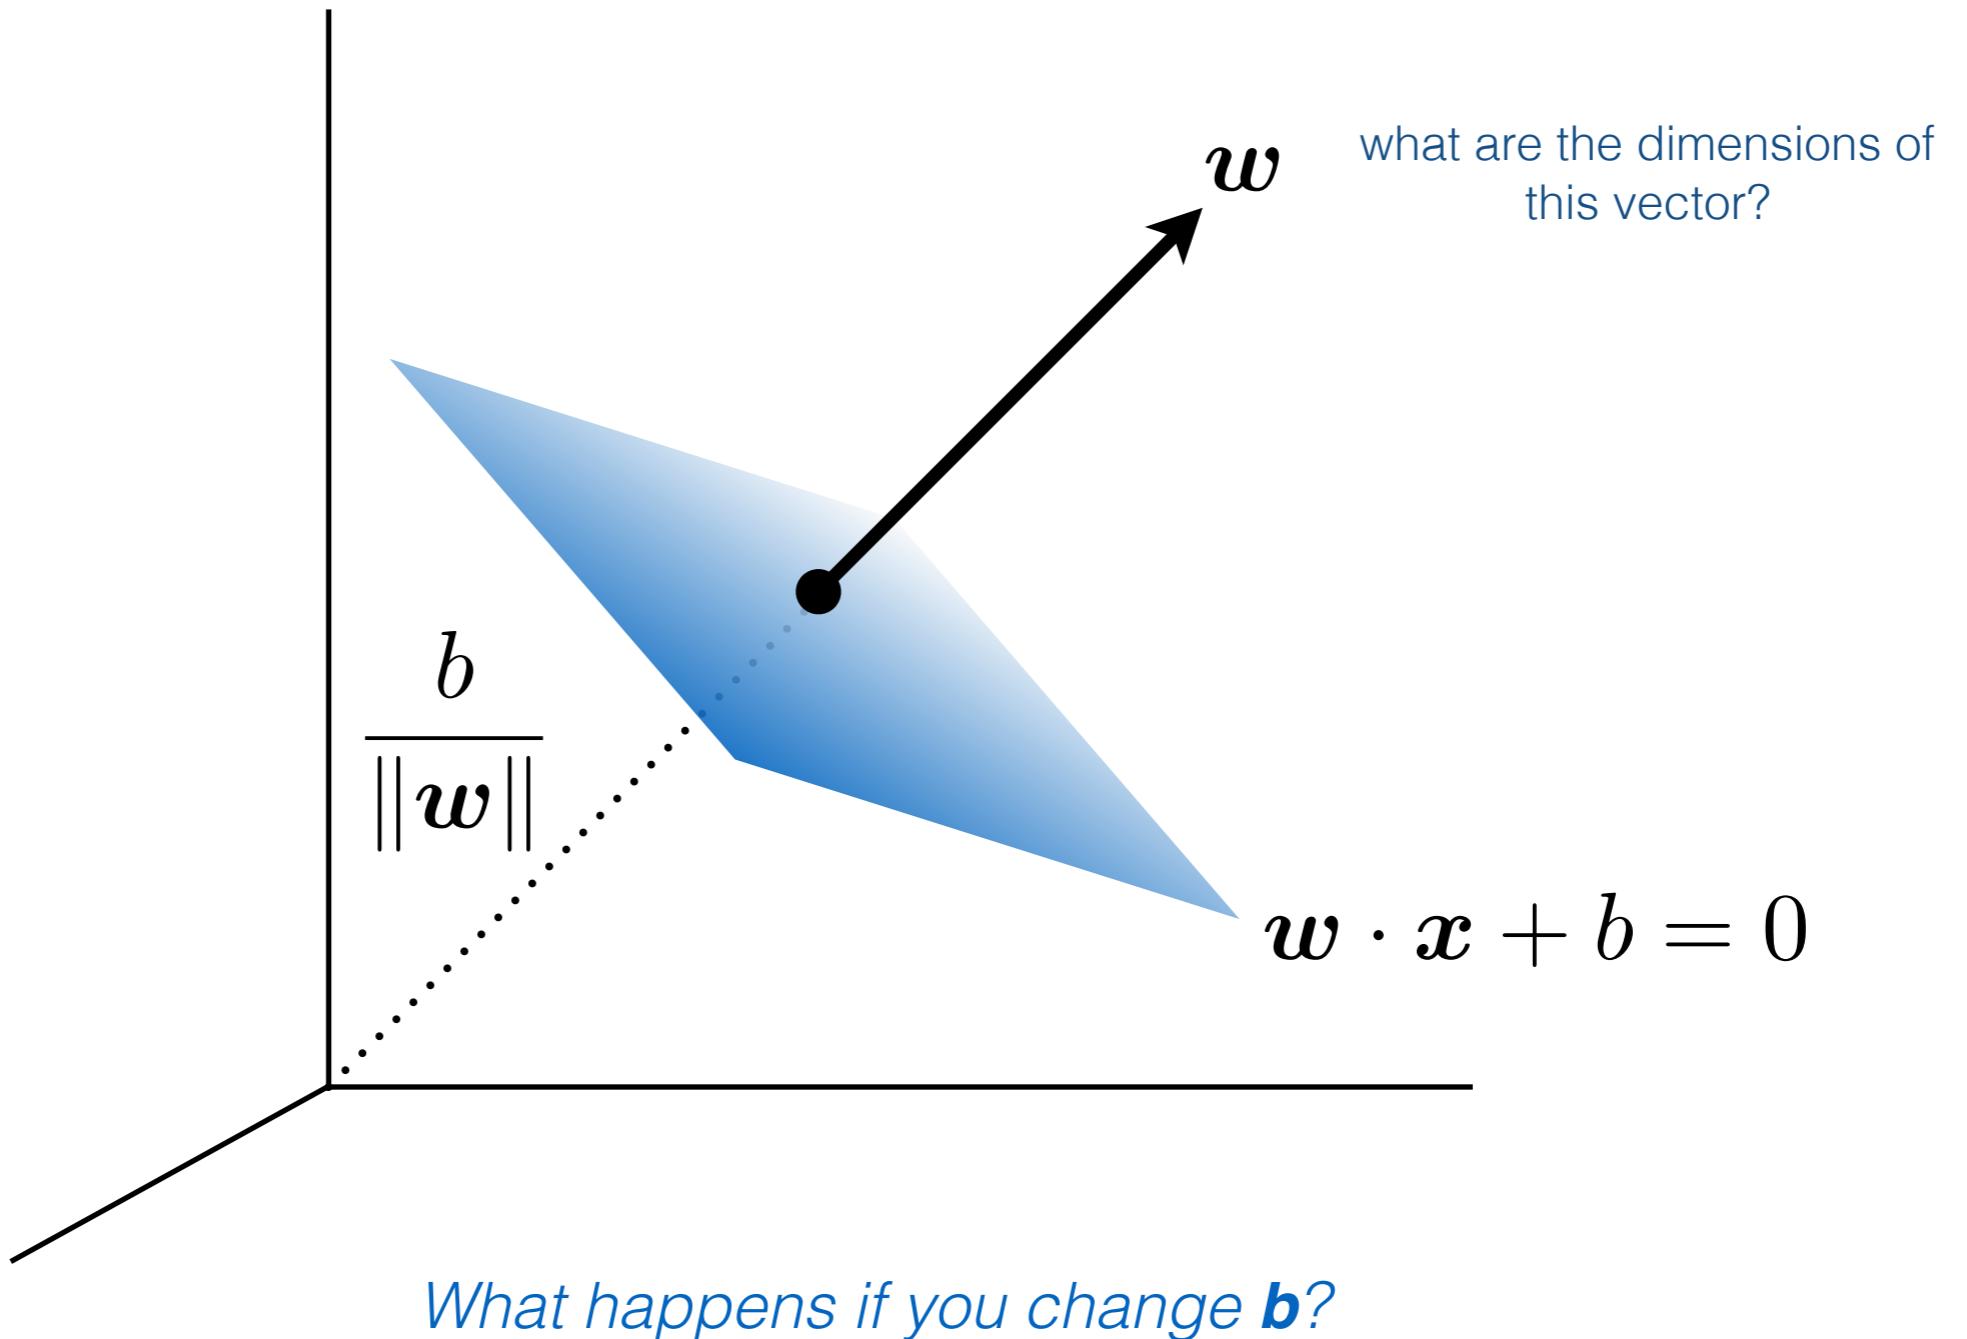


Difference of distance to origin

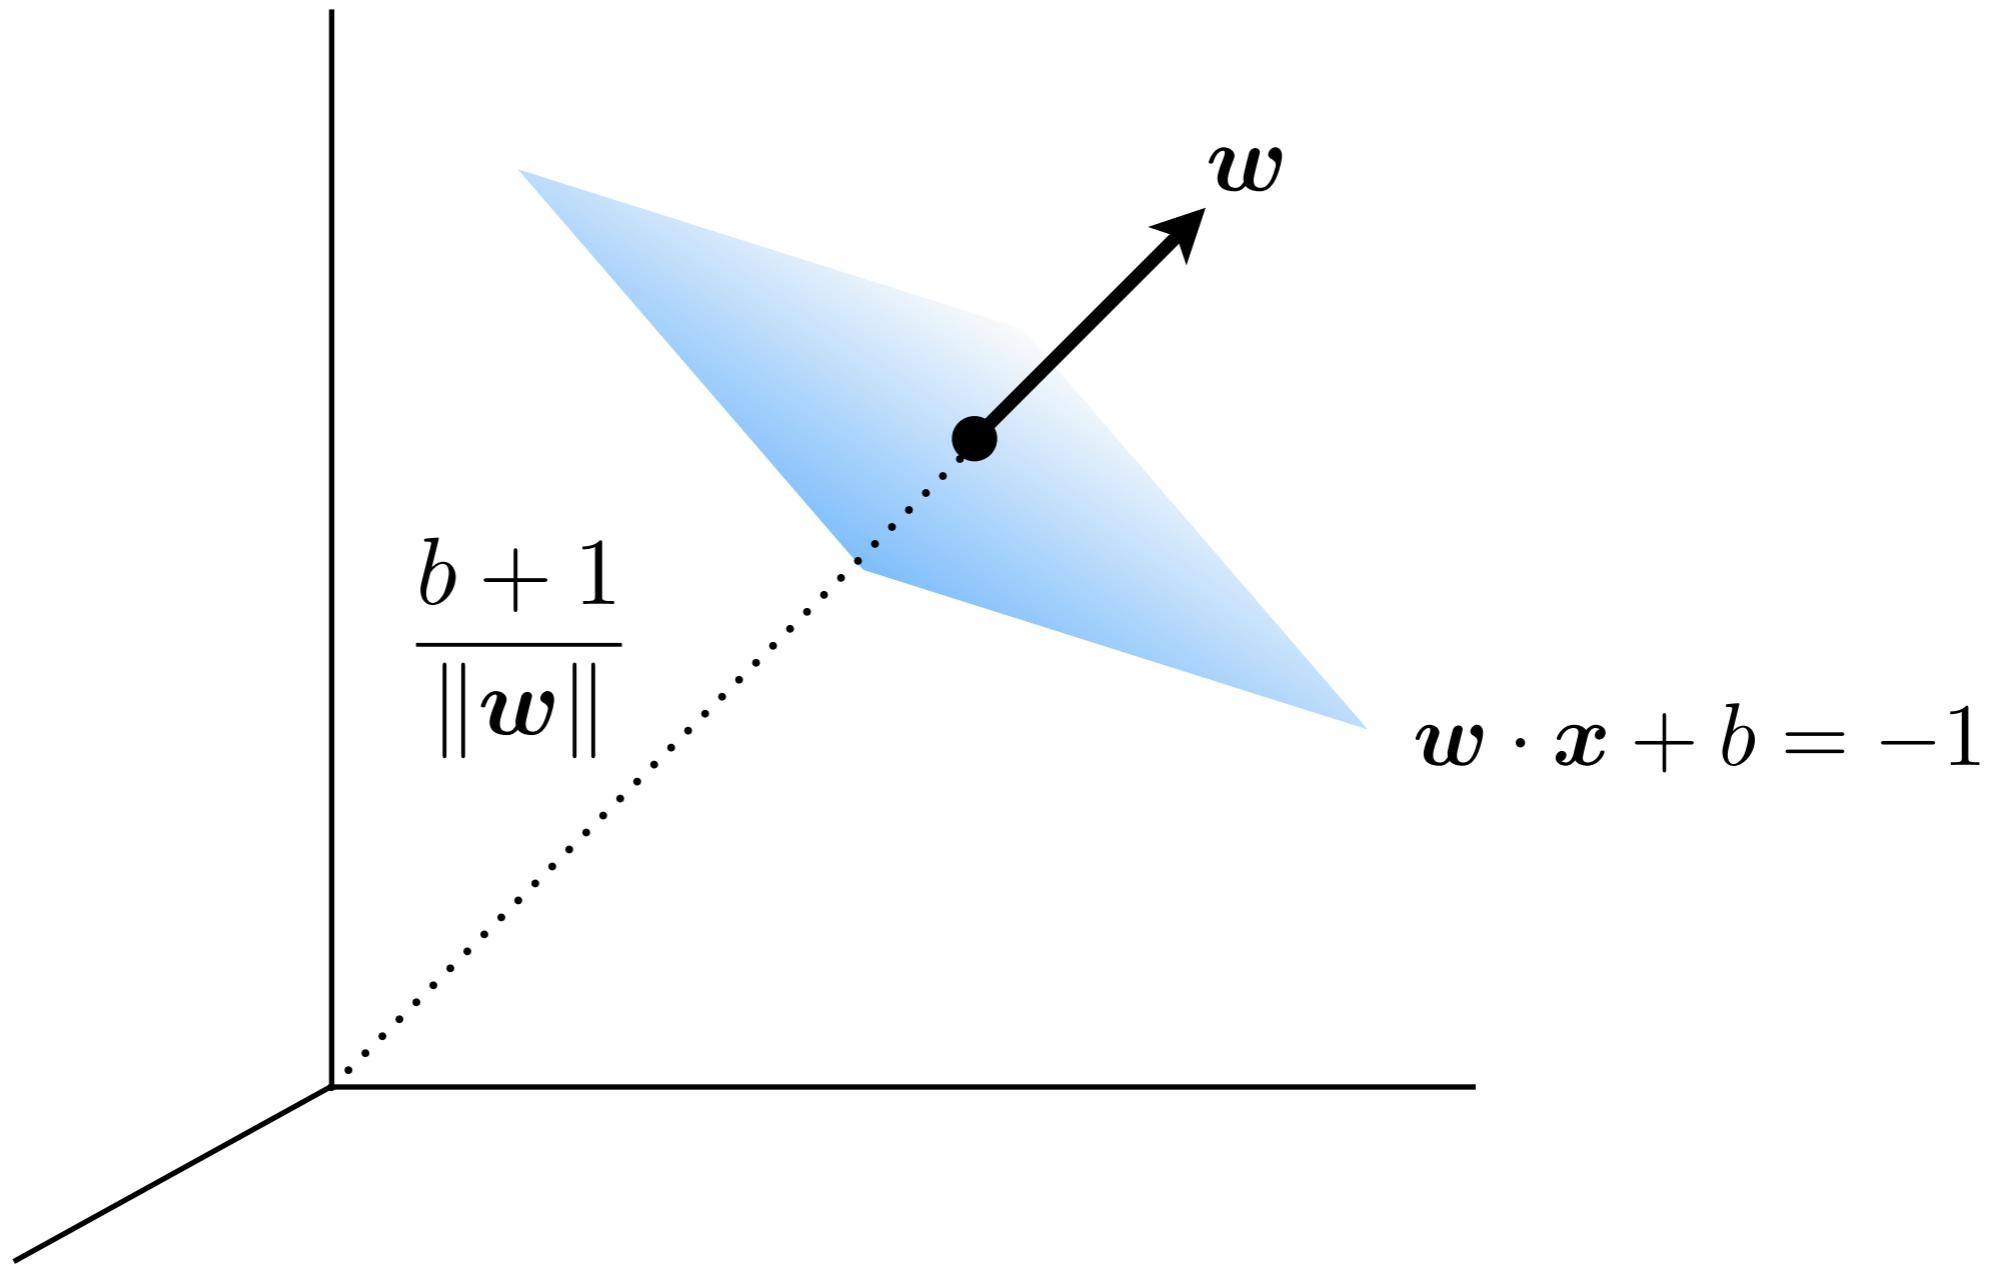
$$\frac{b+1}{\|w\|} - \frac{b}{\|w\|} = \frac{1}{\|w\|}$$

Now we can go to 3D ...

Hyperplanes (planes) in 3D

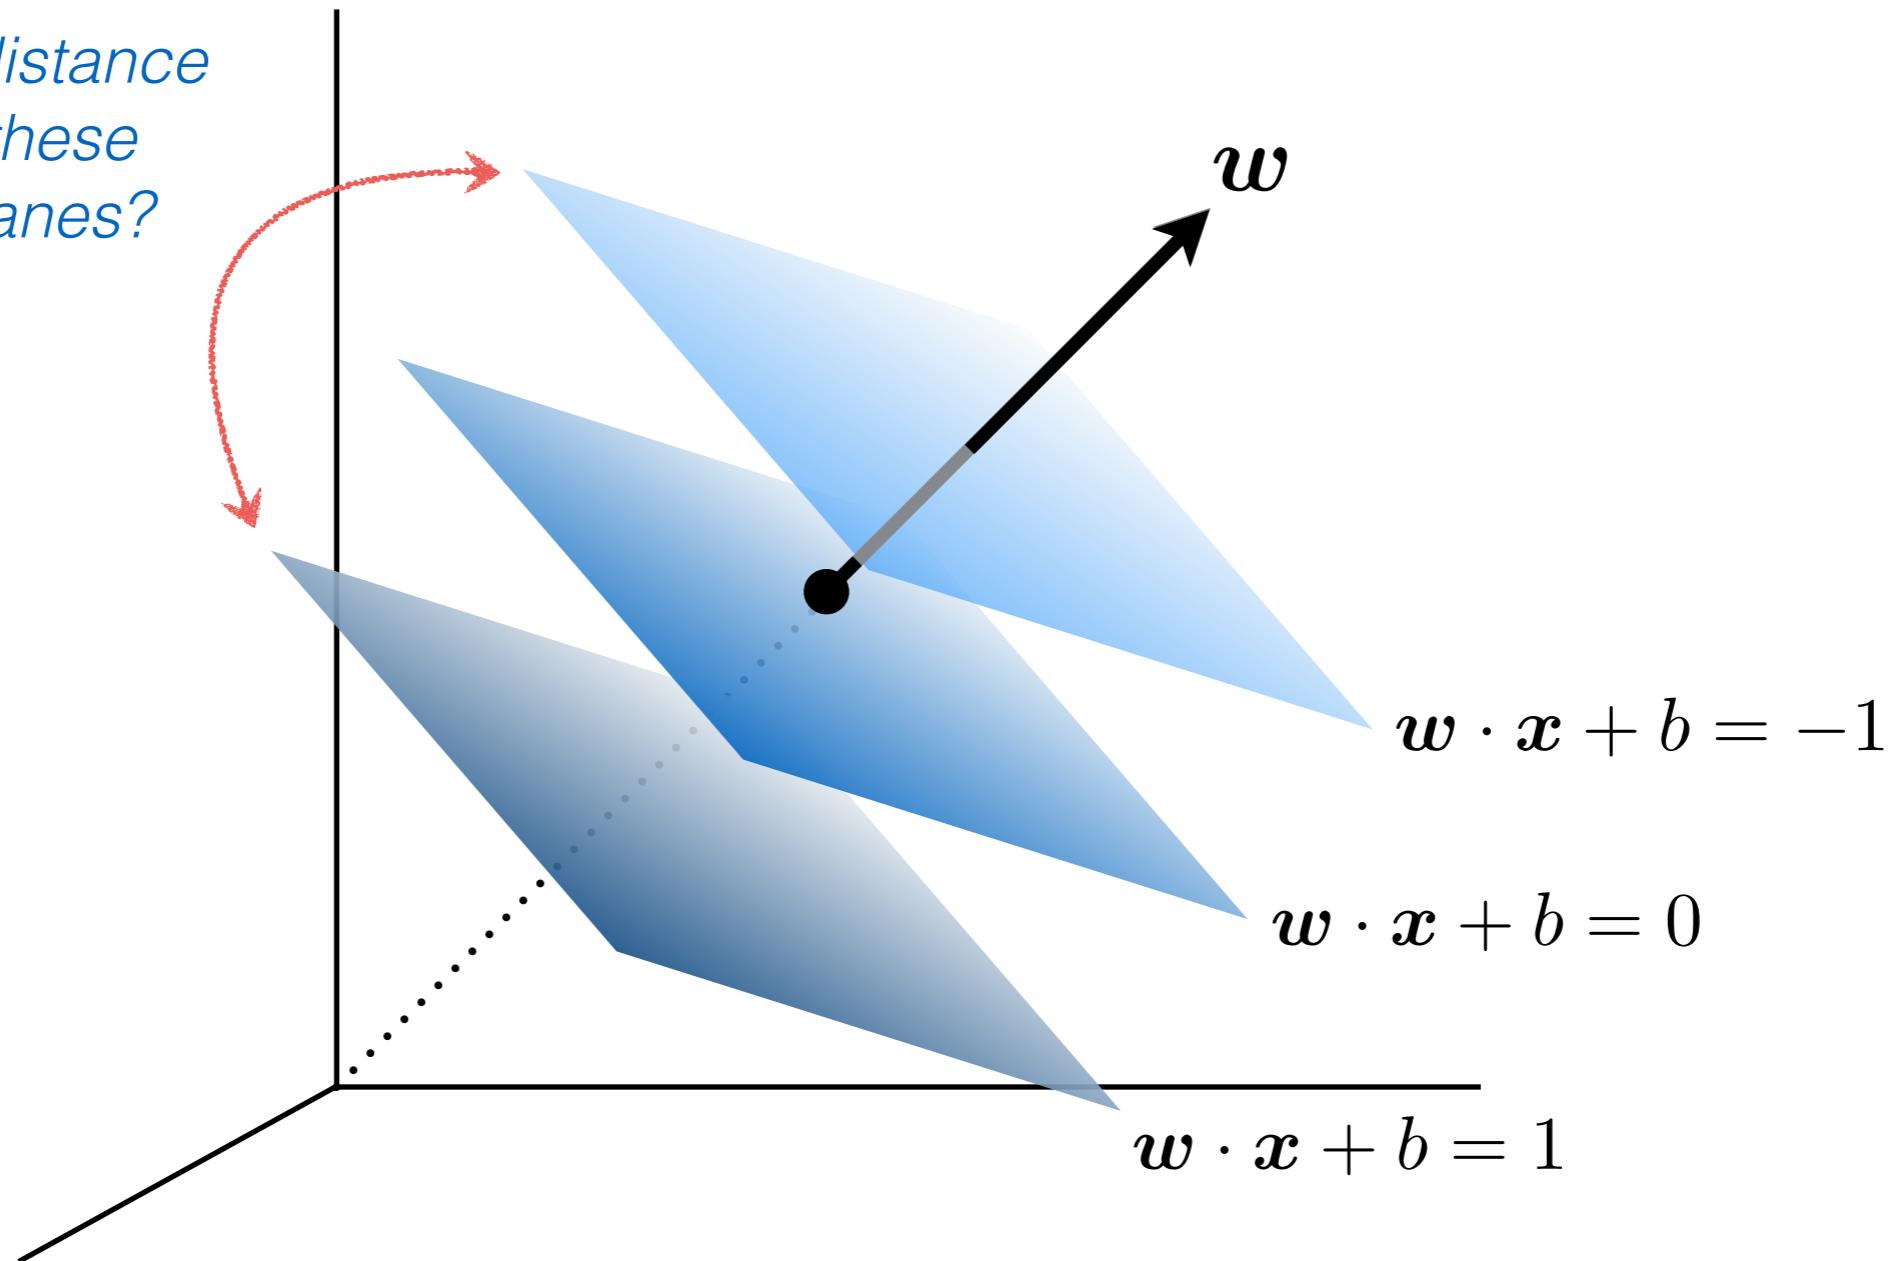


Hyperplanes (planes) in 3D

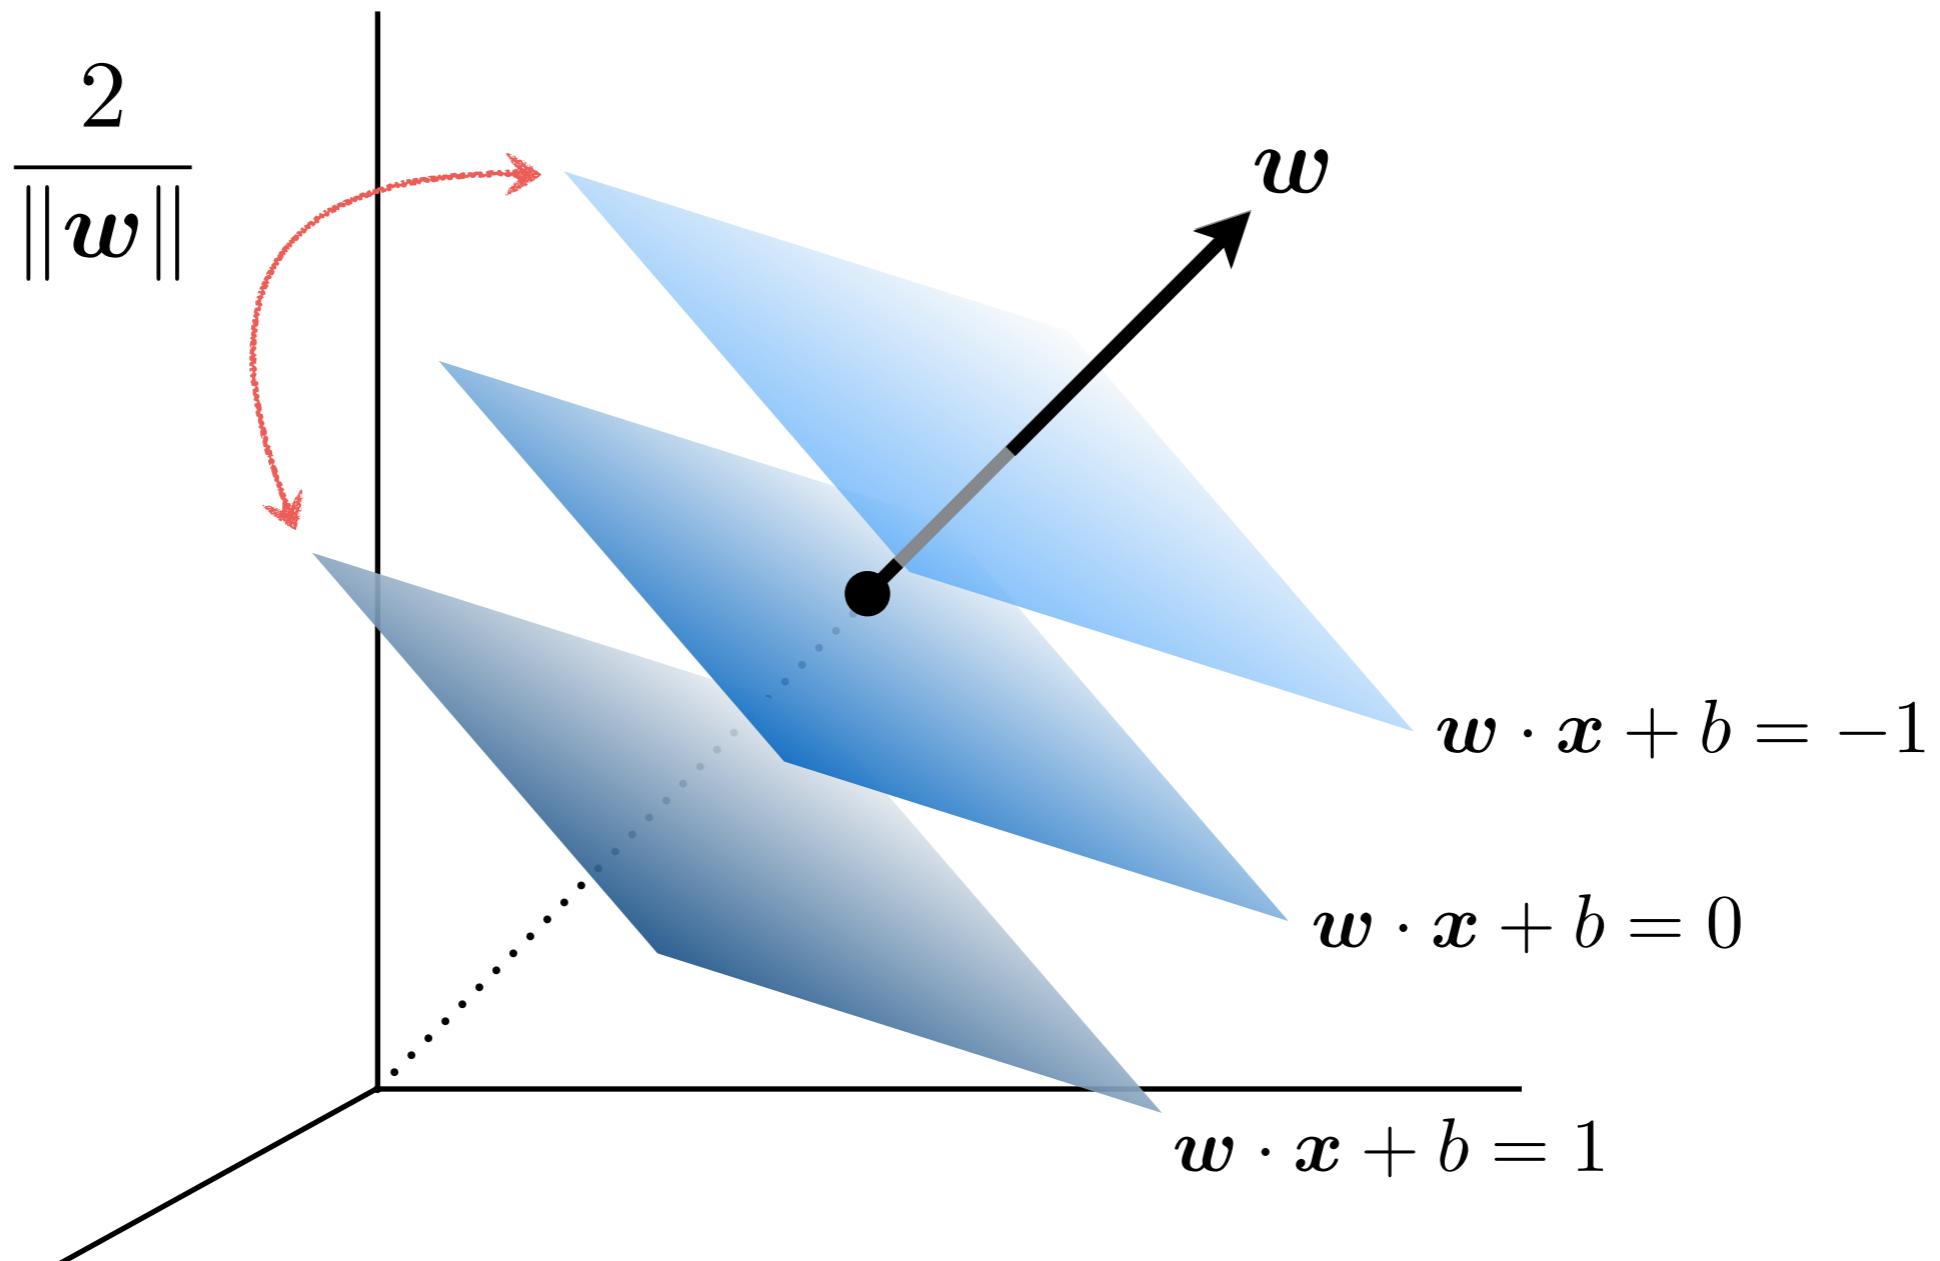


Hyperplanes (planes) in 3D

*What's the distance
between these
parallel planes?*

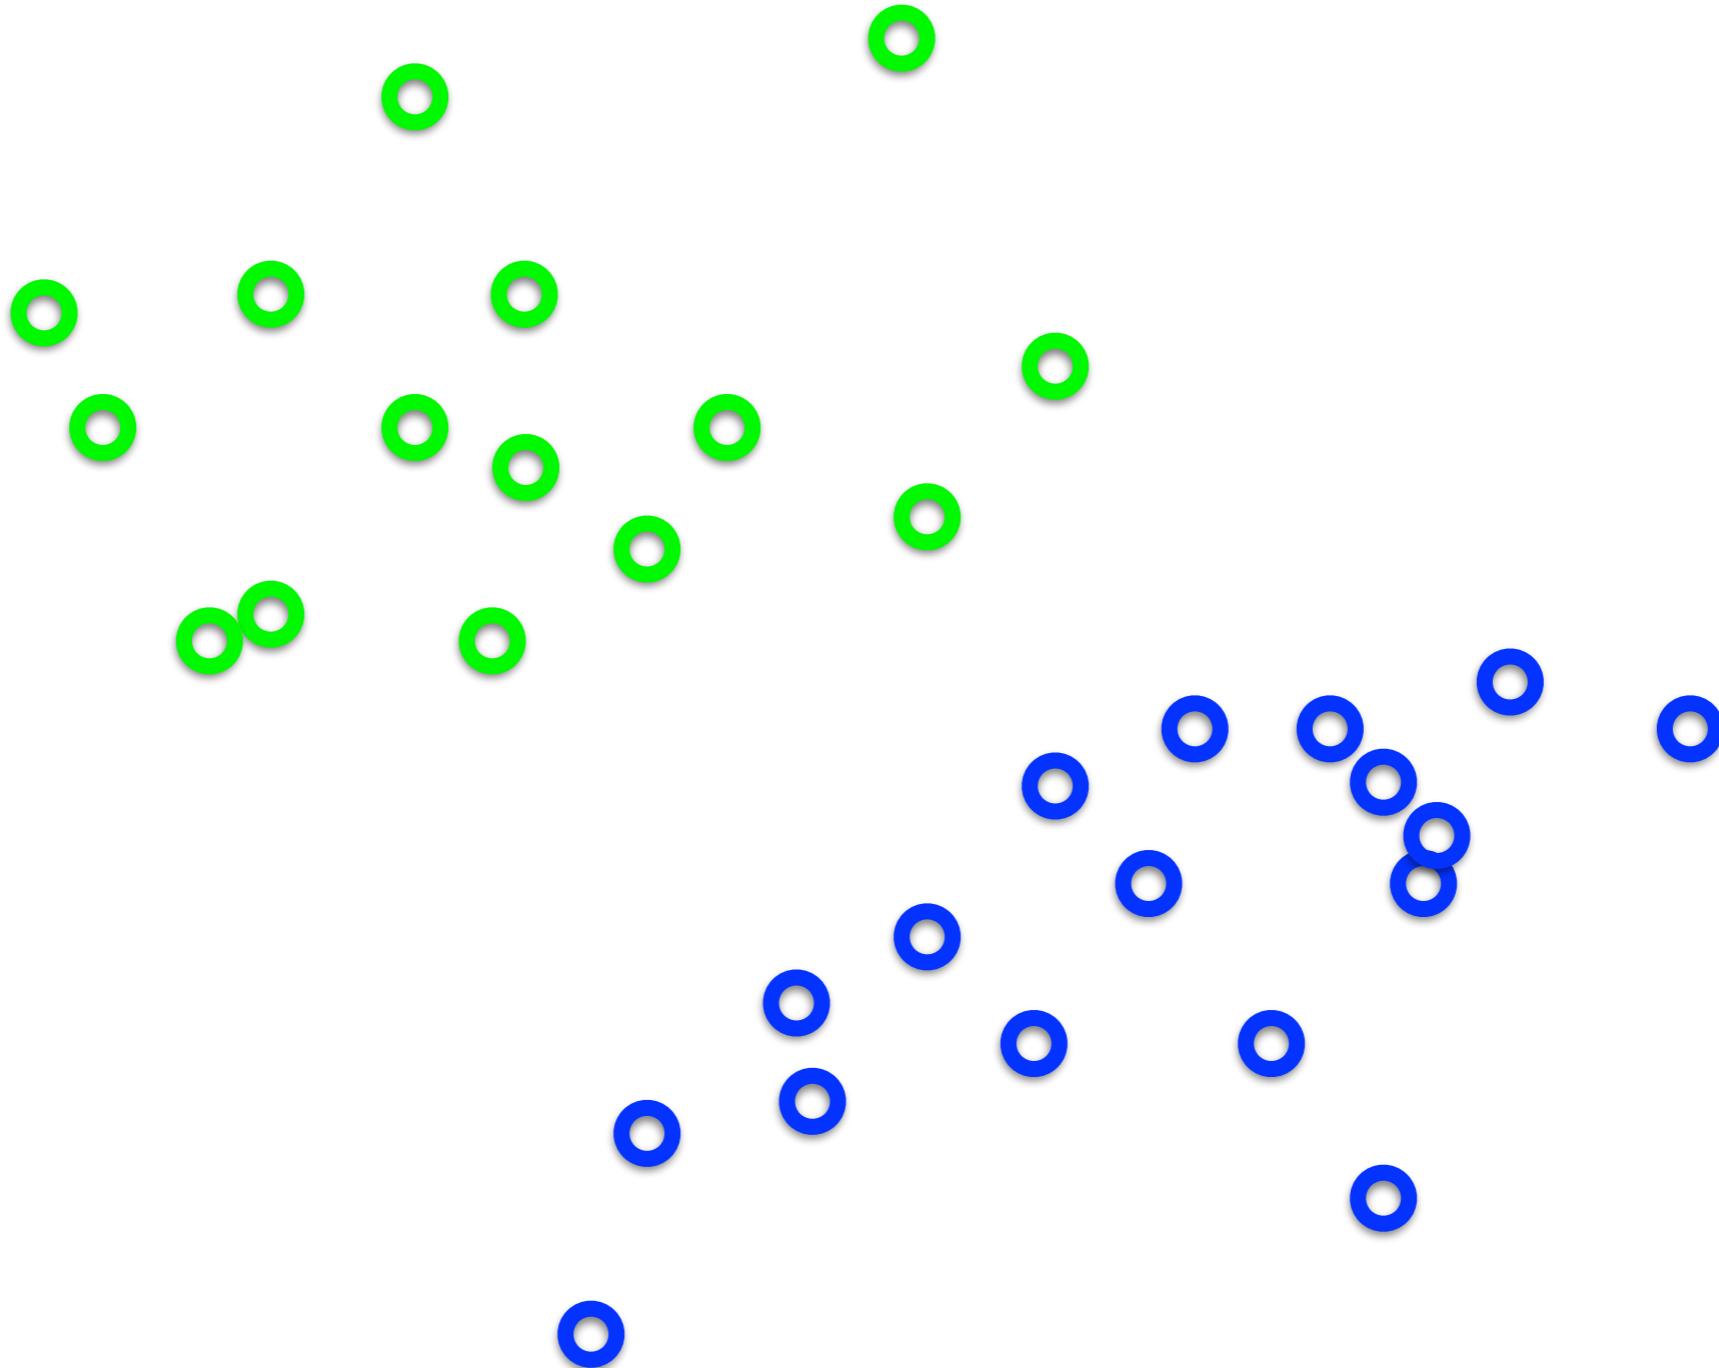


Hyperplanes (planes) in 3D

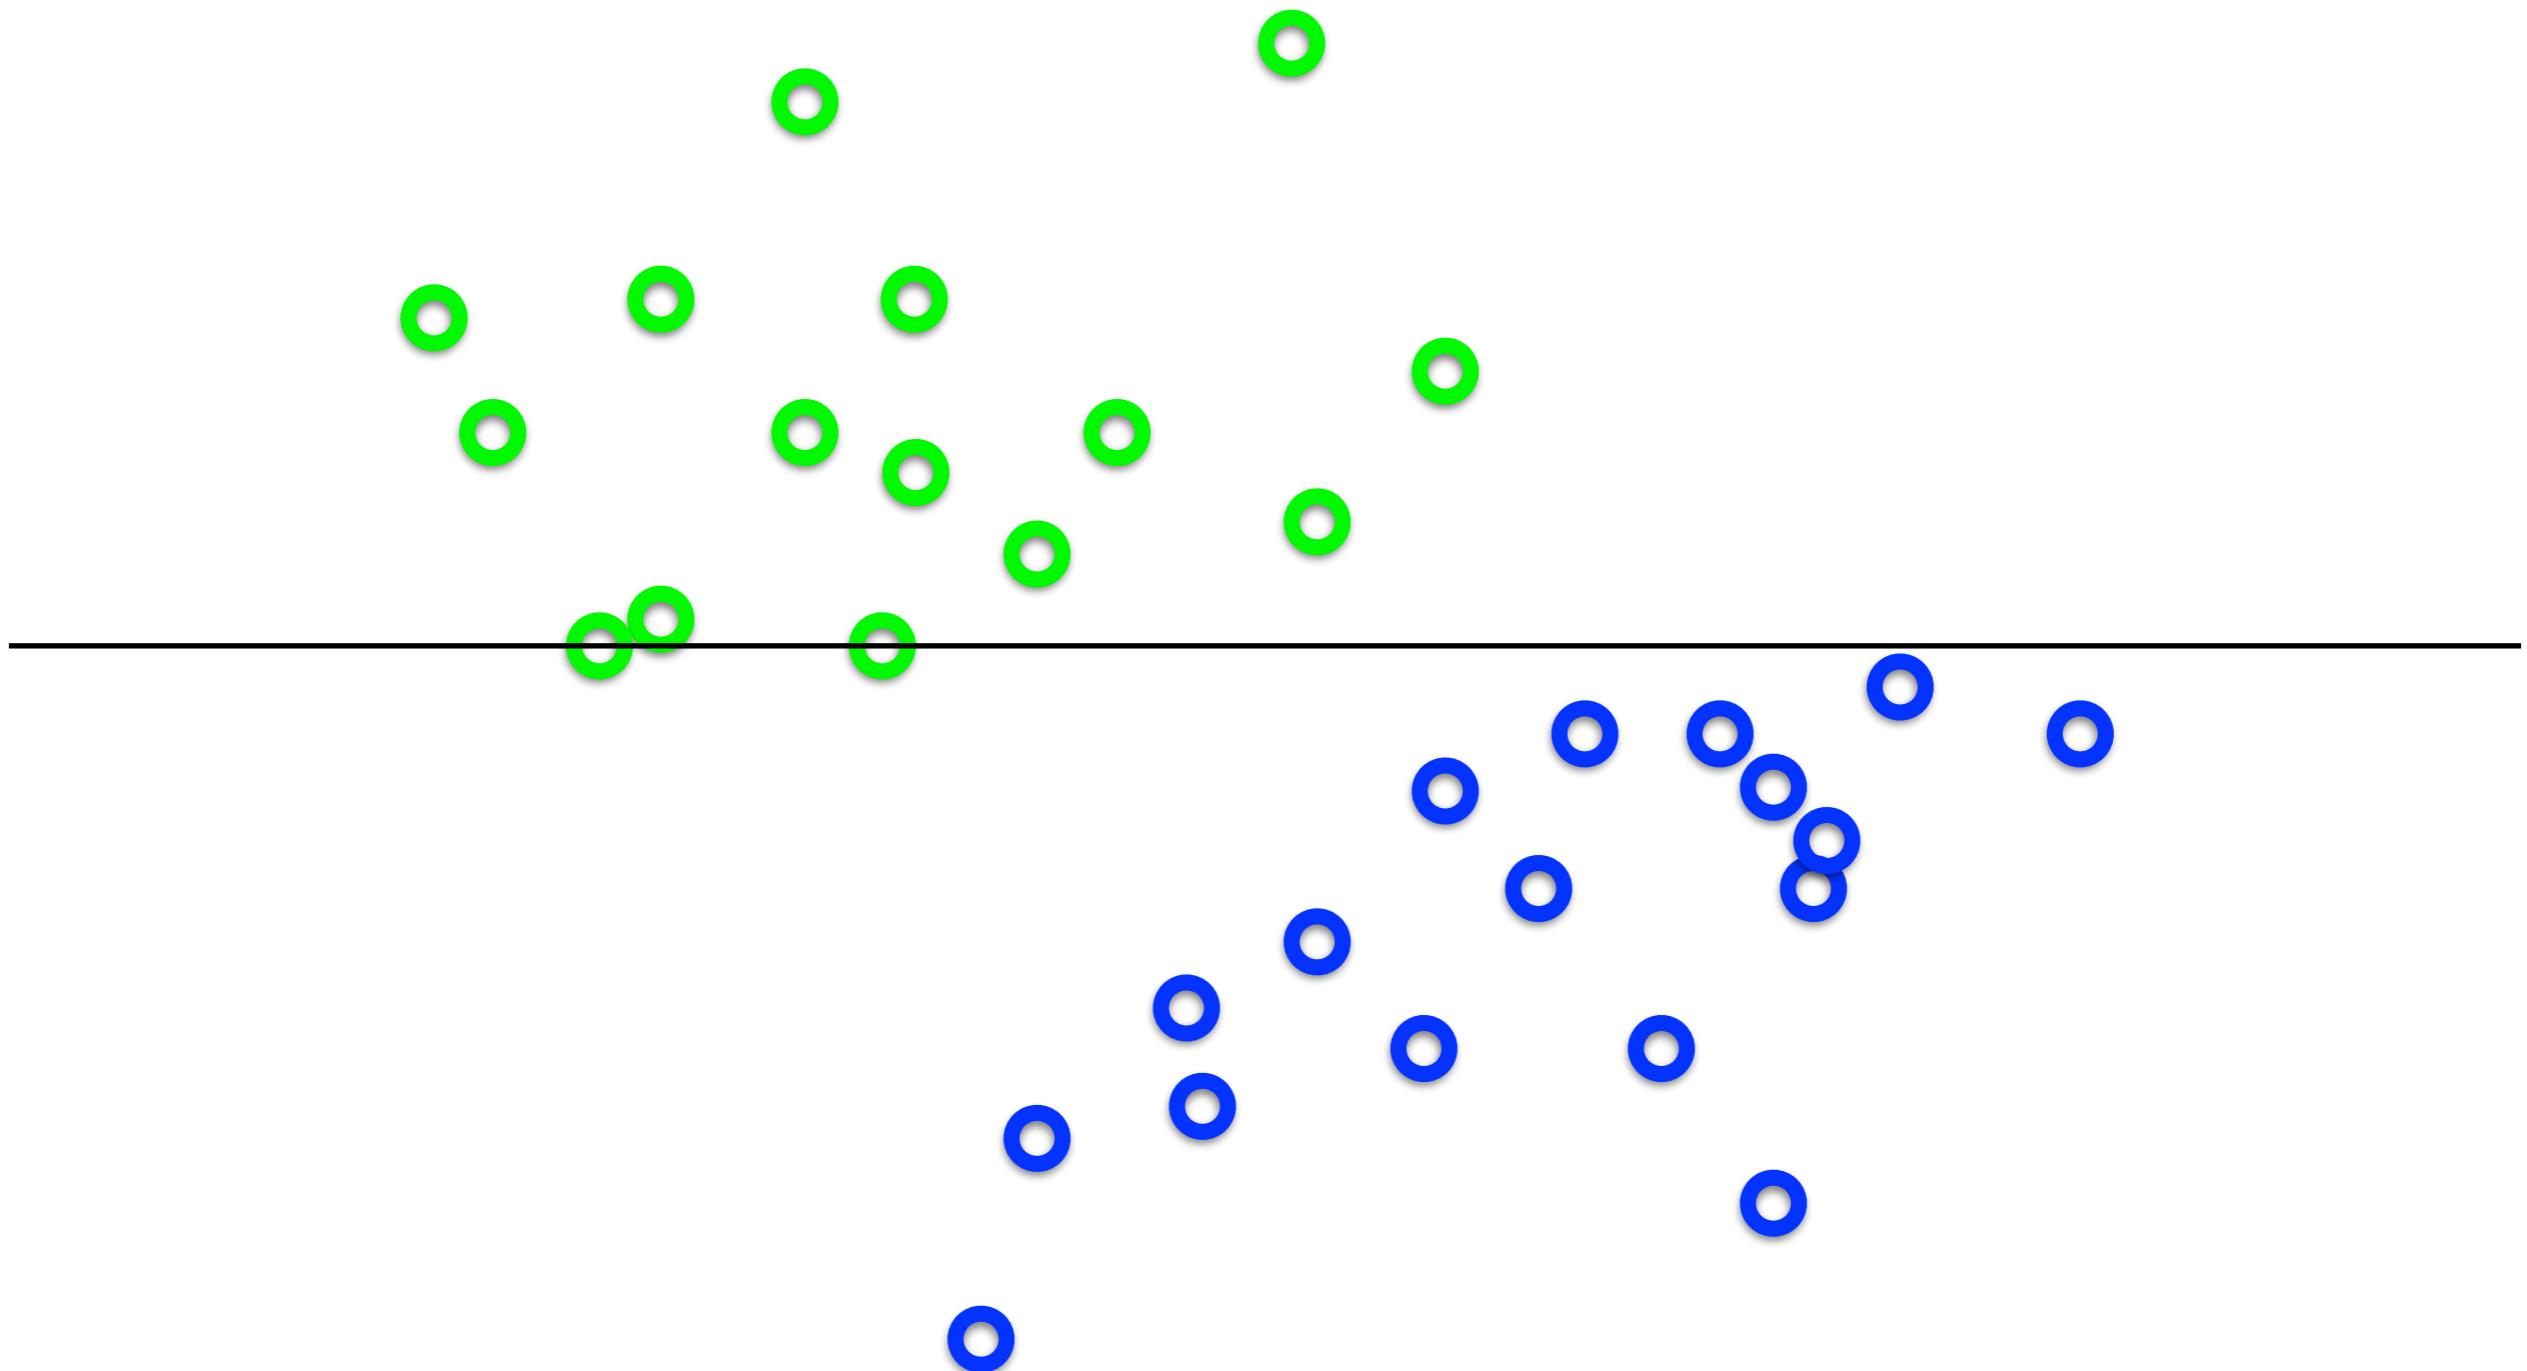


Support Vector Machine

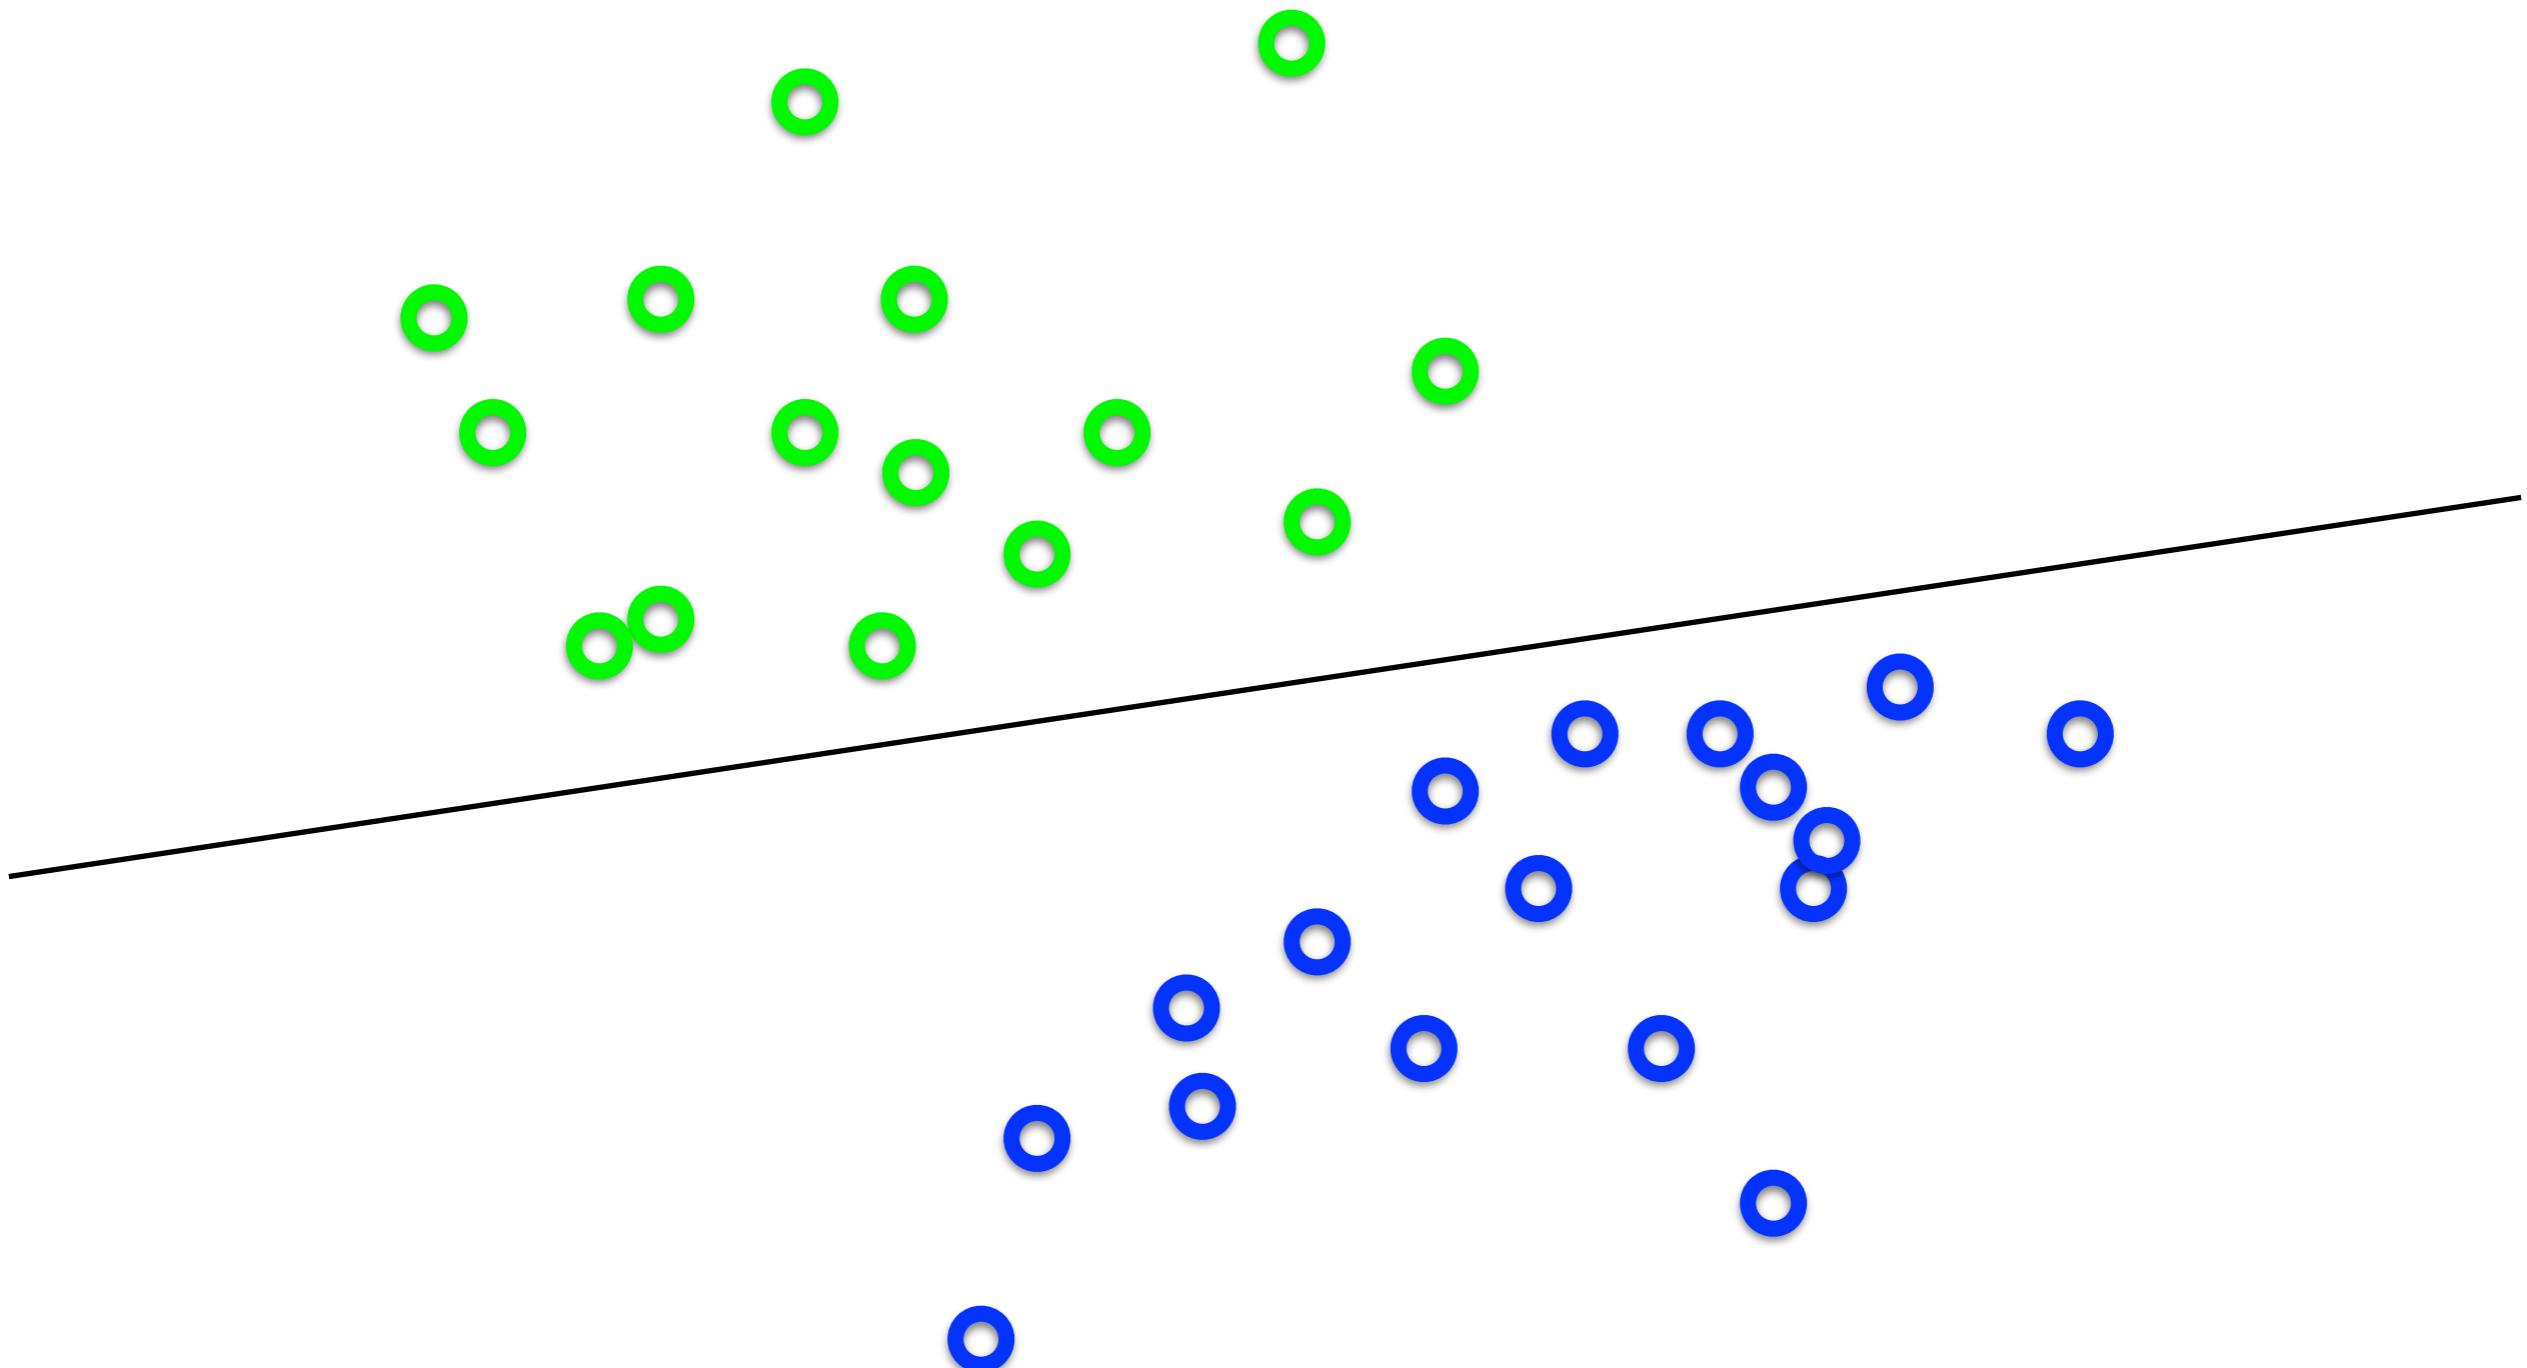
What's the best **w**?



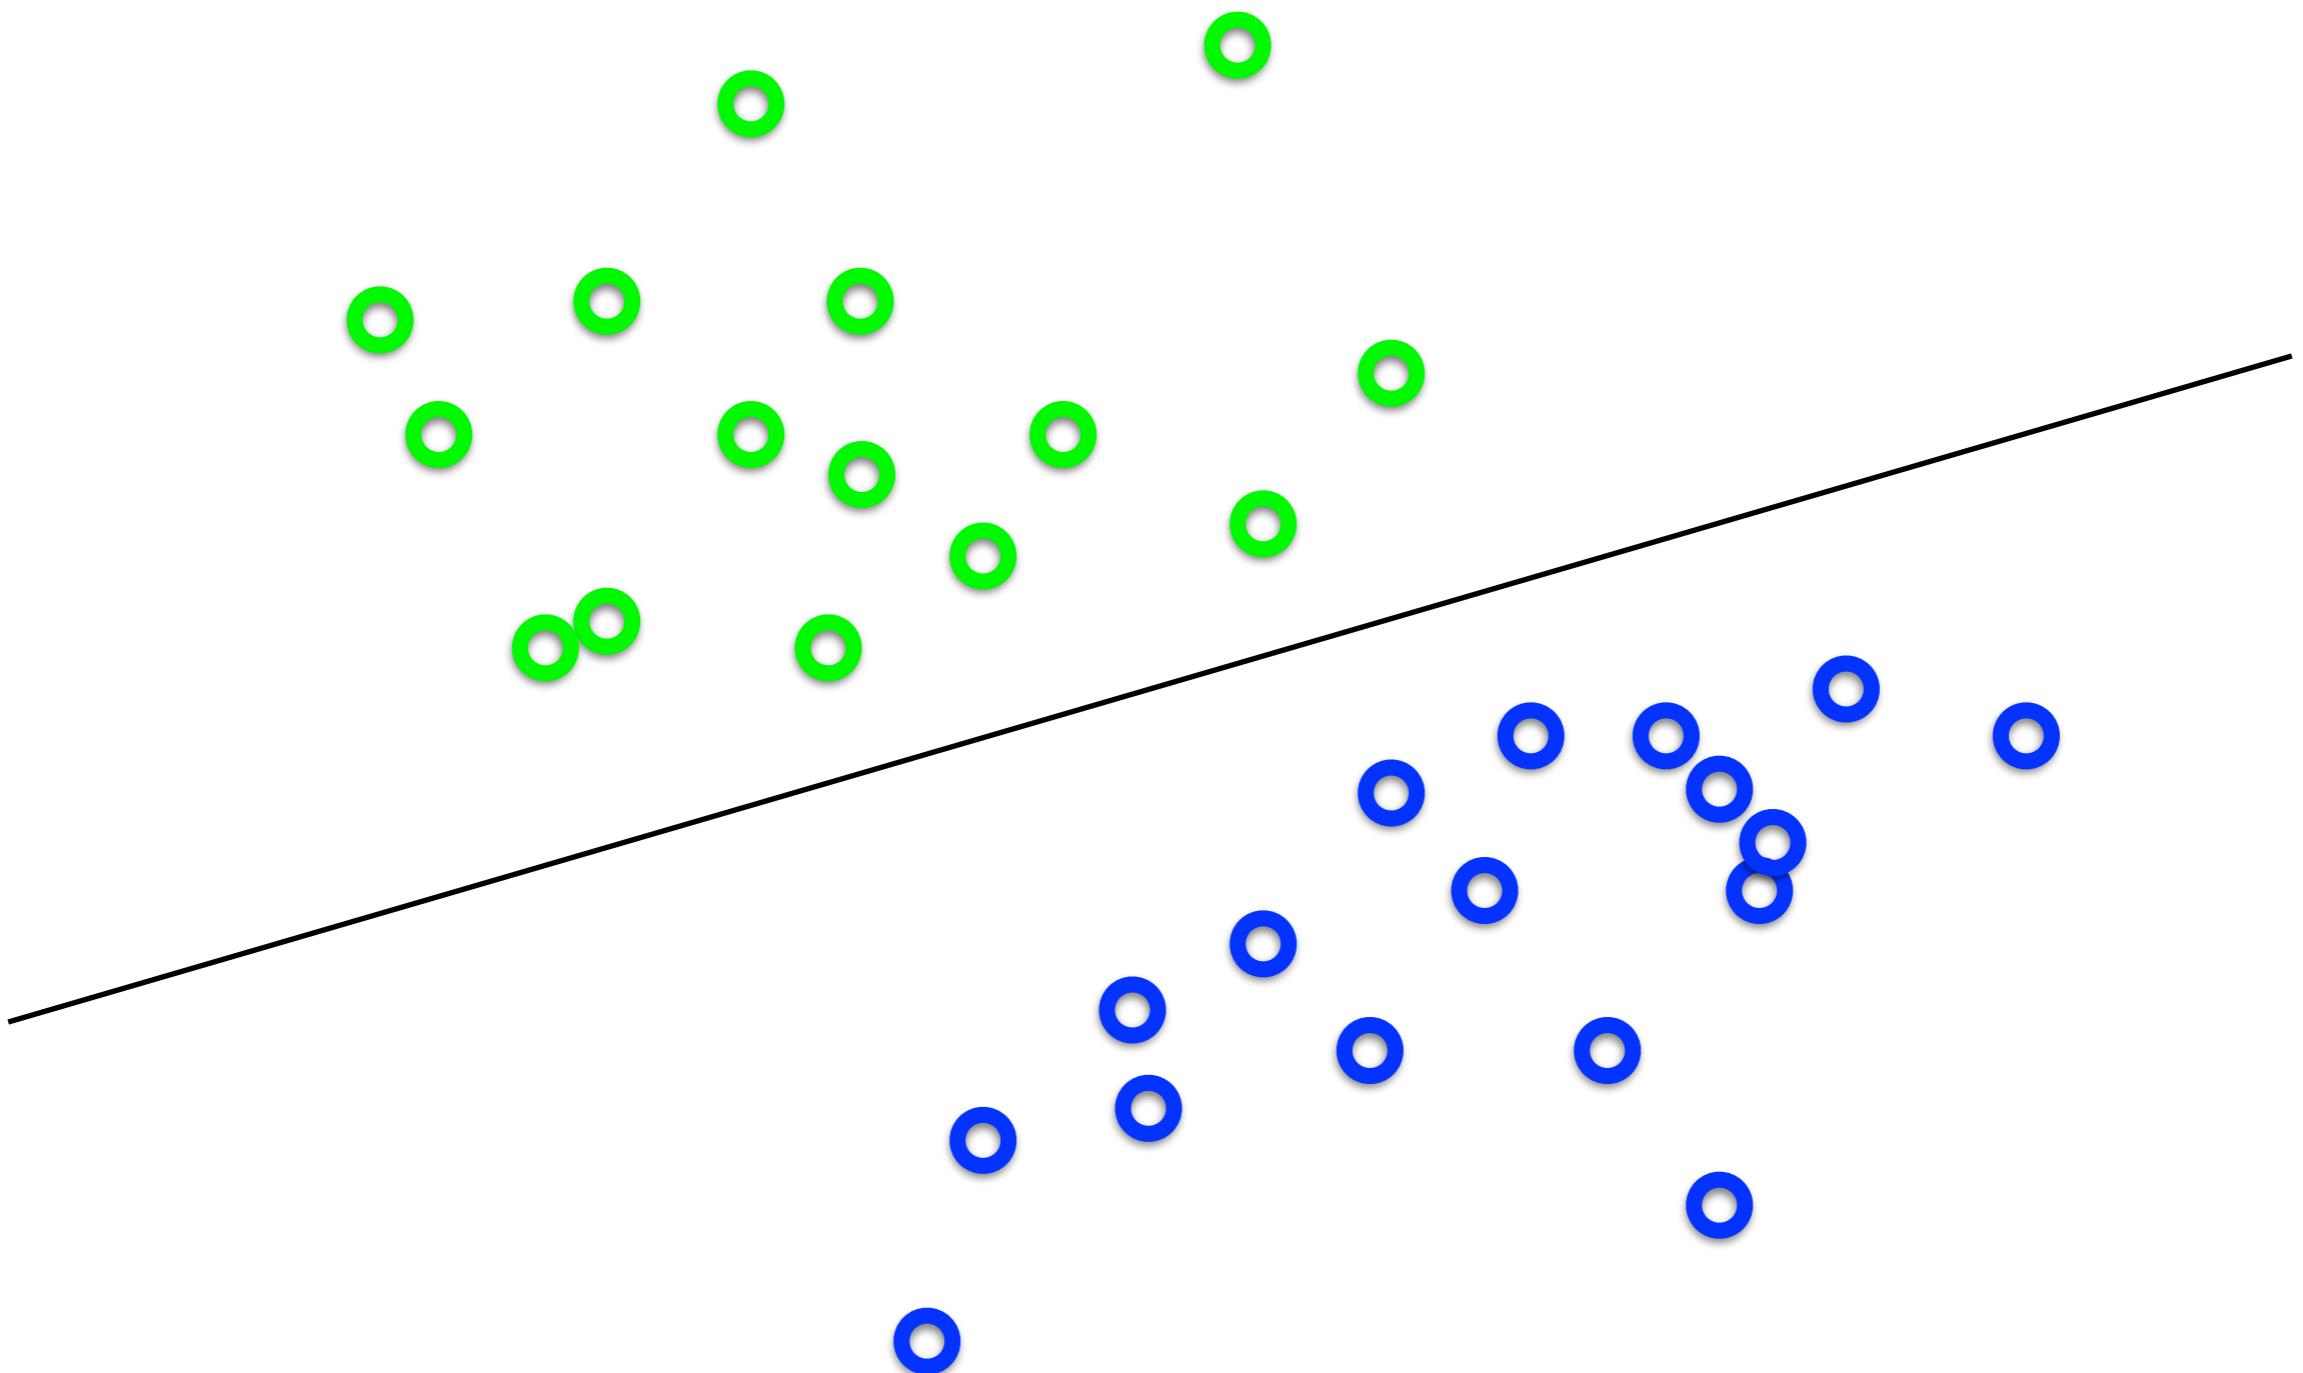
What's the best **w**?



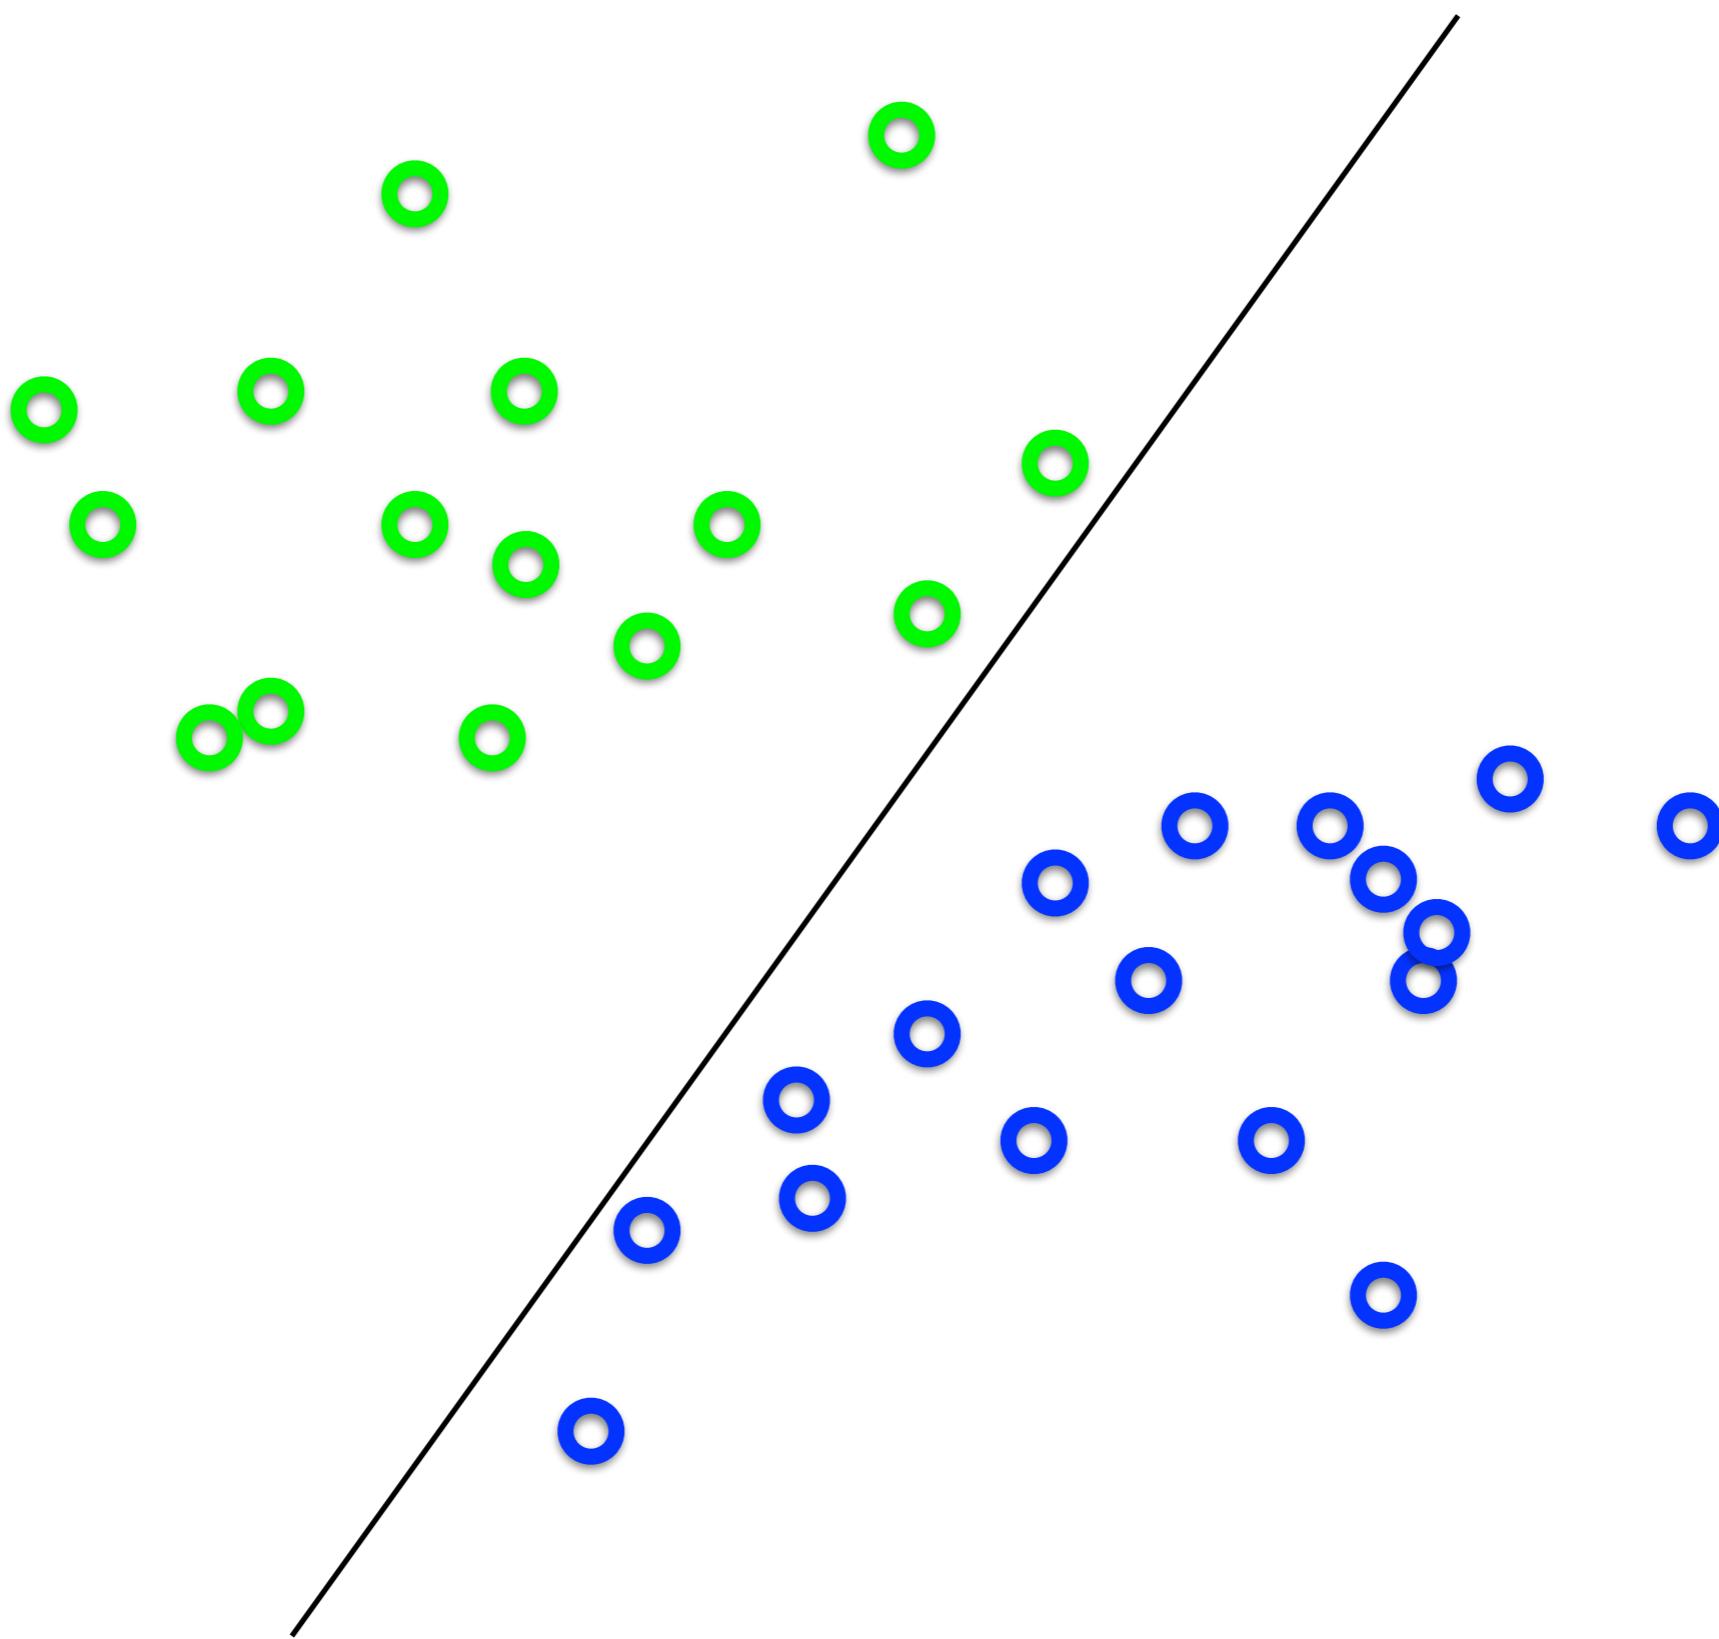
What's the best w ?



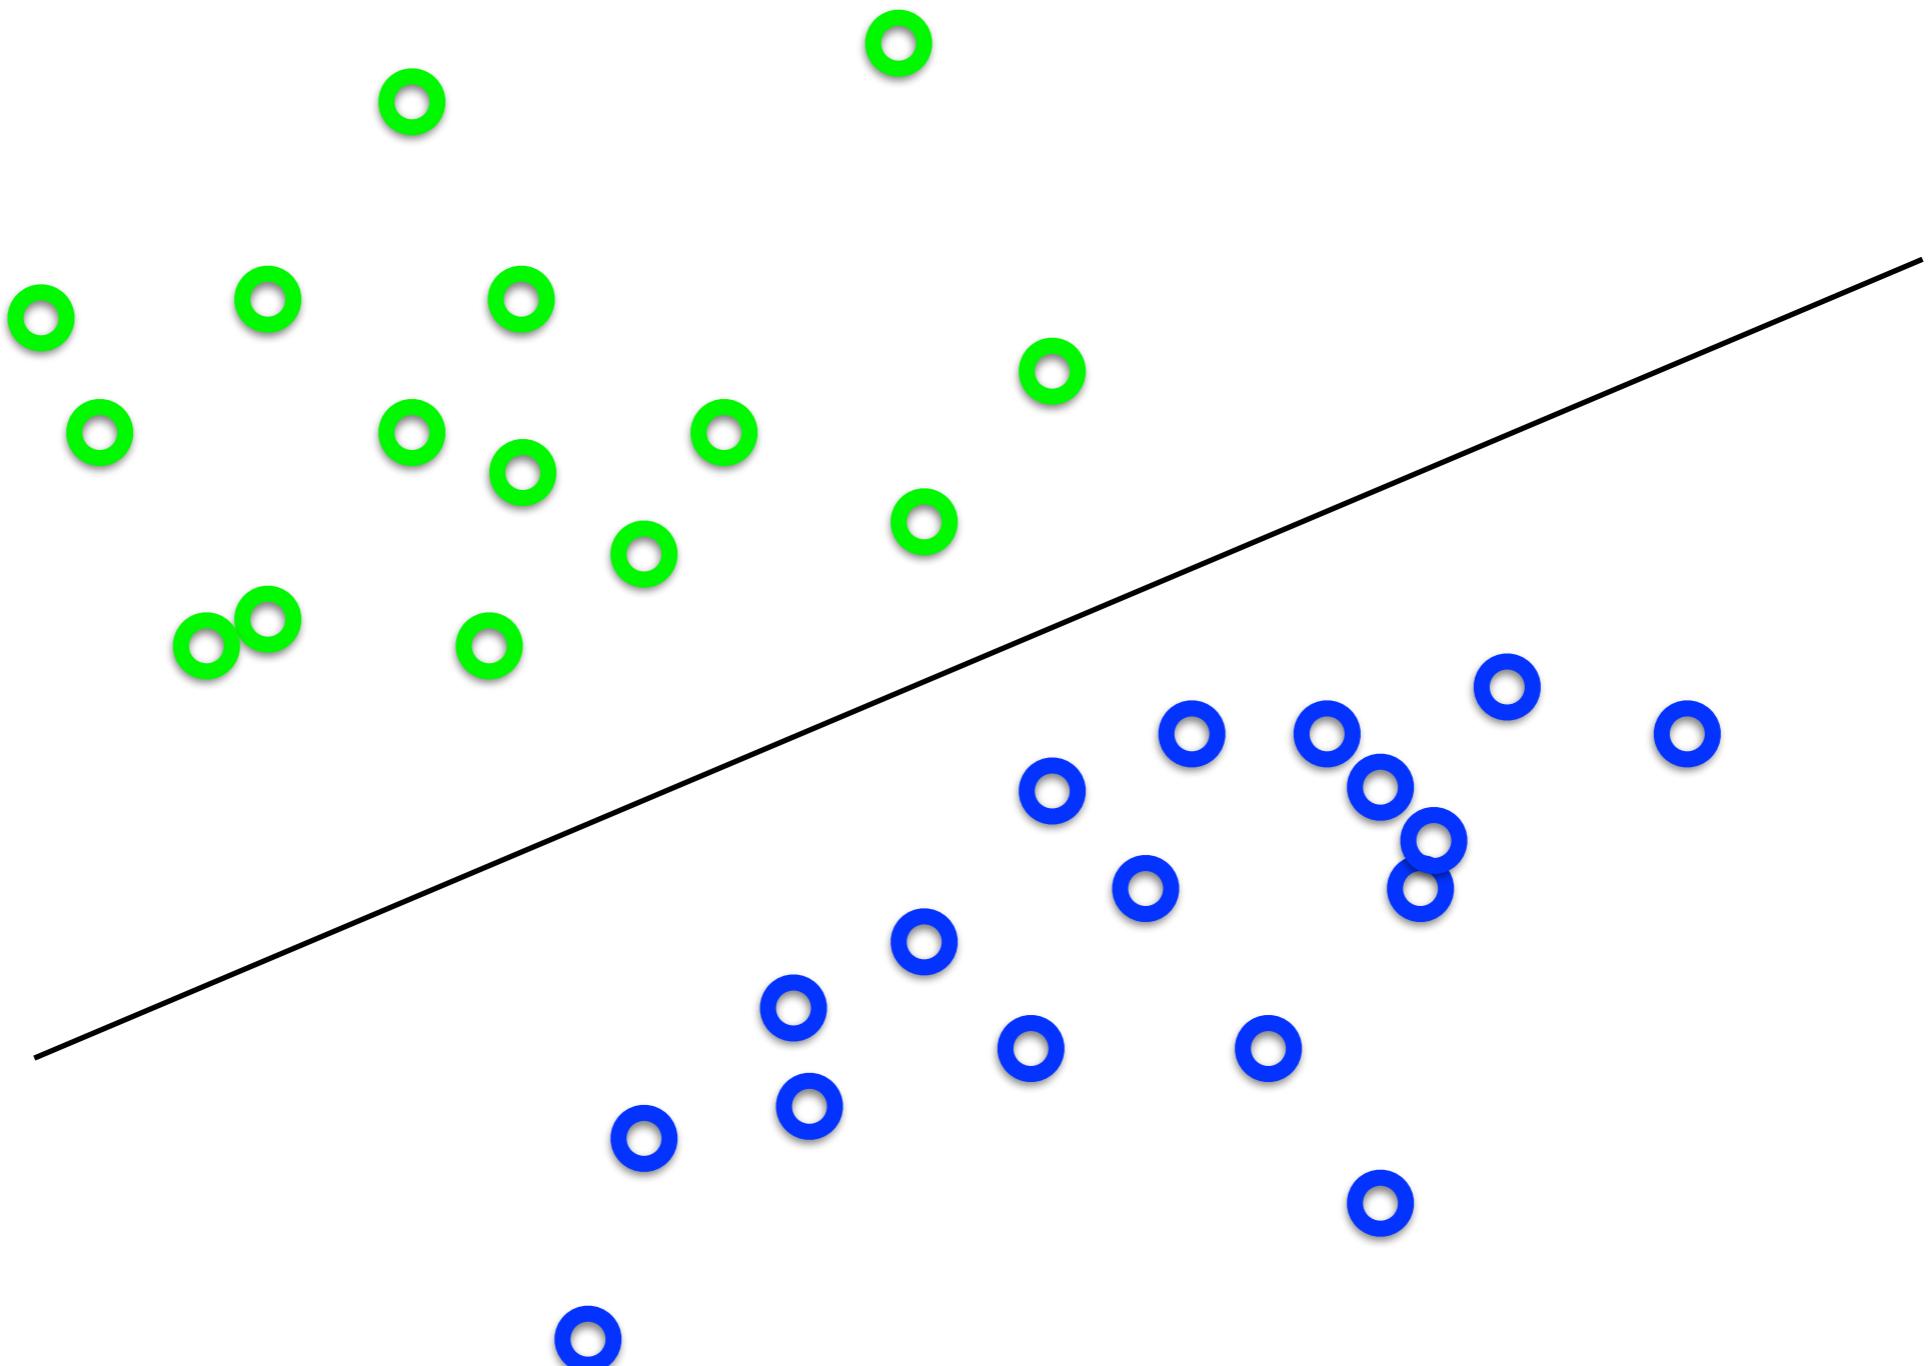
What's the best w ?



What's the best **w**?

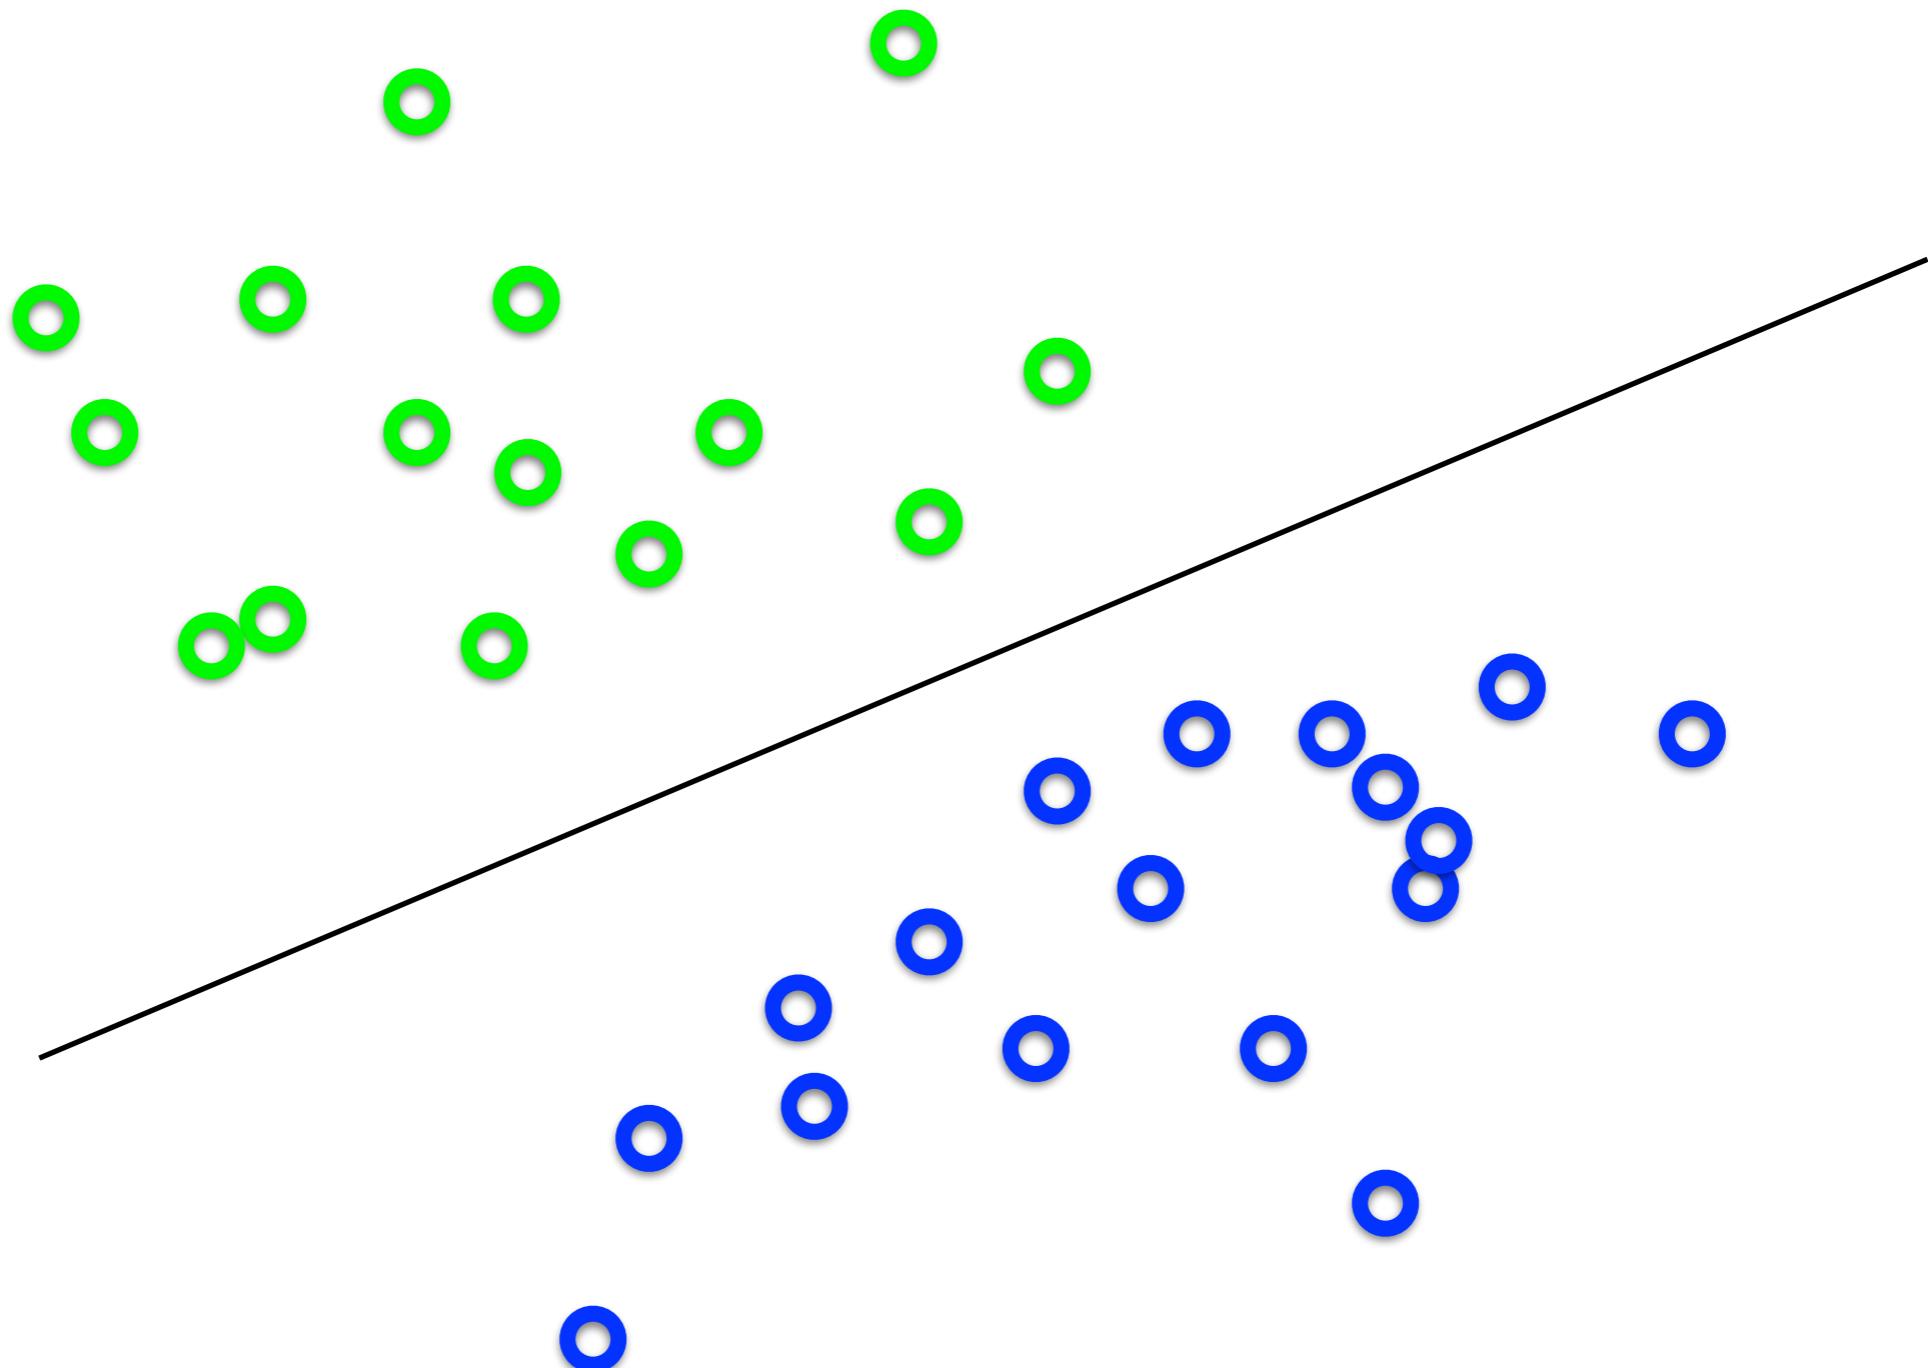


What's the best w ?



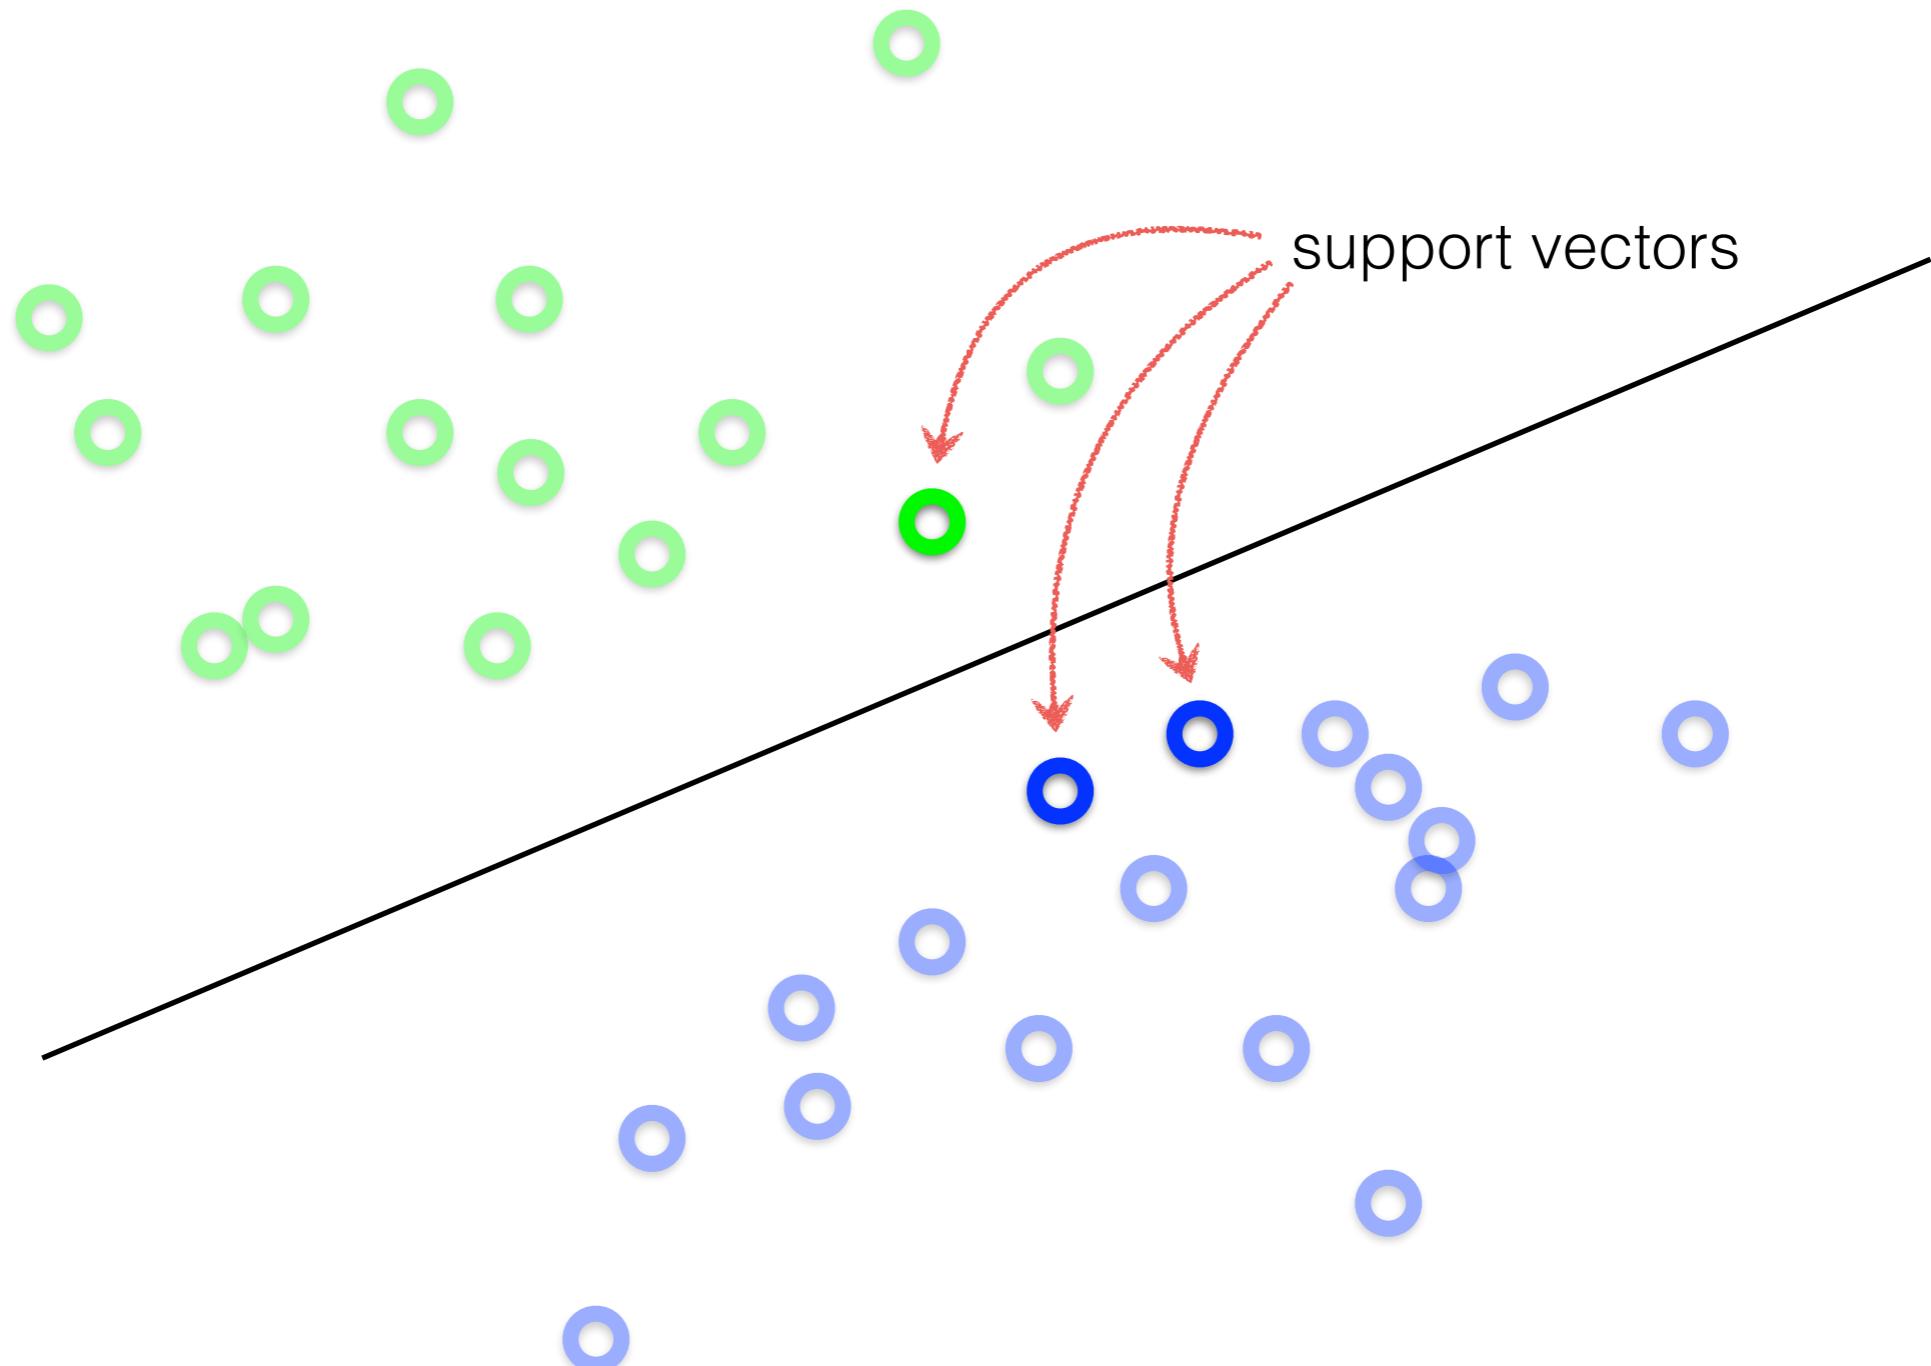
Intuitively, the line that is the
farthest from all interior points

What's the best \mathbf{w} ?



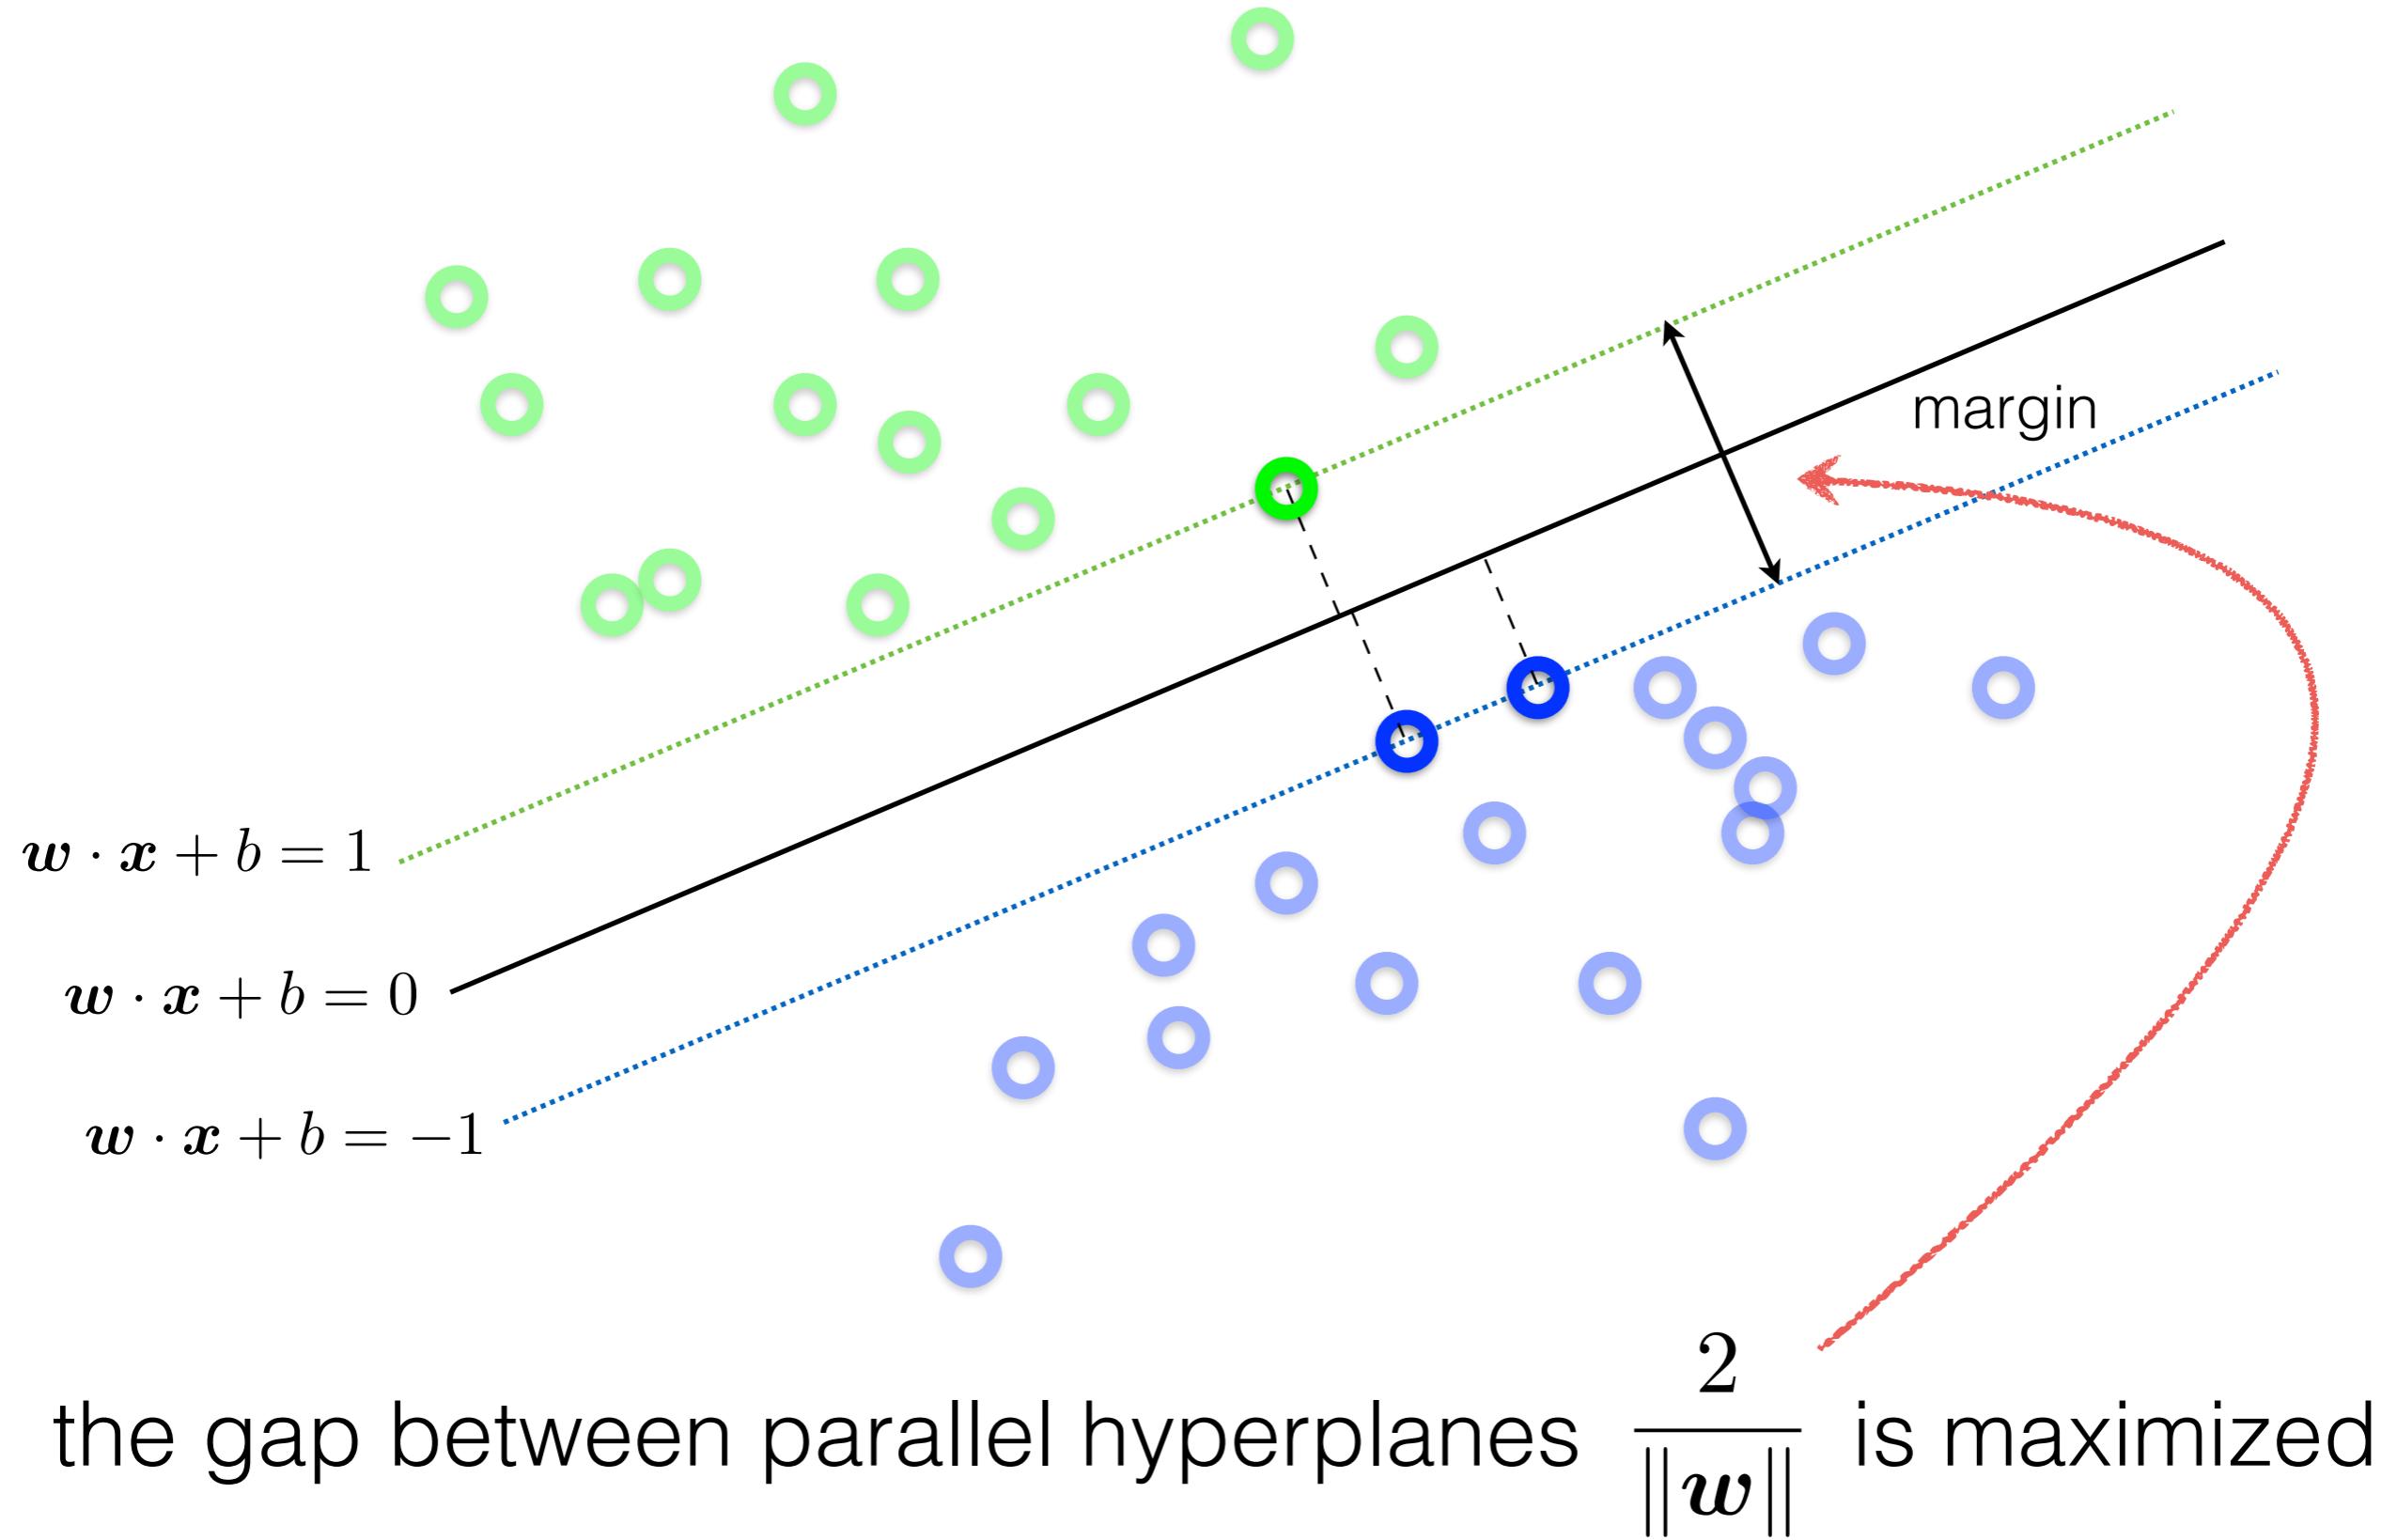
Maximum Margin solution:
most stable to perturbations of data

What's the best \mathbf{w} ?



Want a hyperplane that is far away from ‘inner points’

Find hyperplane \mathbf{w} such that ...

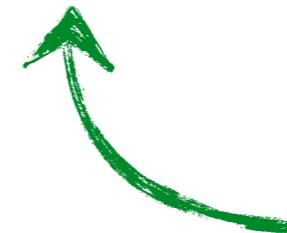


Can be formulated as a maximization problem

$$\max_{\mathbf{w}} \frac{2}{\|\mathbf{w}\|}$$

$$\text{subject to } \mathbf{w} \cdot \mathbf{x}_i + b \begin{cases} \geq +1 & \text{if } y_i = +1 \\ \leq -1 & \text{if } y_i = -1 \end{cases} \text{ for } i = 1, \dots, N$$

What does this constraint mean?



label of the data point

Why is it $+1$ and -1 ?

Can be formulated as a maximization problem

$$\max_{\mathbf{w}} \frac{2}{\|\mathbf{w}\|}$$

subject to $\mathbf{w} \cdot \mathbf{x}_i + b \begin{cases} \geq +1 & \text{if } y_i = +1 \\ \leq -1 & \text{if } y_i = -1 \end{cases}$ for $i = 1, \dots, N$

Equivalently,

Where did the 2 go?

$$\min_{\mathbf{w}} \|\mathbf{w}\|$$

subject to $y_i(\mathbf{w} \cdot \mathbf{x}_i + b) \geq 1$ for $i = 1, \dots, N$

What happened to the labels?

‘Primal formulation’ of a linear SVM

$$\min_{\boldsymbol{w}} \|\boldsymbol{w}\|$$

Objective Function

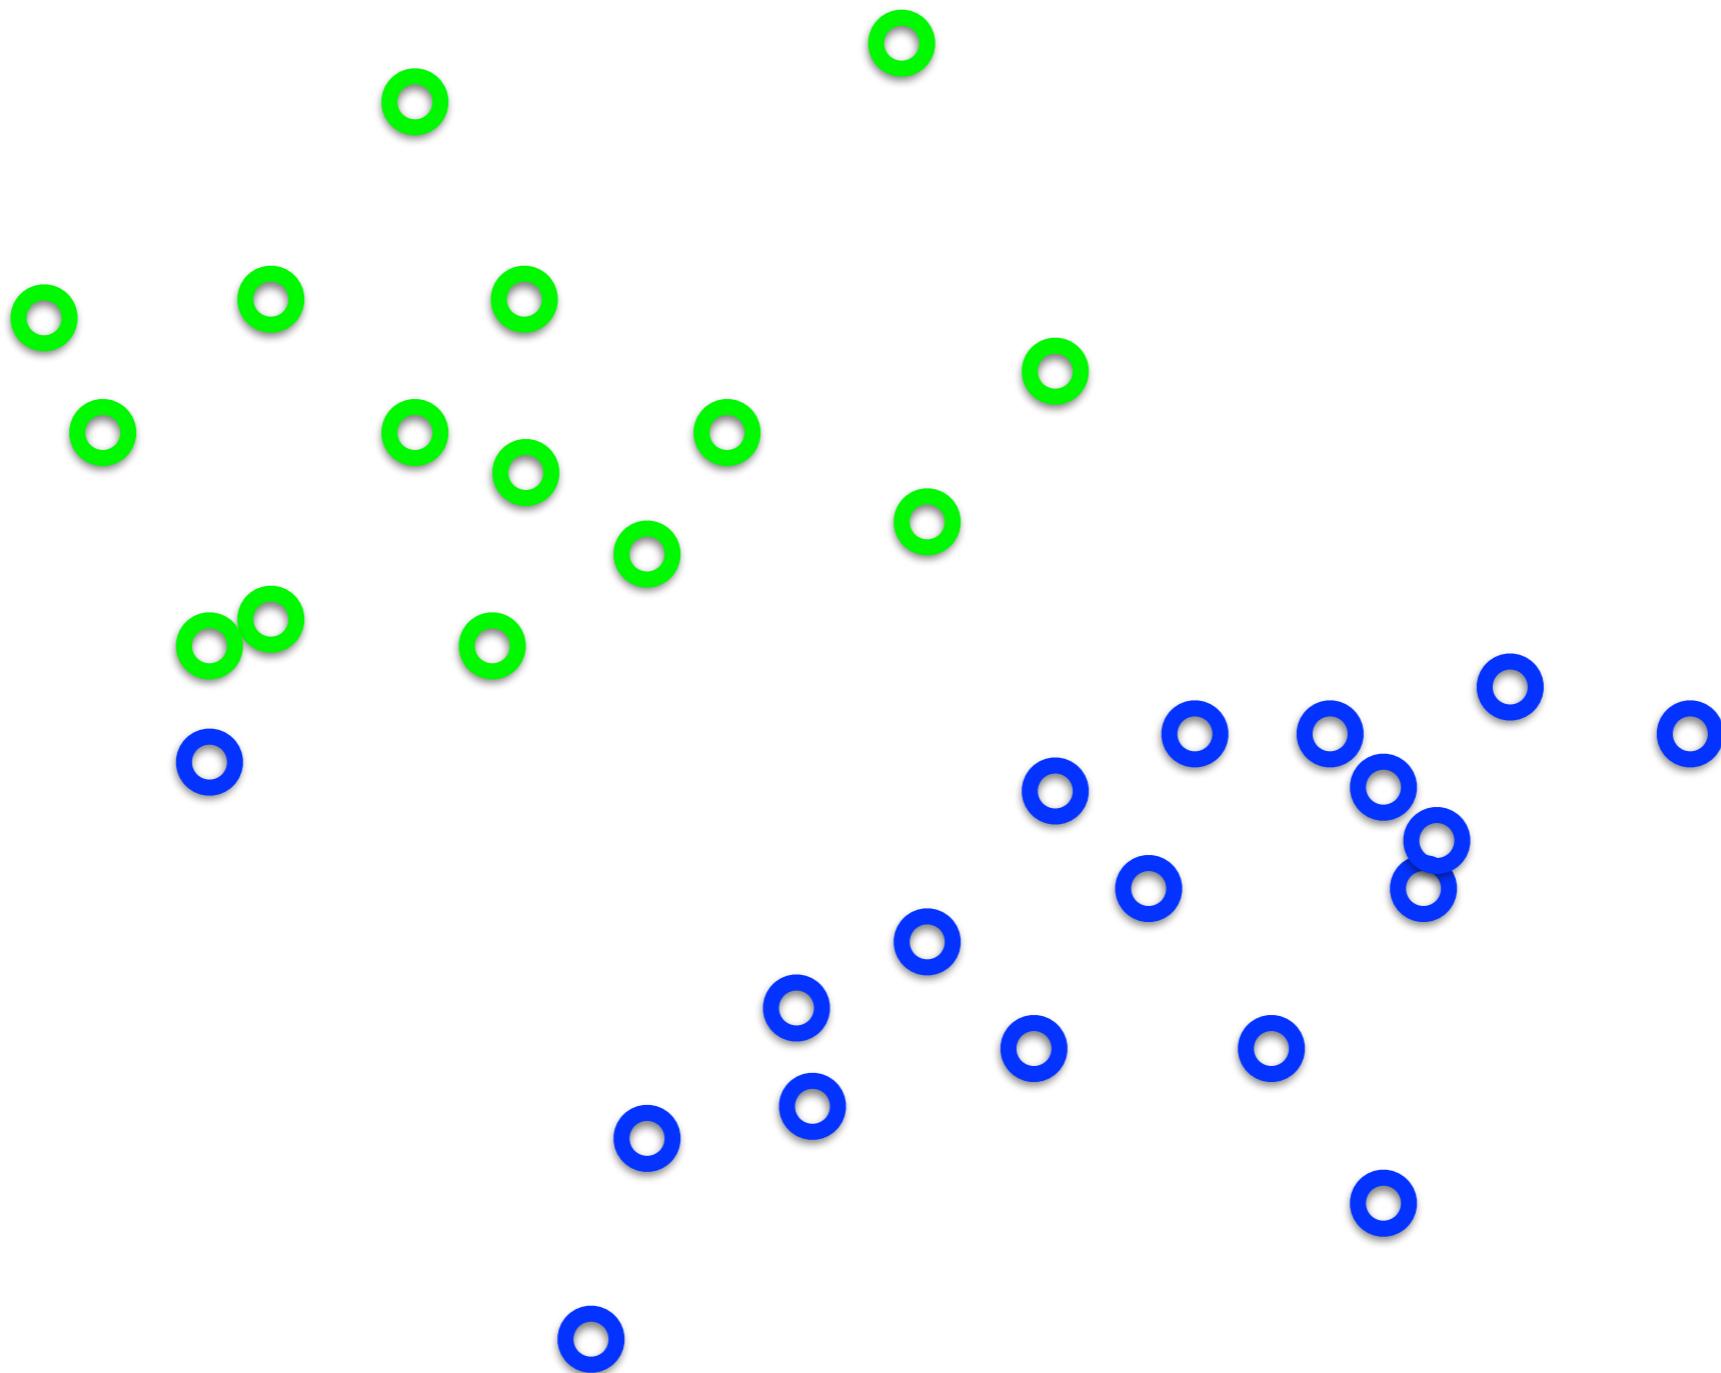
subject to $y_i(\boldsymbol{w} \cdot \boldsymbol{x}_i + b) \geq 1 \quad \text{for } i = 1, \dots, N$

Constraints

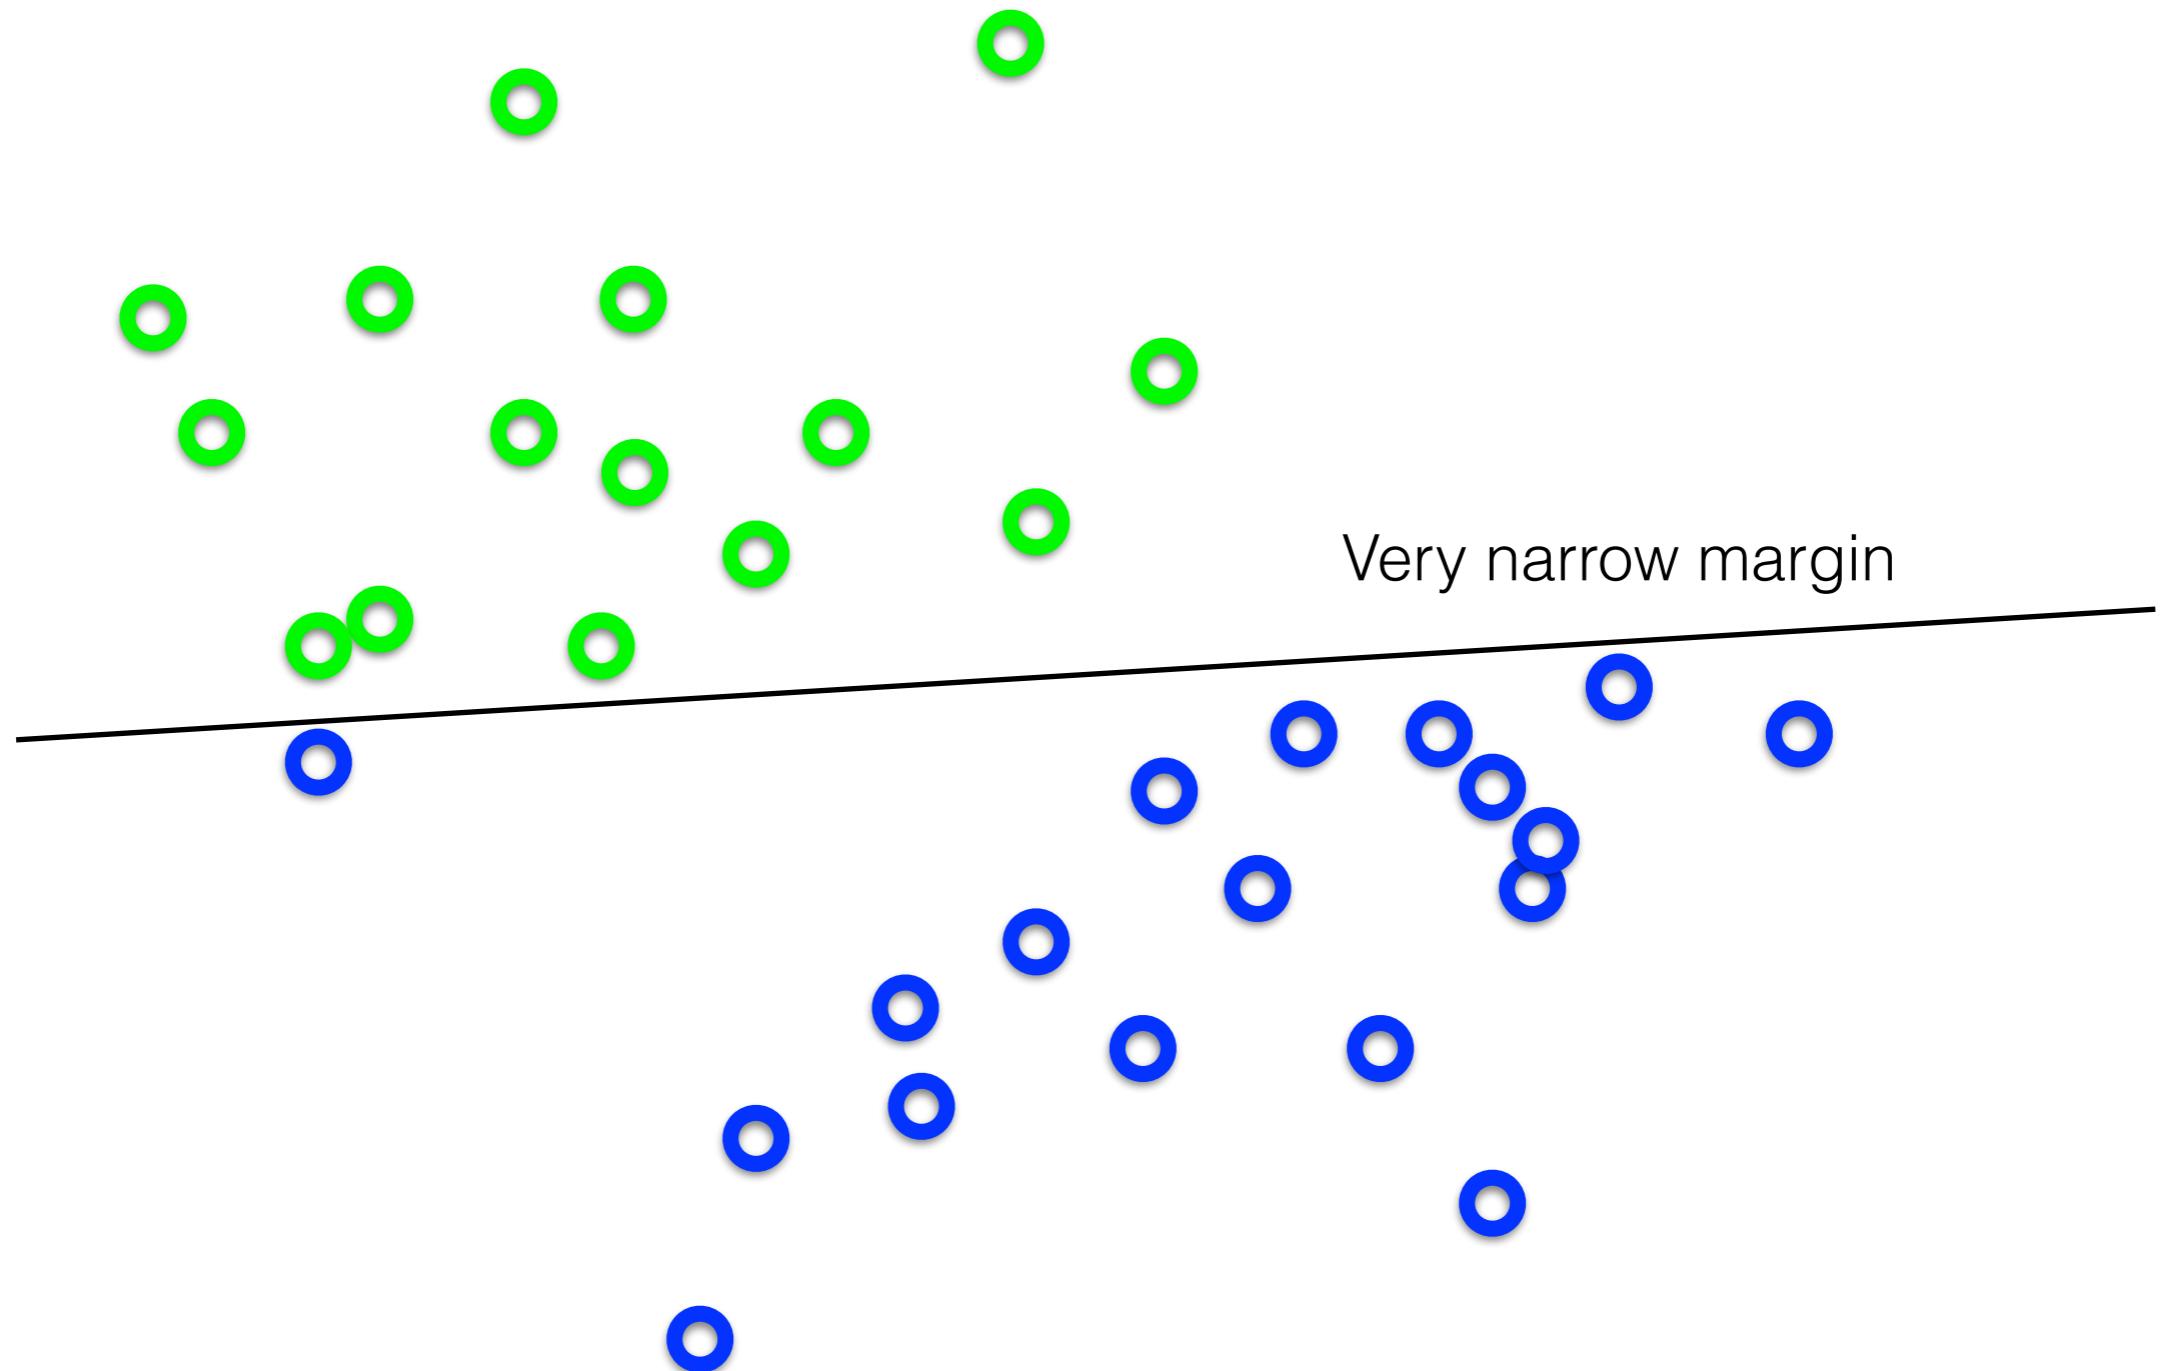
This is a convex quadratic programming (QP) problem
(a unique solution exists)

‘soft’ margin

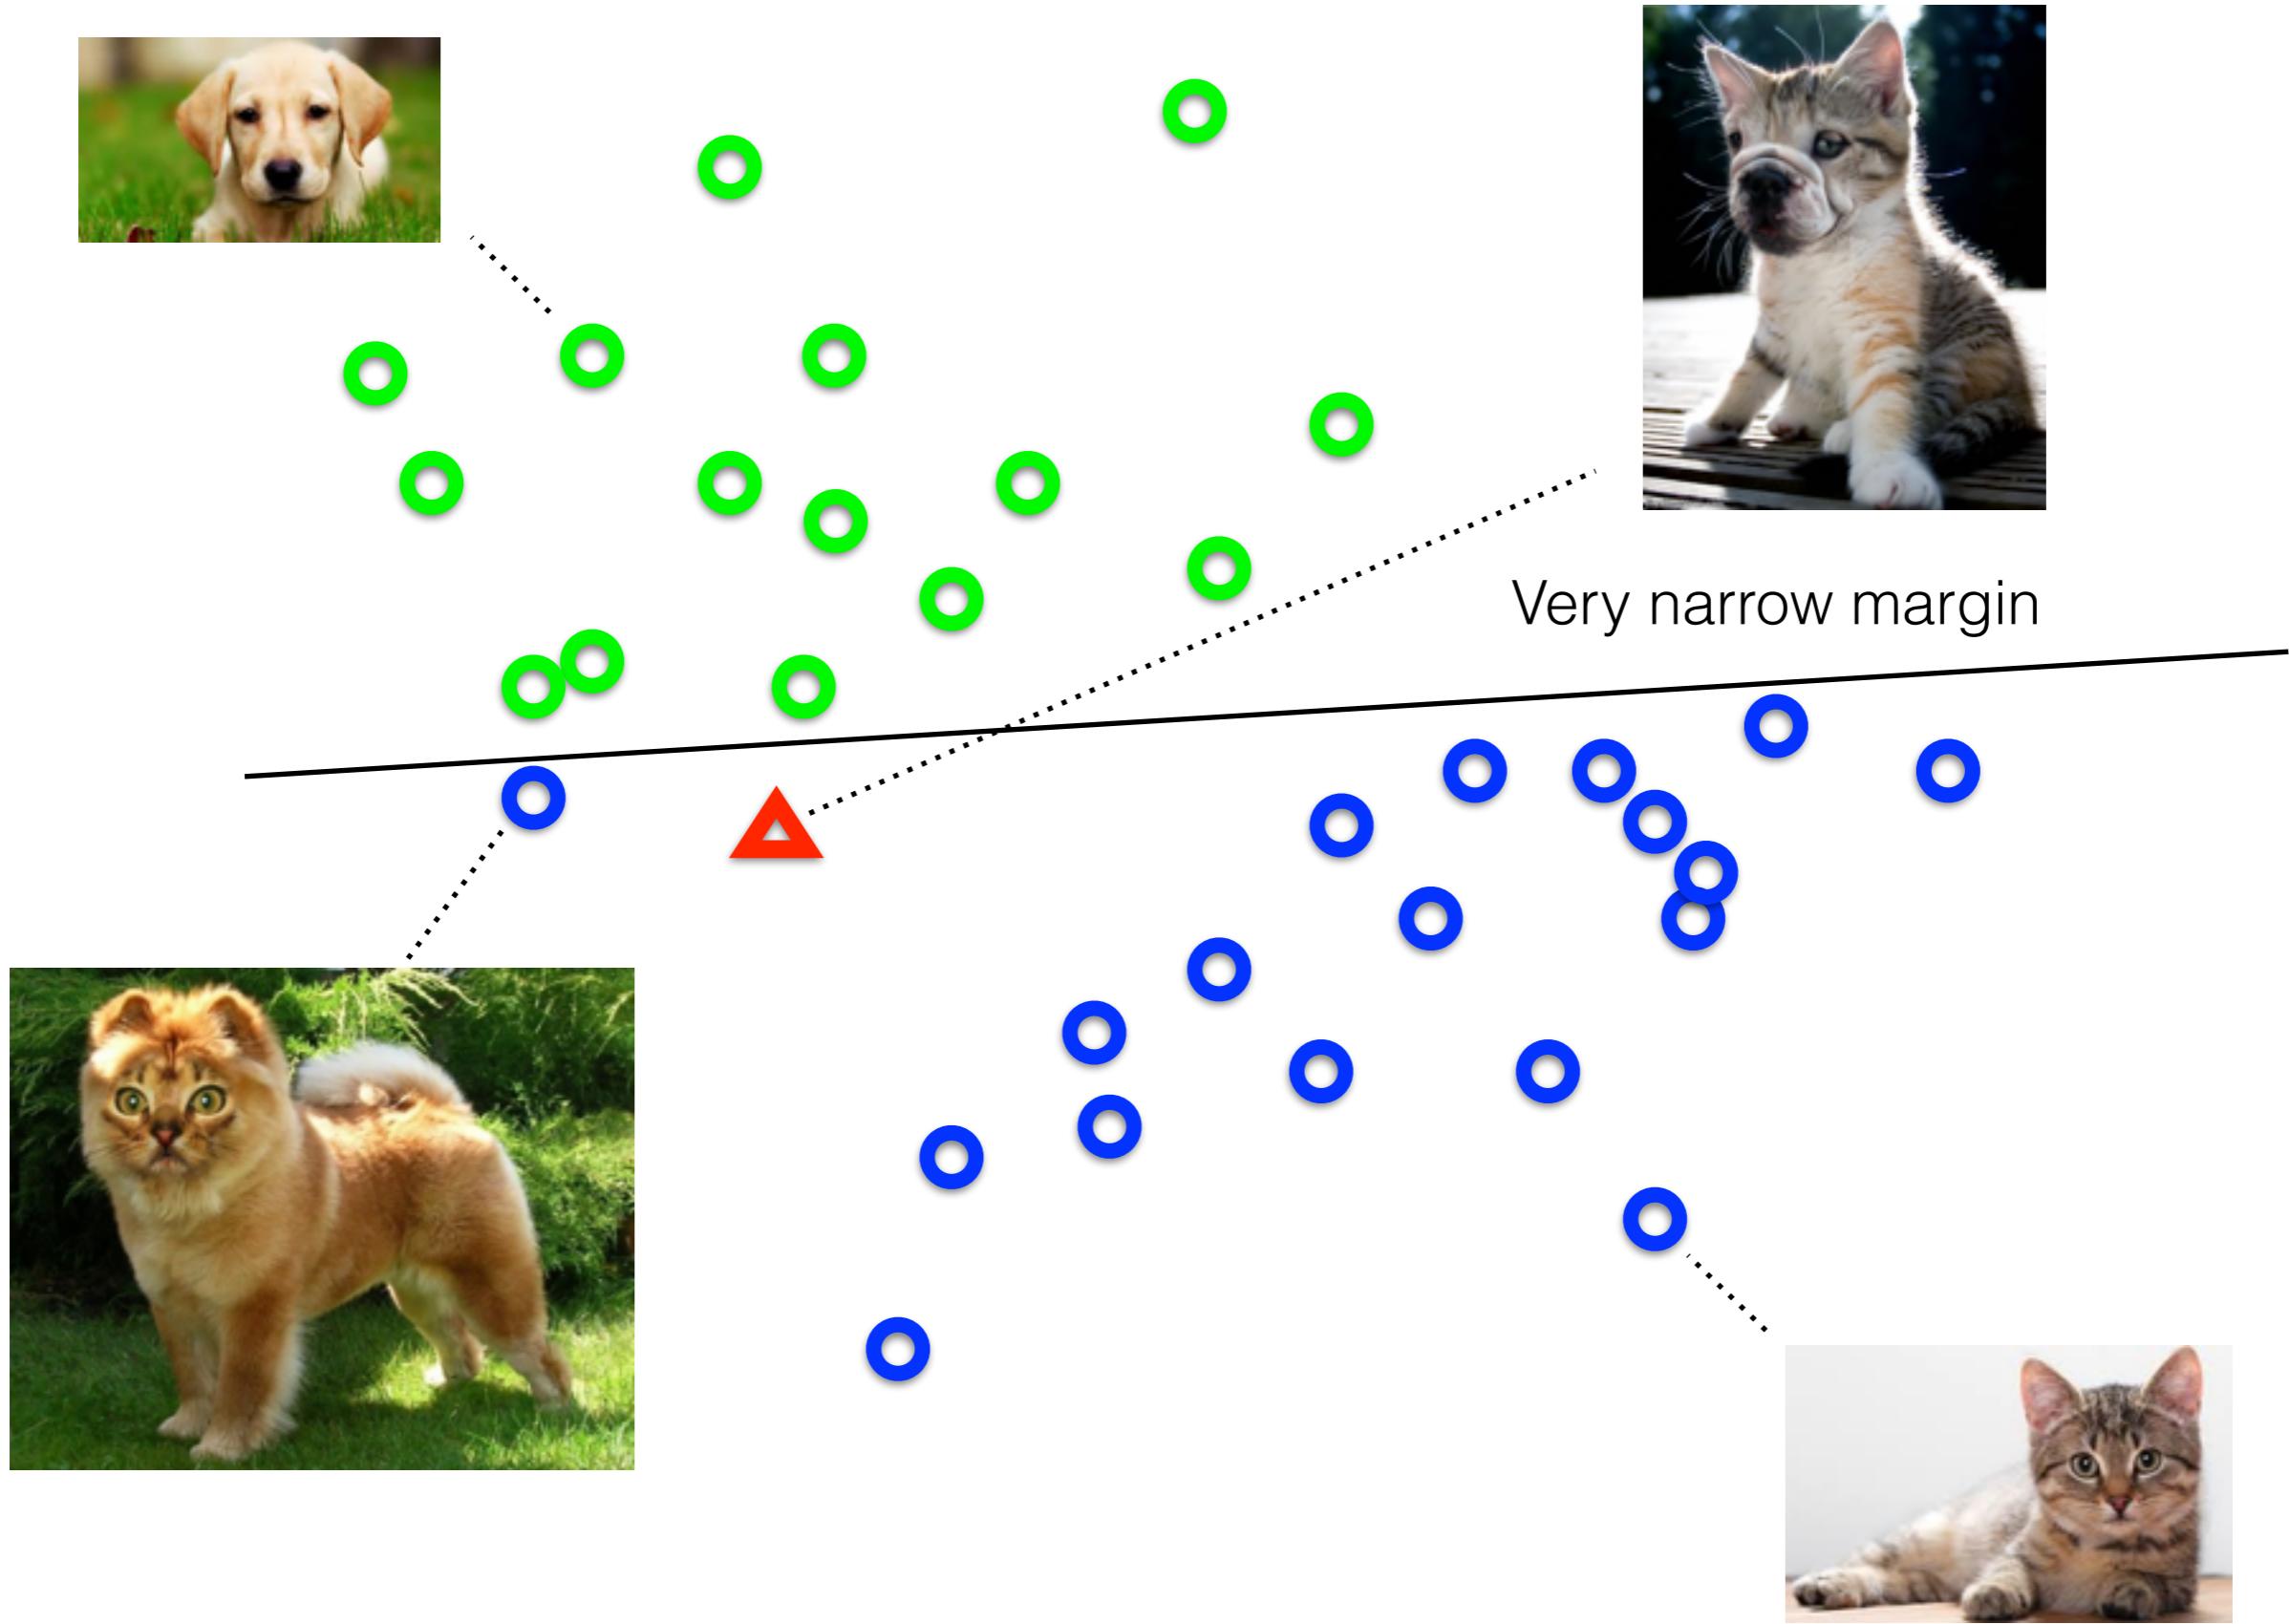
What's the best **w**?



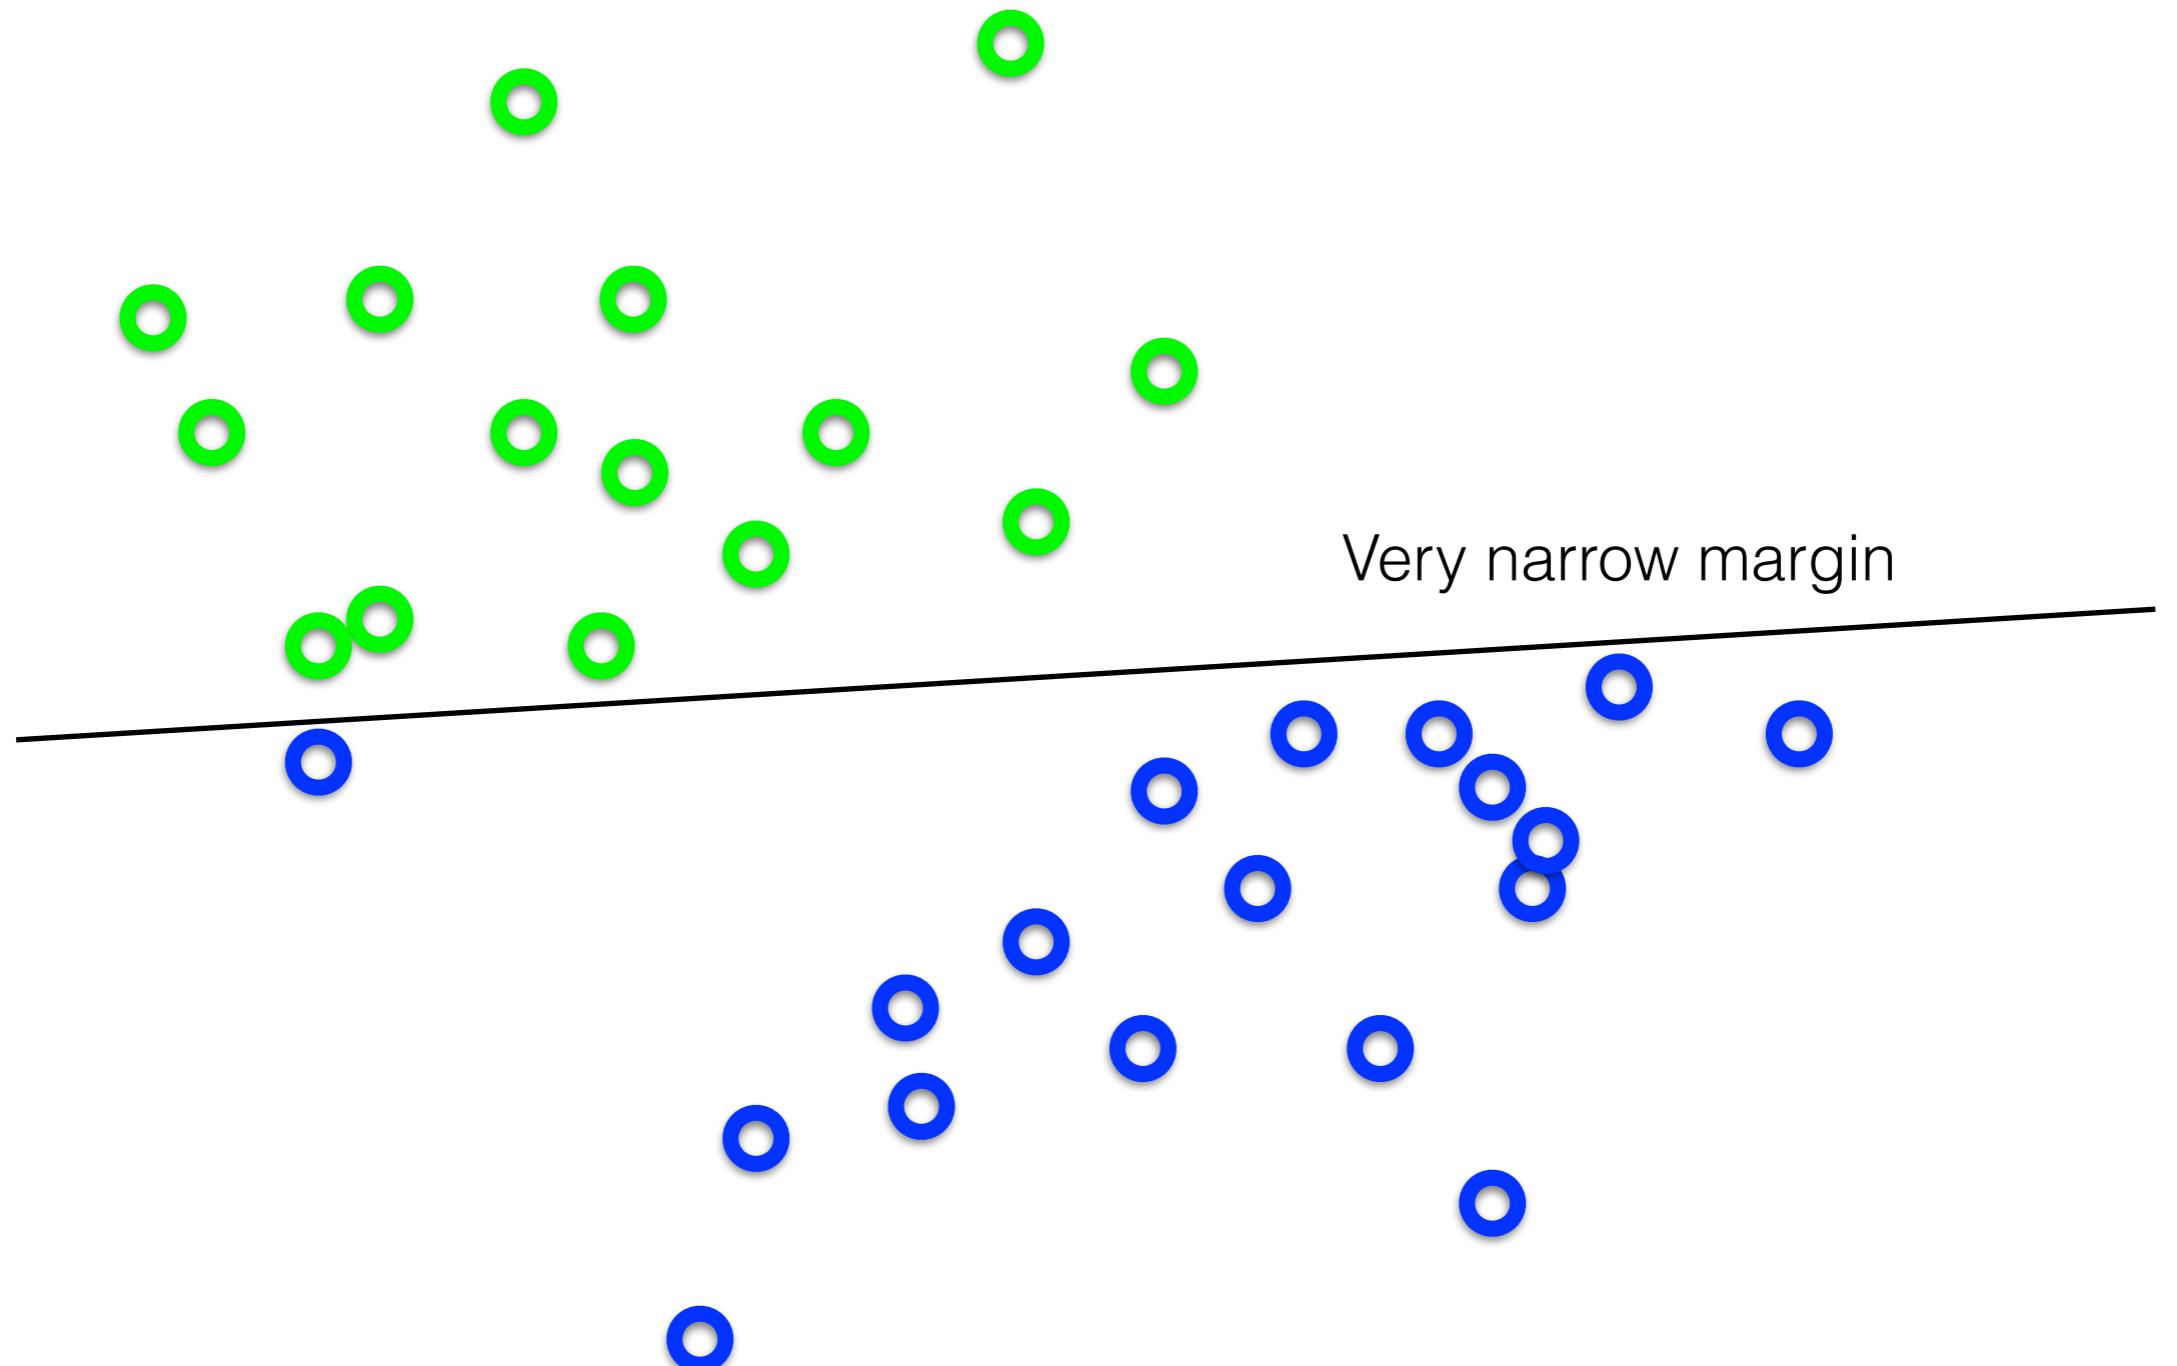
What's the best \mathbf{w} ?



Separating cats and dogs

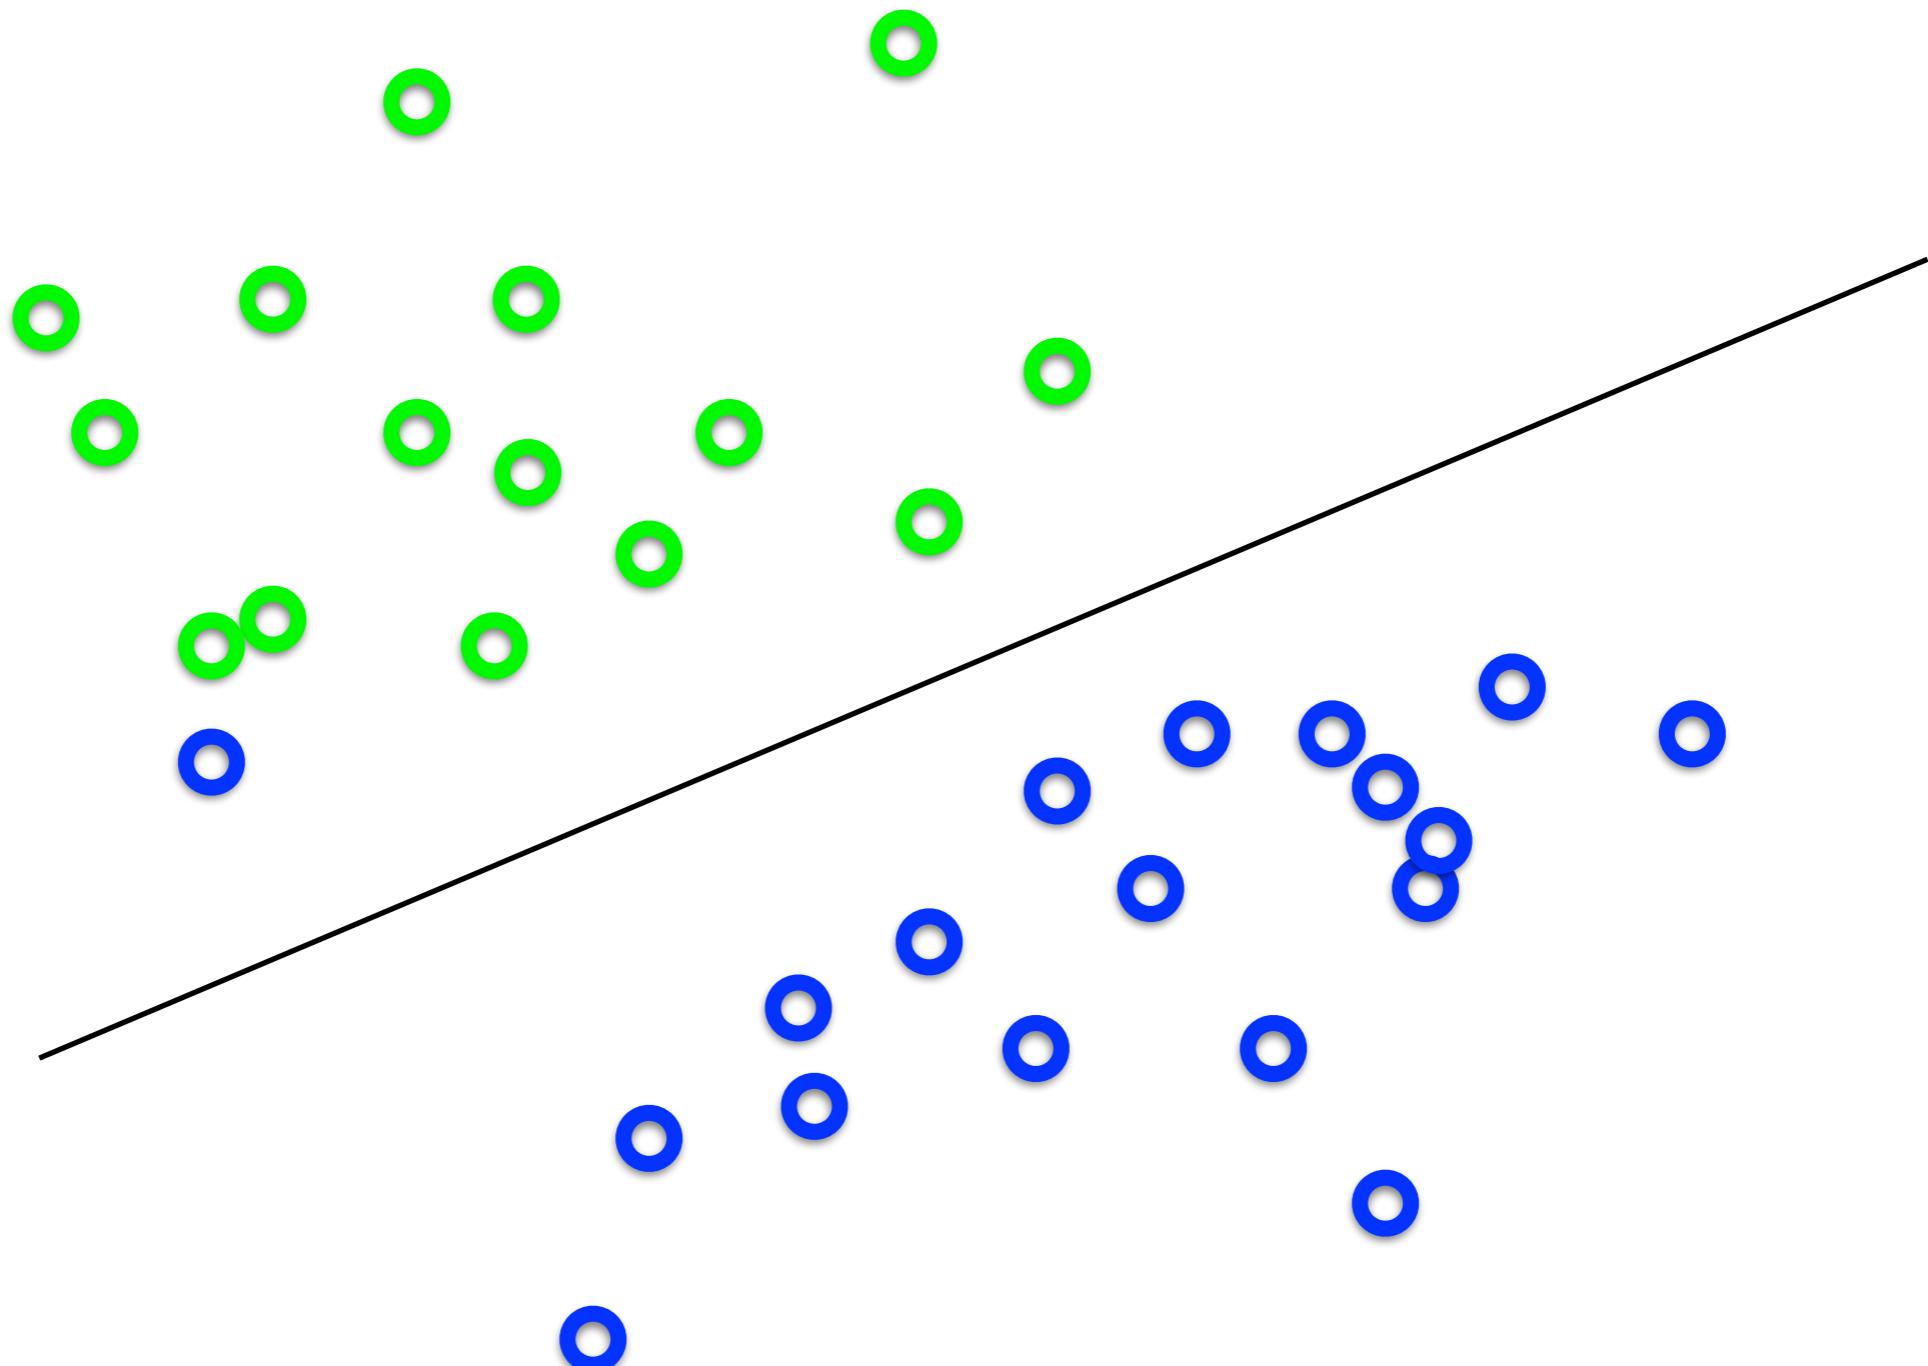


What's the best \mathbf{w} ?



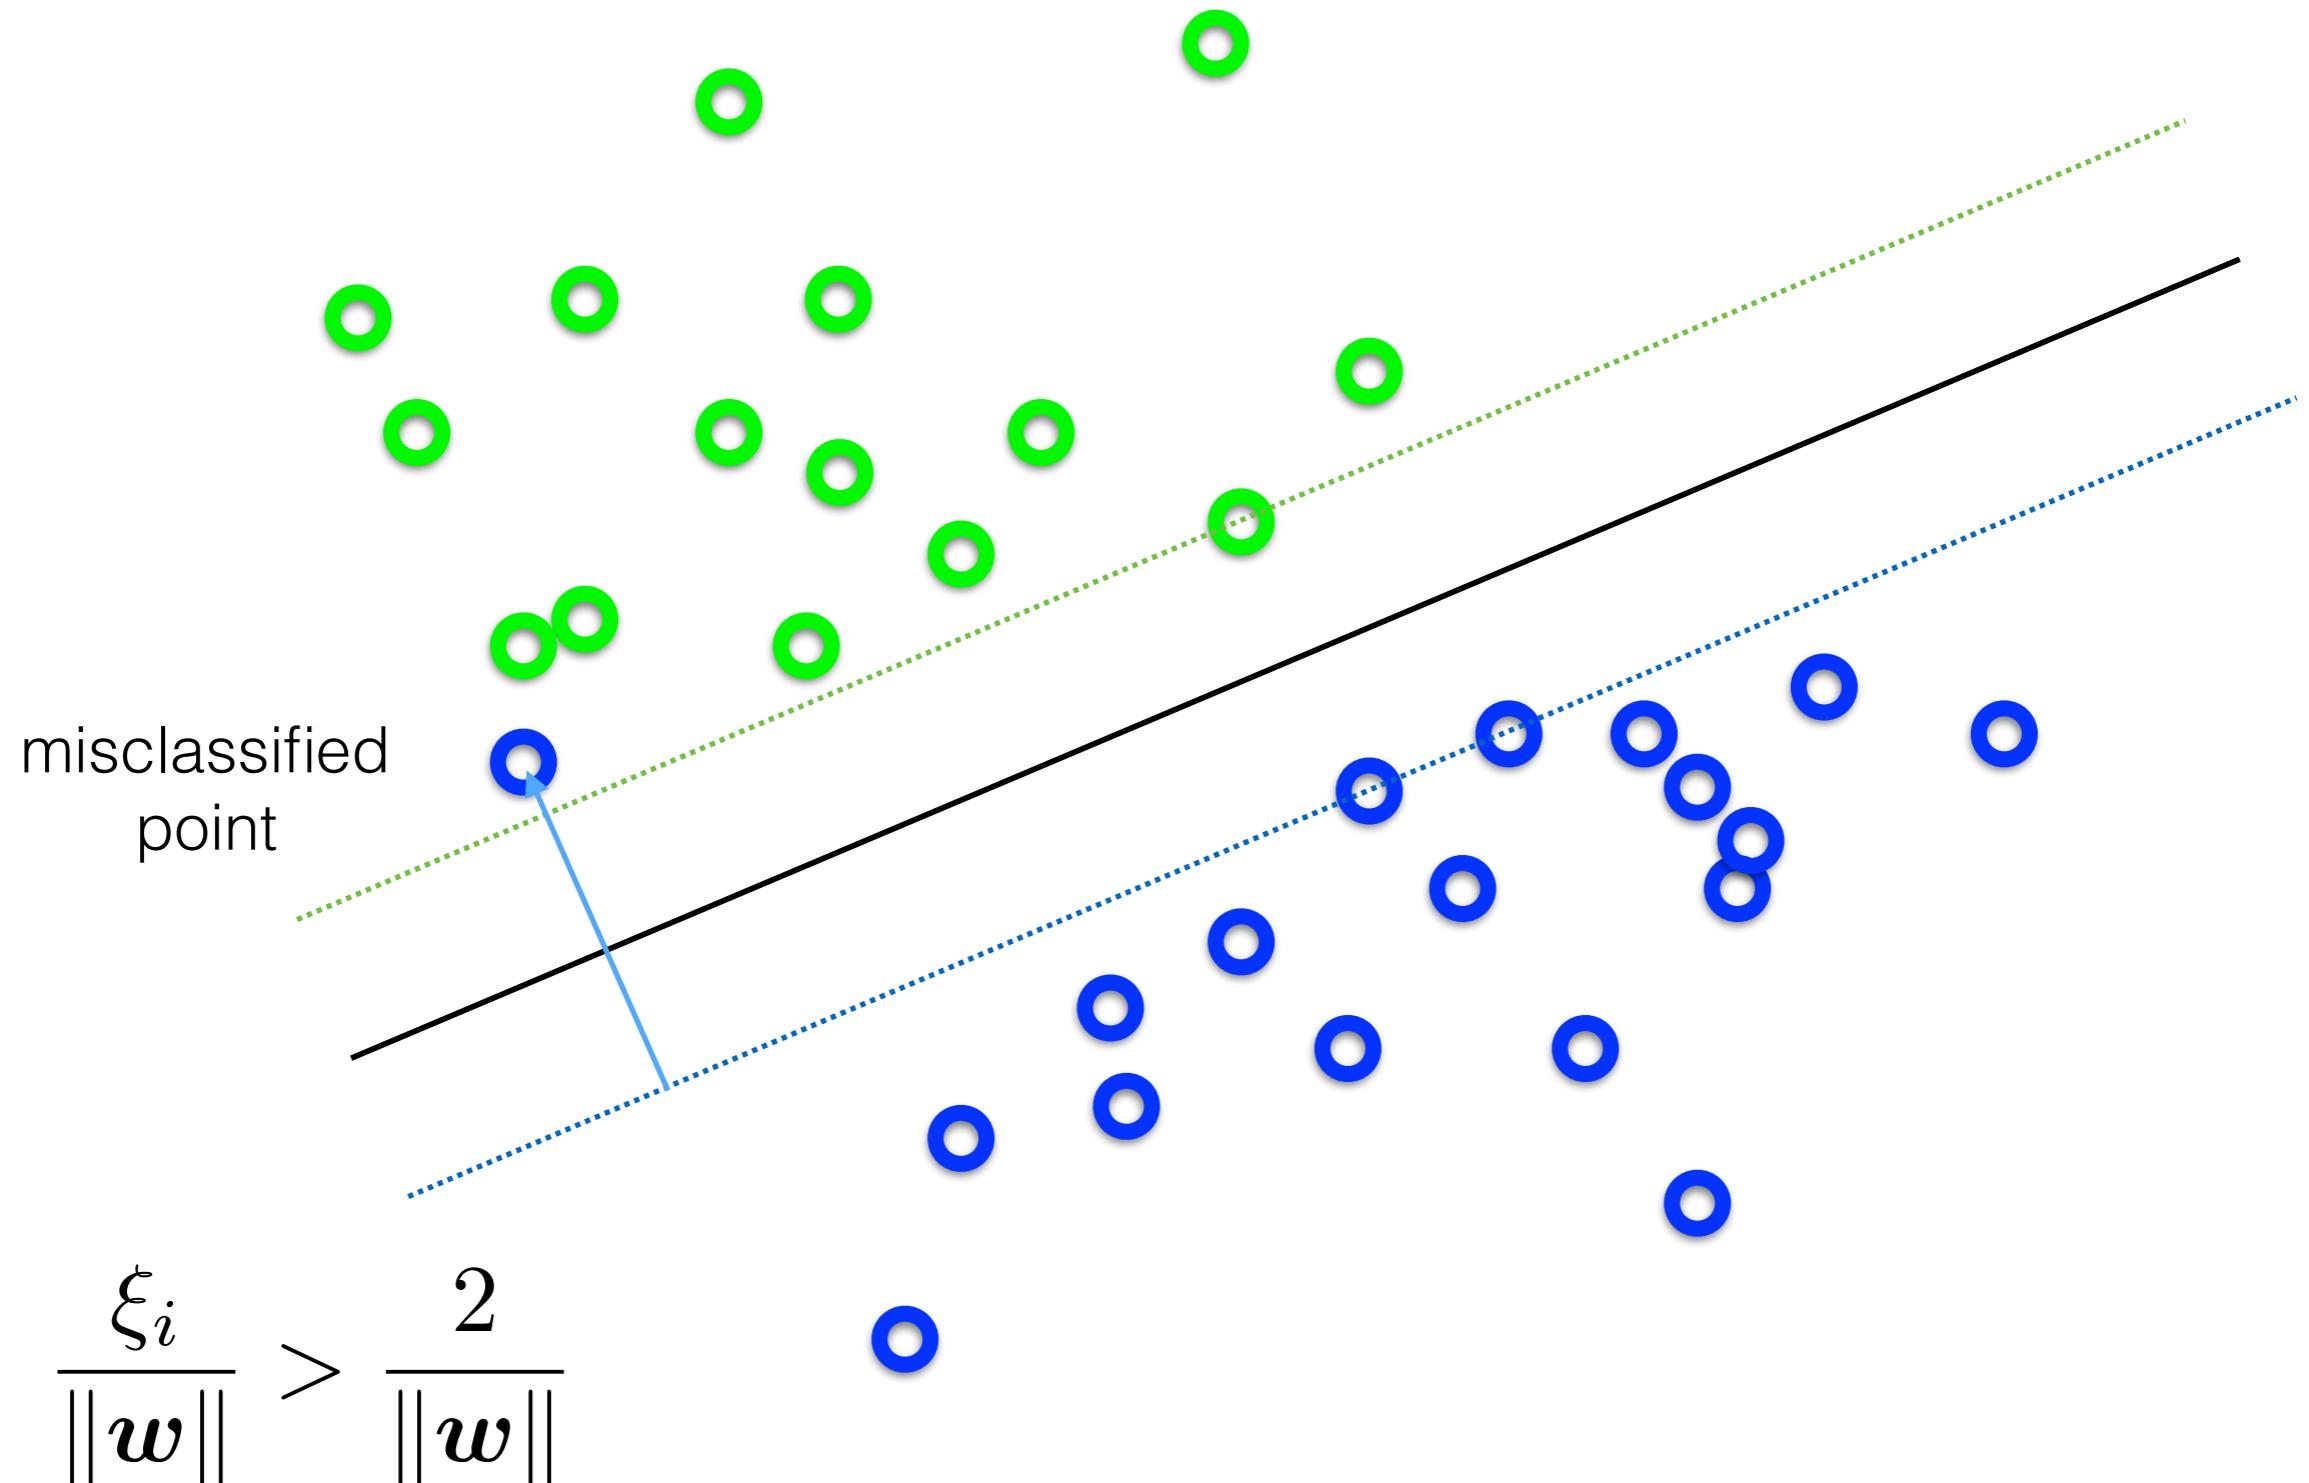
Intuitively, we should allow for some misclassification if we can get more robust classification

What's the best **w**?



Trade-off between the MARGIN and the MISTAKES
(might be a better solution)

Adding slack variables $\xi_i \geq 0$



‘soft’ margin

objective

$$\min_{\boldsymbol{w}, \boldsymbol{\xi}} \|\boldsymbol{w}\|^2 + C \sum_i \xi_i$$

subject to

$$y_i(\boldsymbol{w}^\top \boldsymbol{x}_i + b) \geq 1 - \xi_i$$

for $i = 1, \dots, N$

'soft' margin

objective

$$\min_{\mathbf{w}, \xi} \|\mathbf{w}\|^2 + C \sum_i \xi_i$$

subject to

$$y_i(\mathbf{w}^\top \mathbf{x}_i + b) \geq 1 - \xi_i \quad \text{for } i = 1, \dots, N$$

The slack variable allows for mistakes,
as long as the inverse margin is minimized.

‘soft’ margin

objective

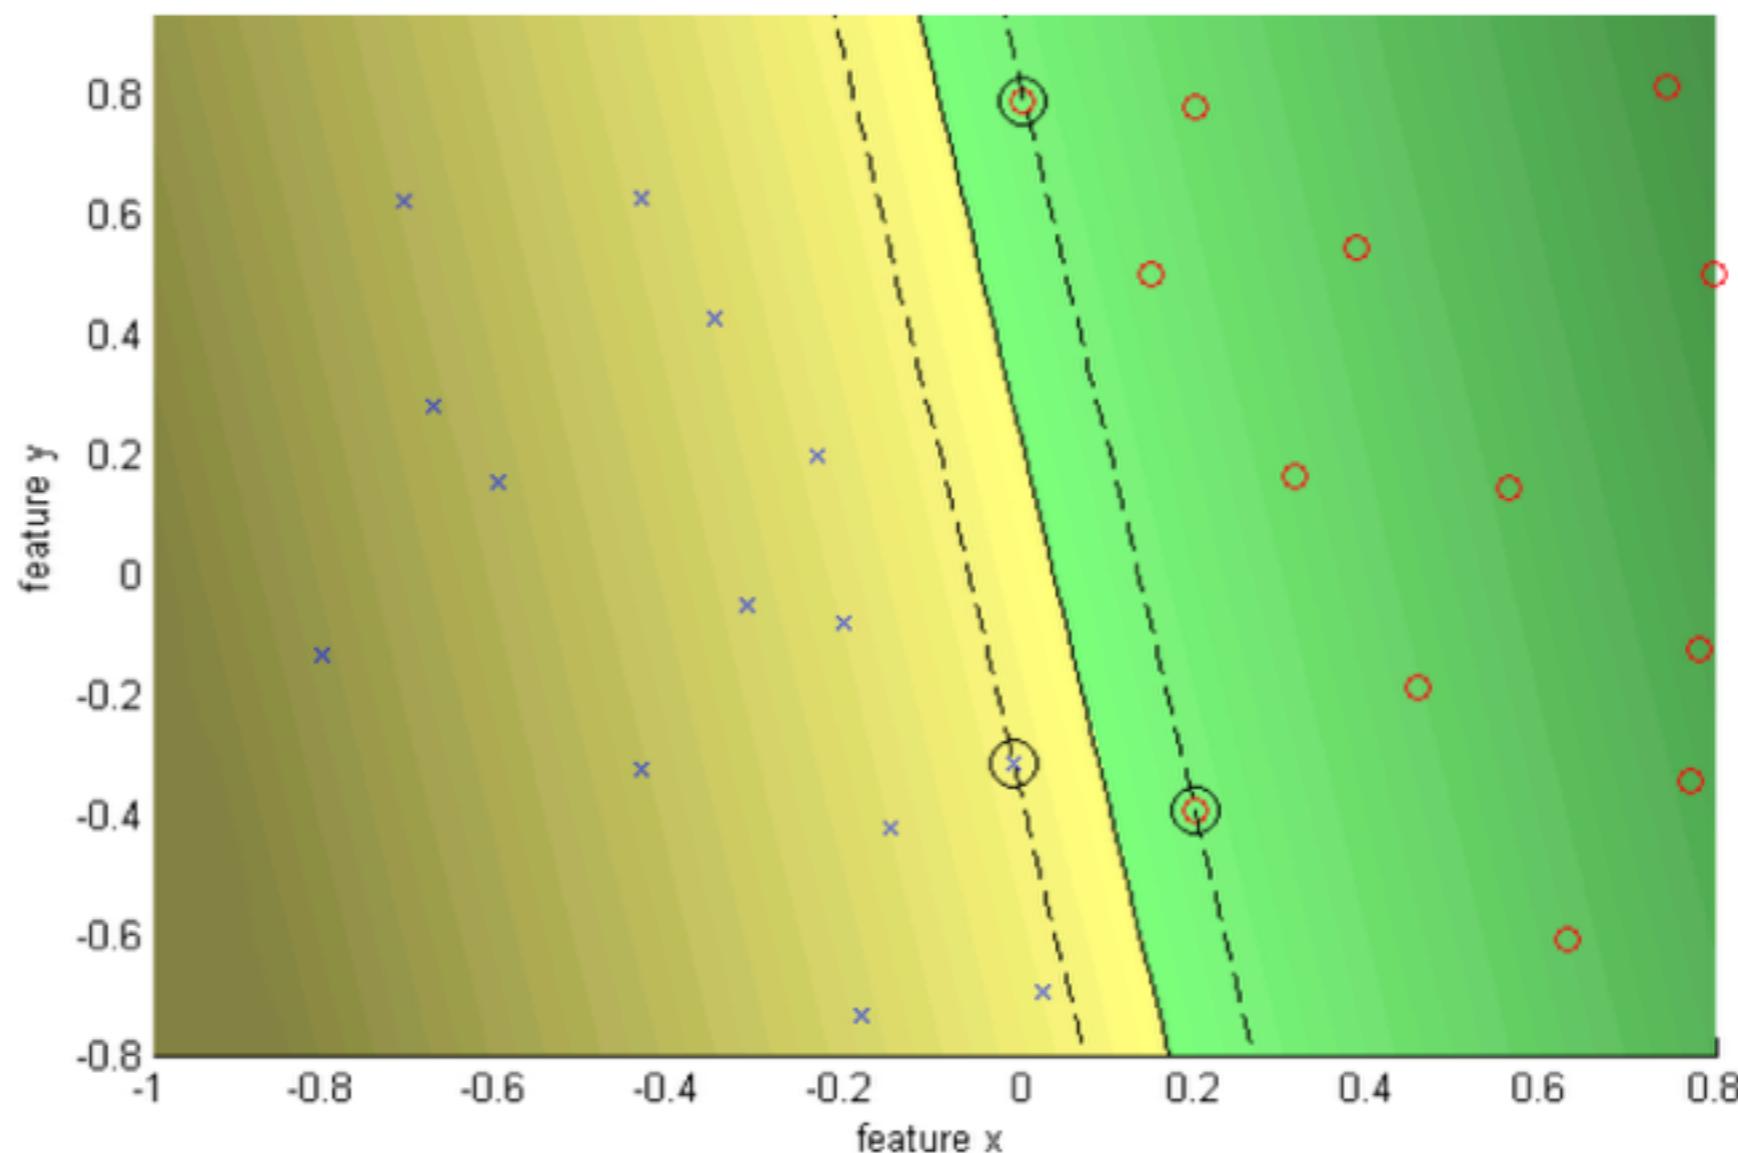
$$\min_{\boldsymbol{w}, \boldsymbol{\xi}} \|\boldsymbol{w}\|^2 + C \sum_i \xi_i$$

subject to

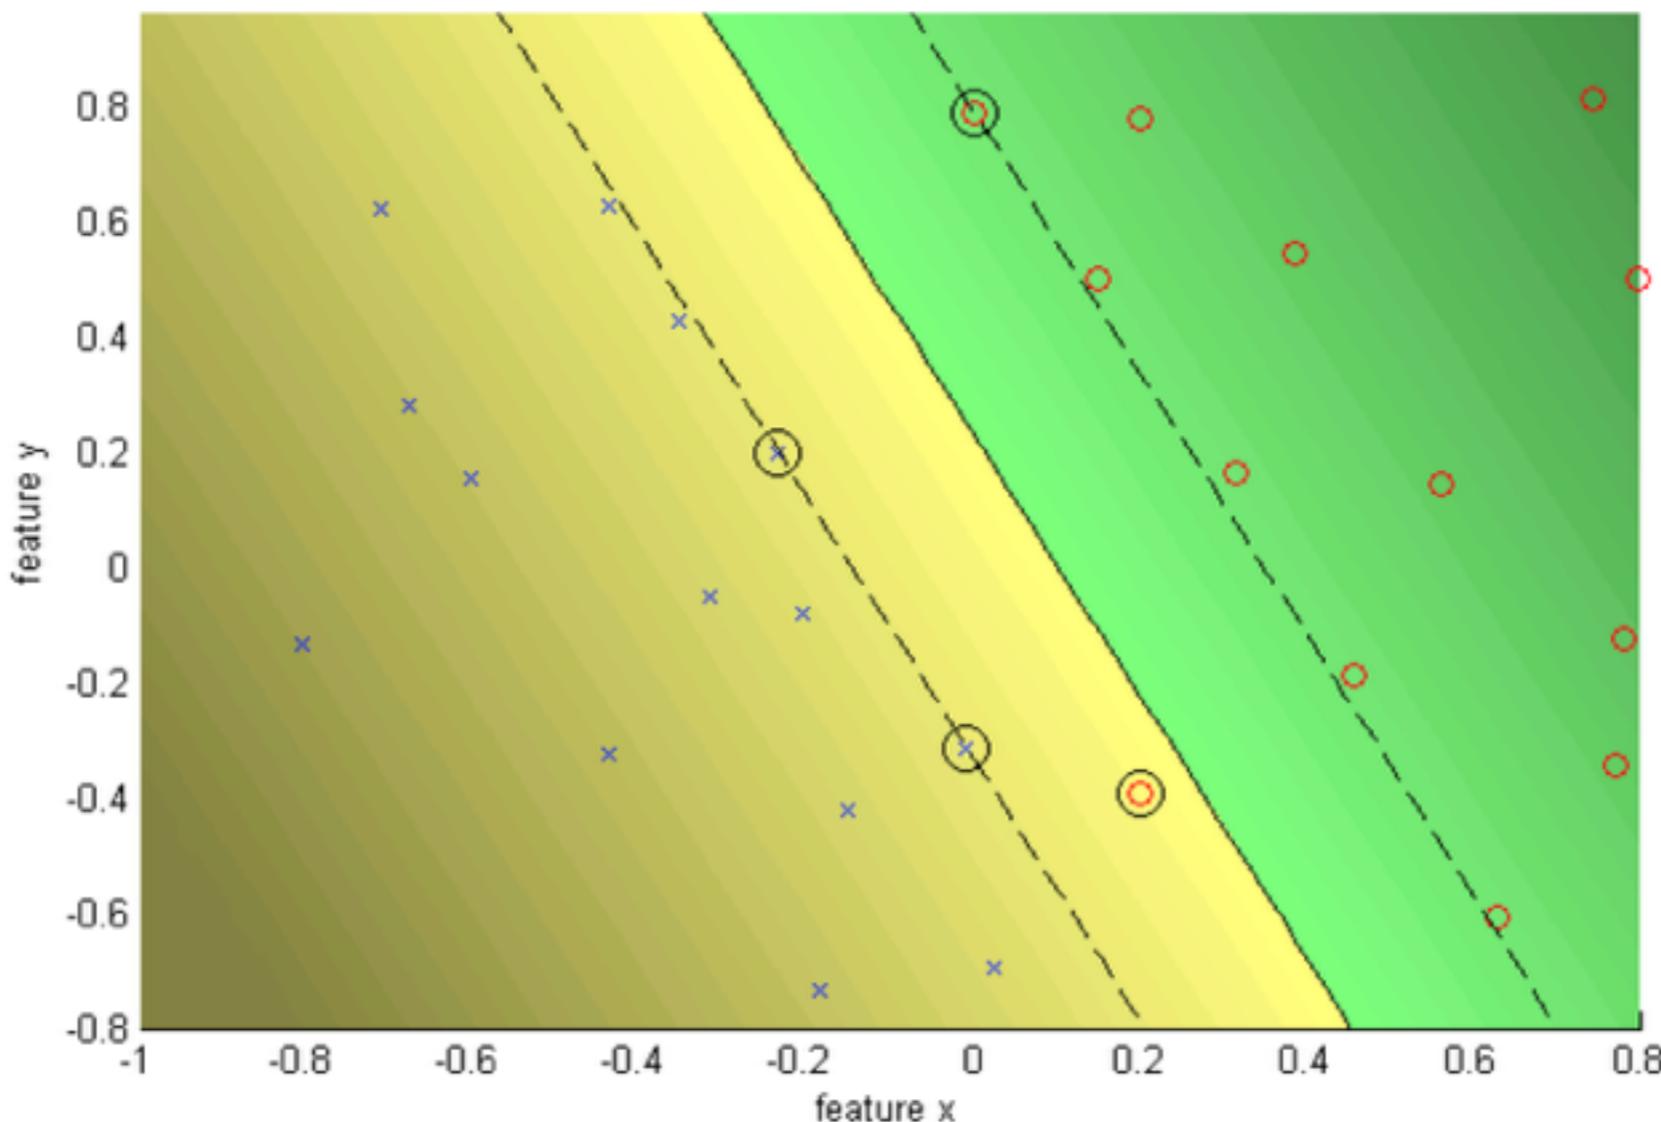
$$y_i(\boldsymbol{w}^\top \boldsymbol{x}_i + b) \geq 1 - \xi_i \quad \text{for } i = 1, \dots, N$$

- Every constraint can be satisfied if slack is large
- C is a regularization parameter
 - Small C: ignore constraints (larger margin)
 - Big C: constraints (small margin)
- Still QP problem (unique solution)

$C = \text{Infinity}$ hard margin



C = 10 soft margin



Comment Window

SVM (L1) by Sequential Minimal Optimizer
Kernel: linear (-), C: 10.0000
Kernel evaluations: 2645
Number of Support Vectors: 4
Margin: 0.2265
Training error: 3.70%